

Software transformations to reduce instruction memory power consumption using a loop buffer

Tom Vander Aa
Murali Jayapala
Francisco Barat
Geert Deconinck

Henk Corporaal
Francky Catthoor

ESAT,
K.U.Leuven, Belgium


IMEC, Leuven,
Belgium

Overview

- Introduction and motivation
 - ☐ Low power
 - ☐ Multimedia applications and the loop buffer
- Optimization problem
 - ☐ Map loops onto the loop buffer
 - ☐ Energy function
- Software transformations
 - ☐ Goal
 - ☐ Examples and results
- Conclusions and future work
 - ☐ Related issues

Context

Low Energy Embedded systems



Low Power Embedded Systems

- ☐ Battery operated (low energy)
 - ⇒ 10-50 MOPS/mW
- ☐ Small
- ☐ Low cost
- ☐ Flexible
- ☐ Multimedia Applications
 - ⇒ Video, audio, wireless
 - ⇒ High performance
 - ✓ 10-100 GOPS
 - ✓ real-time constraints

Context

Embedded systems: Programmable Processor Based

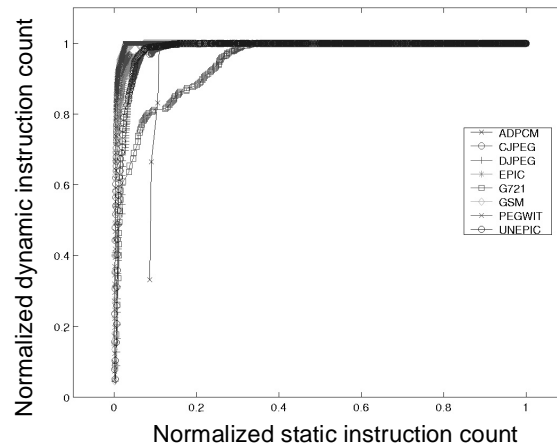


Embedded processors

- Power Breakdown
 - ☐ 43 % of power in on-chip Memory
 - ⇒ StrongARM SA110: A 160MHz 32b 0.5W CMOS ARM processor
 - ☐ 40 % of power in internal memory
 - ⇒ C6x, Texas Instruments Inc.

25-30% of power in Instruction Memory

Instruction Behavior




⇒ Within a program, few basic blocks or instructions take up most of the execution time ($IC_{dynamic}$)

5

Optimisations for DSP and embedded systems, March 2003

Loop Buffer




R.S. Bajwa et.al, "Instruction Buffering to Reduce Power in Processors for Signal Processing", IEEE Trans VLSI Systems, vol 5, no 4, 1997

L. H. Lee et.al, (M-CORE), "Instruction Fetch Energy Reduction Using Loop Caches for Applications with Small and Tight Loops", ISLPED 1999

- L0 Buffer: Buffer (< 1KB) + Local Controller (LC); [no tags]
- L0 / L1 access latency: 1 cycle
- Used only for *specific program segments* (innermost loops)
- Software control:

Special instruction (*lbon*, *sbb*) to map program segments to L0 buffer

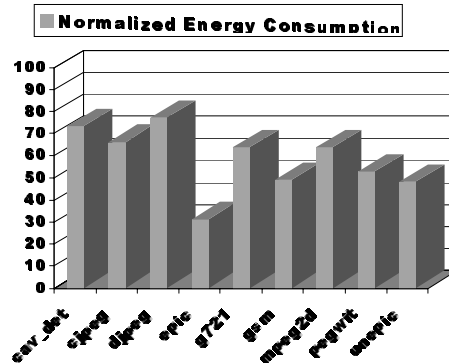


6

Optimisations for DSP and embedded systems, March 2003


Software controlled L0 buffers

- Assumed Architecture
 - MIPS 4000 ISA
 - Single Issue Processor
 - L1 Cache
 - 16KB Direct Mapped
 - Loop Buffer (2KB)
 - Depth = 128 instructions
 - Width = 16 Bytes
- Tools
 - SimpleScalar 2.0
 - Wattch Power estimator



- Loops with less than 128 instructions were hand-mapped onto the loop buffer

Energy Behaviour

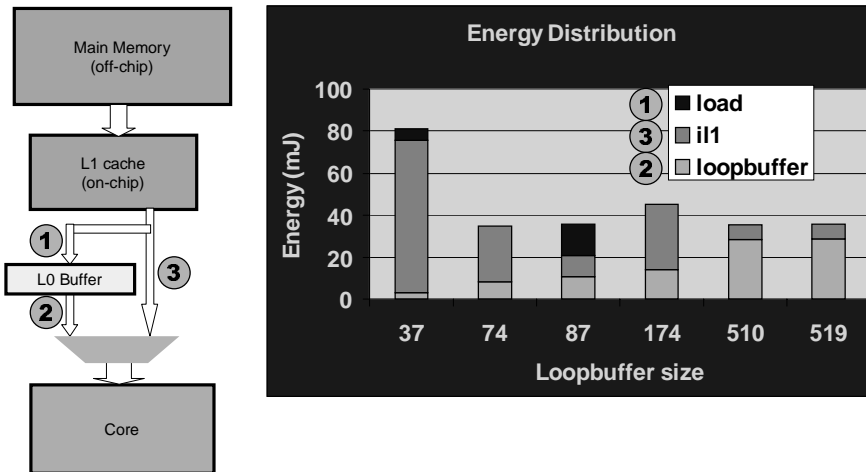


- Energy of the instruction memory hierarchy (lb and il1)

$$E = \sum_{l \in \text{mapped}} N_{\text{loads}}(l) \times E_{\text{access}}(il1) + N_{\text{exec}}(l) \times E_{\text{access}}(lb) + \sum_{l \in \text{unmapped}} N_{\text{exec}}(l) \times E_{\text{access}}(il1)$$


- $E_{\text{access}}(lb) < E_{\text{access}}(il1)$
- $E_{\text{access}}(lb)$ depends on the mapped loops
- $N_{\text{loads}}(l)$ depends on the Control Flow Graph

Energy Behaviour



Code Transformations


- goal: reduce energy
 - ❑ reduce the loop sizes
 - ❑ reduce accesses to il1
- to reduce the active code size
 - ⇒ loop peeling
 - ⇒ loop splitting
 - ⇒ code hoisting
 - ⇒ loop creation
- to reduce loads from l1 cache
 - ⇒ factorisation
 - ⇒ loop interchange
- Classical code optimisations borrowed from HPC



Loop Peeling (I)

- Loop peeling
 - ❑ Removes boundary cases
- Benefits
 - ❑ Reduces the loop size
 - ❑ Reduces the needed loop buffer size
- Example:
 - ❑ Cavity Detector
 - ⇒ Medical imaging application

Boundary cases




```
FOR i = 0 to no_of_rows
  FOR j = 0 to no_of_columns
    IF i < 10 THEN
      /* special case */
    END IF
    /* normal case */
  END FOR
END FOR
```

```
FOR i = 0 to 9
  FOR j = 0 to no_of_columns
    /* special case */
  END FOR
END FOR
FOR i = 10 to no_of_rows
  FOR j = 0 to no_of_columns
    /* normal case */
  END FOR
END FOR
```

12

Optimisations for DSP and embedded systems, March 2003

Loop Peeling (II)




13

Optimisations for DSP and embedded systems, March 2003


Loop Rerolling (I)

- Make trade off between
 - ❑ Software pipelining
 - ⇒ No code duplication inside the loop
 - ⇒ Limited by recurring dependencies
 - ❑ Loop Unrolling
 - ⇒ Code size explosion
 - ⇒ Better ILP, if recurring dependencies



- Benefits
 - ❑ Reduces the loop buffer size
 - ❑ Reduces $E_{\text{access}}(\text{lb})$
- Example
 - ❑ Blowfish encryption

Loop Rerolling




Factorization (I)

- Factorization
 - Finds common code in the loop buffer
 - ⇒ Make them equal with if/then/else constructs
 - ⇒ Create a function call
- Benefits
 - Reduces $N_{\text{loads}}(I)$
- Example
 - 3D Image Reconstruction Algorithm



Factorization (II)



Related Work

- Data memory hierarchy power
 - ⇒ Should also be handled
 - ⇒ Methods and tools are known (DTSE)
- Compiler support for loop buffers
 - ❑ Code layout [Bellas, et al]
 - ❑ Complementary transformation [Sias, et al]
- Impact of the transformations
 - ❑ On data memory
 - ⇒ Unchanged, since accesses are not reordered (yet)
 - ❑ On processor performance
 - ❑ Evaluation framework
 - ⇒ Enhanced VLIW compiler
 - ⇒ Based on Trimaran framework
 - ⇒ Wattch power models

Conclusion

- L0 Buffer Organization
 - ❑ Multimedia applications have high locality in small program segments
 - ❑ An additional small L0 buffer should be used
 - ❑ Standard implementations: loop buffer not efficiently used
- Software transformations are needed
 - ❑ Loop peeling
 - ❑ Loop rerolling
 - ❑ Factorisation
 - ❑ Loop splitting
- Up to 80% of energy reduction in the instruction memory hierarchy
 - ❑ Evaluation framework
- What's next?
 - ❑ Automation of the transformations