
ARM ISA SIMD Extensions

John.Rayfield@arm.com

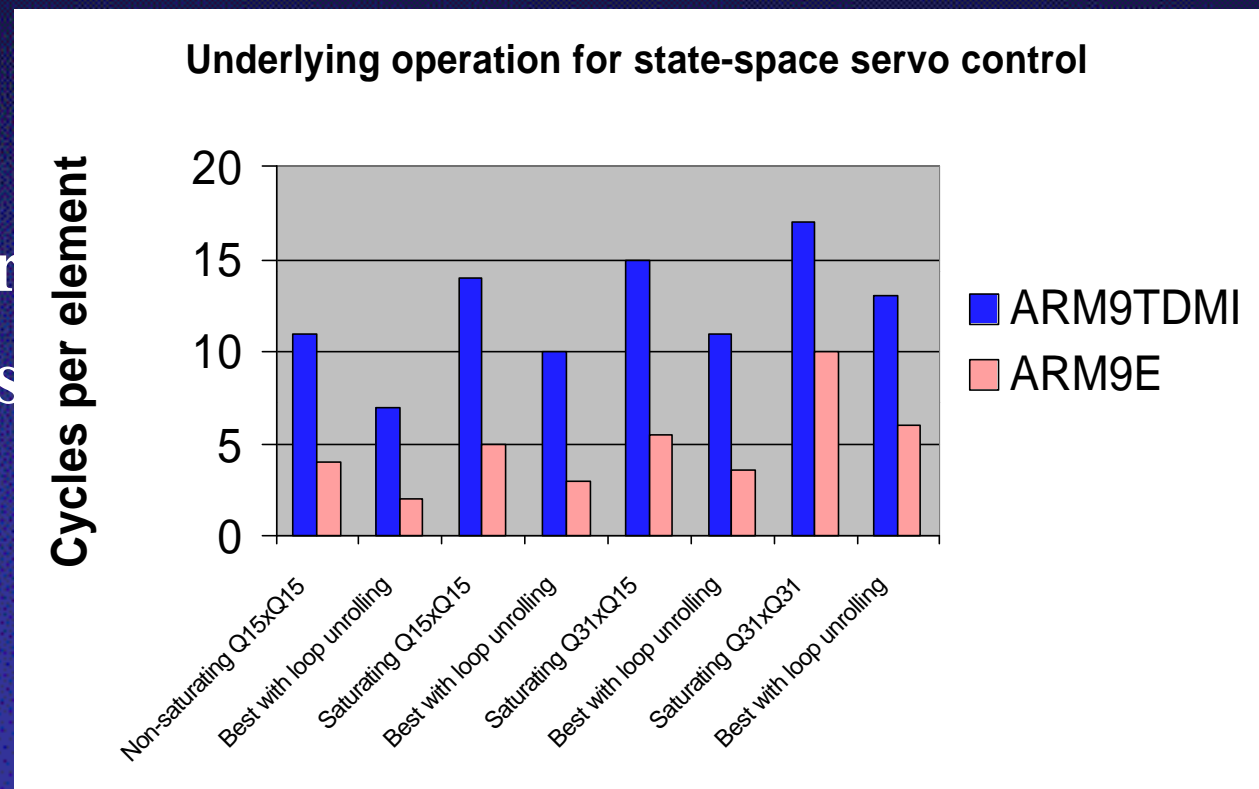
www.arm.com

Markets


- ARM implementations are increasingly being used in a broad range of embedded and platform computing applications:
 - Wireless, mass-storage, automotive, industrial, digital imaging, secure, consumer entertainment
- Many of these are increasing the need for signal processing capabilities in the base architecture

The ARM ISA v5TE extensions

- ARM added considerable DSP functionality into the architecture with the introduction of ARM ISA V5TE
- More than 2x performance for some algorithms



ARM Architecture roadmap



ARM

THE ARCHITECTURE FOR THE 21ST CENTURY

Examples of new market needs:

- Video capabilities in 3G cellular
 - MPEG4 QCIF decode and encode at low power
- Audio encoders/decoders in PDA
 - MP3 and others
- Automotive collision avoidance systems
 - Radar systems with extensive FFT needs
- All pushing DSP/Media performance
 - and suffer from the 90/10 rule

Why SIMD?

- With short data-types, most of a 32-bit data-path is left computing on sign-extension bits
- SIMD is the most power and area efficient architectural model if you have data regularity
 - 8-bit and 16-bit data can be more efficiently manipulated in the 32-bit pipeline
- Very carefully architected so as not to compromise ARM implementations

Summary of the new instructions

■ SIMD arithmetic

- 16-bit addition and subtraction
- 8-bit addition and subtraction
- Selection

ADD16, SUB16, ADDSUB, SUBADD
ADD8, SUB8
SEL8, SEL16

■ Multiply instructions

- Dual 16-bit multiply
- 32-bit media multiplies


SMLAD, SMLSD, SMLALD, SMLSLD
SMMLA, SMMLAR SMMLS, SMMLSR

■ Flexible saturation

SSAT, USAT

Detail of the new instructions

New SIMD status flags



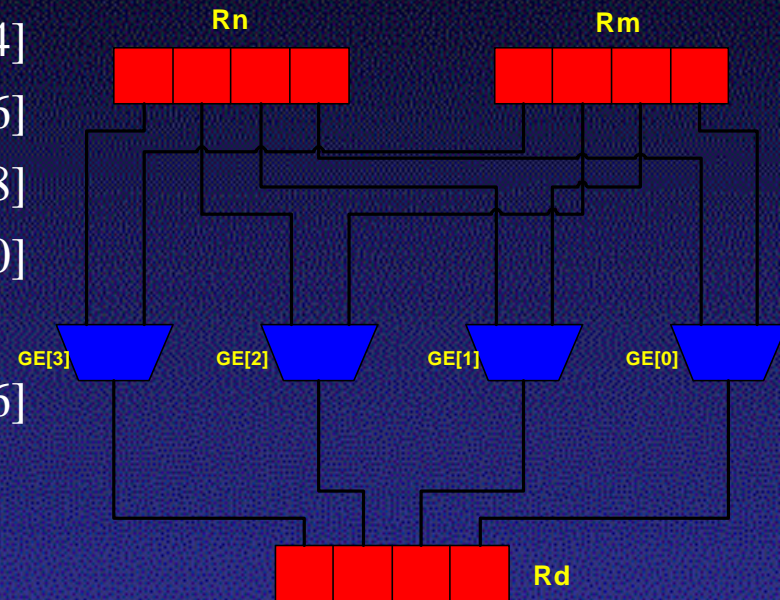
- Four new flags GE[3:0] added into the CPSR
 - GE[3] = ‘greater than or equal’ for bit 31 (8-bit,16-bit SIMD)
 - GE[2] = ‘greater than or equal’ for bit 23 (8-bit SIMD)
 - GE[1] = ‘greater than or equal’ for bit 15 (8-bit,16-bit SIMD)
 - GE[0] = ‘greater than or equal’ for bit 7 (8-bit SIMD)
- Flags affected by all SIMD operations
- Flags used by SIMD select operations

SIMD Add and Subtract

- ADD16 Rd,Rn,Rm, SUB16 Rd,Rn,Rm
- ADDSUB Rd,Rn,Rm, SUBADD Rd,Rn,Rm
- ADD8 Rd,Rn,Rm, SUB8 Rd,Rn,Rm
- Other variants
 - saturating
 - versions that swaps two halves of one source (16-bit)
 - version that halve the results to prevent overflow
 - unsigned versions
- Applications
 - General use in all algorithms using 16-bit precision
 - 16-bit FFT - audio, speech recognition
 - DCT video and still-image coding

SIMD Selection Operations

- SEL8 Rd,Rn,Rm
 - $Rd[31:24] = GE[3] ? Rn[31:24] : Rm[31:24]$
 - $Rd[23:16] = GE[2] ? Rn[23:16] : Rm[23:16]$
 - $Rd[15:08] = GE[1] ? Rn[15:08] : Rm[15:08]$
 - $Rd[07:00] = GE[0] ? Rn[07:00] : Rm[07:00]$
- SEL16 Rd,Rn,Rm
 - $Rd[31:16] = GE[3] ? Rn[31:16] : Rm[31:16]$
 - $Rd[15:00] = GE[1] ? Rn[15:0] : Rm[15:0]$
- Applications
 - 8-bit and 16-bit SIMD
 - minimum, maximum, absolute, Viterbi (ACS)



Dual 16x16 multiply

- Result = [Accumulator +] 16-bits x 16-bits +/- 16-bits x 16-bits

- SMUA{X}D Rd, Rm, Rs
- SMUS{X}D Rd, Rm, Rs
- SMLA{X}D Rd, Rm, Rs, Rn
- SMLS{X}D Rd, Rm, Rs, Rn
- SMLALD Rlo, Rhi, Rm, Rs

$$Rd = Rmlo * Rslo / hi + Rmhi * Rshi / lo$$

$$Rd = Rmlo * Rslo / hi - Rmhi * Rshi / lo$$

$$Rd = Rn + Rmlo * Rslo / hi + Rmhi * Rshi / lo$$

$$Rd = Rn + Rmlo * Rslo / hi - Rmhi * Rshi / lo$$

$$Rhi.Rlo += Rmlo * Rslo + Rmhi * Rshi$$

- Option to swap one 16-bit input operand position
- Option on 32-bit or 64-bit extended accumulation
- Applications
 - 16-bit FFT
 - 16-bit real/complex filters/dot-products with 32-bit accumulator
 - 16-bit filters/dot-products with 64-bit accumulator

32x32 fractional multiplies

■ Result = [Accumulator +/-] ((32-bits x 32-bits + round)>>32)

- SMMLA{R} Rd, Rm, Rs, Rn ;Rd = Rn + ((Rm*Rs + round)>>32)
- SMMLS{R} Rd, Rm, Rs, Rn ;Rd = Rn - ((Rm*Rs + round)>>32)
- SMMUL{R} Rd, Rm, Rs ;Rd = ((Rm*Rs) + round)>>32

■ Applications

- 32-bit FFT, complex filters
- 32-bit real/complex filters with 32-bit accumulator
- 32-bit filters with 64-bit accumulator

Saturation operation

- SSAT Rd,#BitPosition,Rm,{ Shift}
- USAT Rd,#BitPosition,Rm,{ Shift}
- Flexible saturation in signed and unsigned form at arbitrary bit-position in the source word (Rm)
- Applications
 - Graphics, pixel plotting, audio algorithms ISO/ETSI standards

Example uses of the new instructions

16-bit Complex multiply

- ;ARM ISA v5TE implementation

SMULTT Real,Ra,Rb ;Real = Ra.real*Rb.real

SMULBB Temp,Ra,Rb ;Temp = Ra.imag*Rb.imag

SUB Real,Real,Temp ;Real = Ra.real*Rb.real - Ra.imag*Rb.imag

SMULTB Imag,Ra,Rb ;Imag = Ra.real*Rb.imag

SMLABT Imag,Ra,Rb ;Imag = Ra.real*Rb.imag + Ra.imag*Rb.real

;Total of 5-cycle in a single-cycle implementation

- ;New implementation

SMUSD Real,Ra,Rb ;Real = Ra.real*Rb.real - Ra.imag*Rb.imag

SMUADX Imag,Ra,Rb ;Imag = Ra.real*Rb.imag + Ra.imag*Rb.real

- ;Total of 2-cycles in a single cycle implementation

Add-compare-select

- Underlying operation in the Viterbi algorithm

Equalization, convolutional code decoding, recognition engines

- ;New implementation

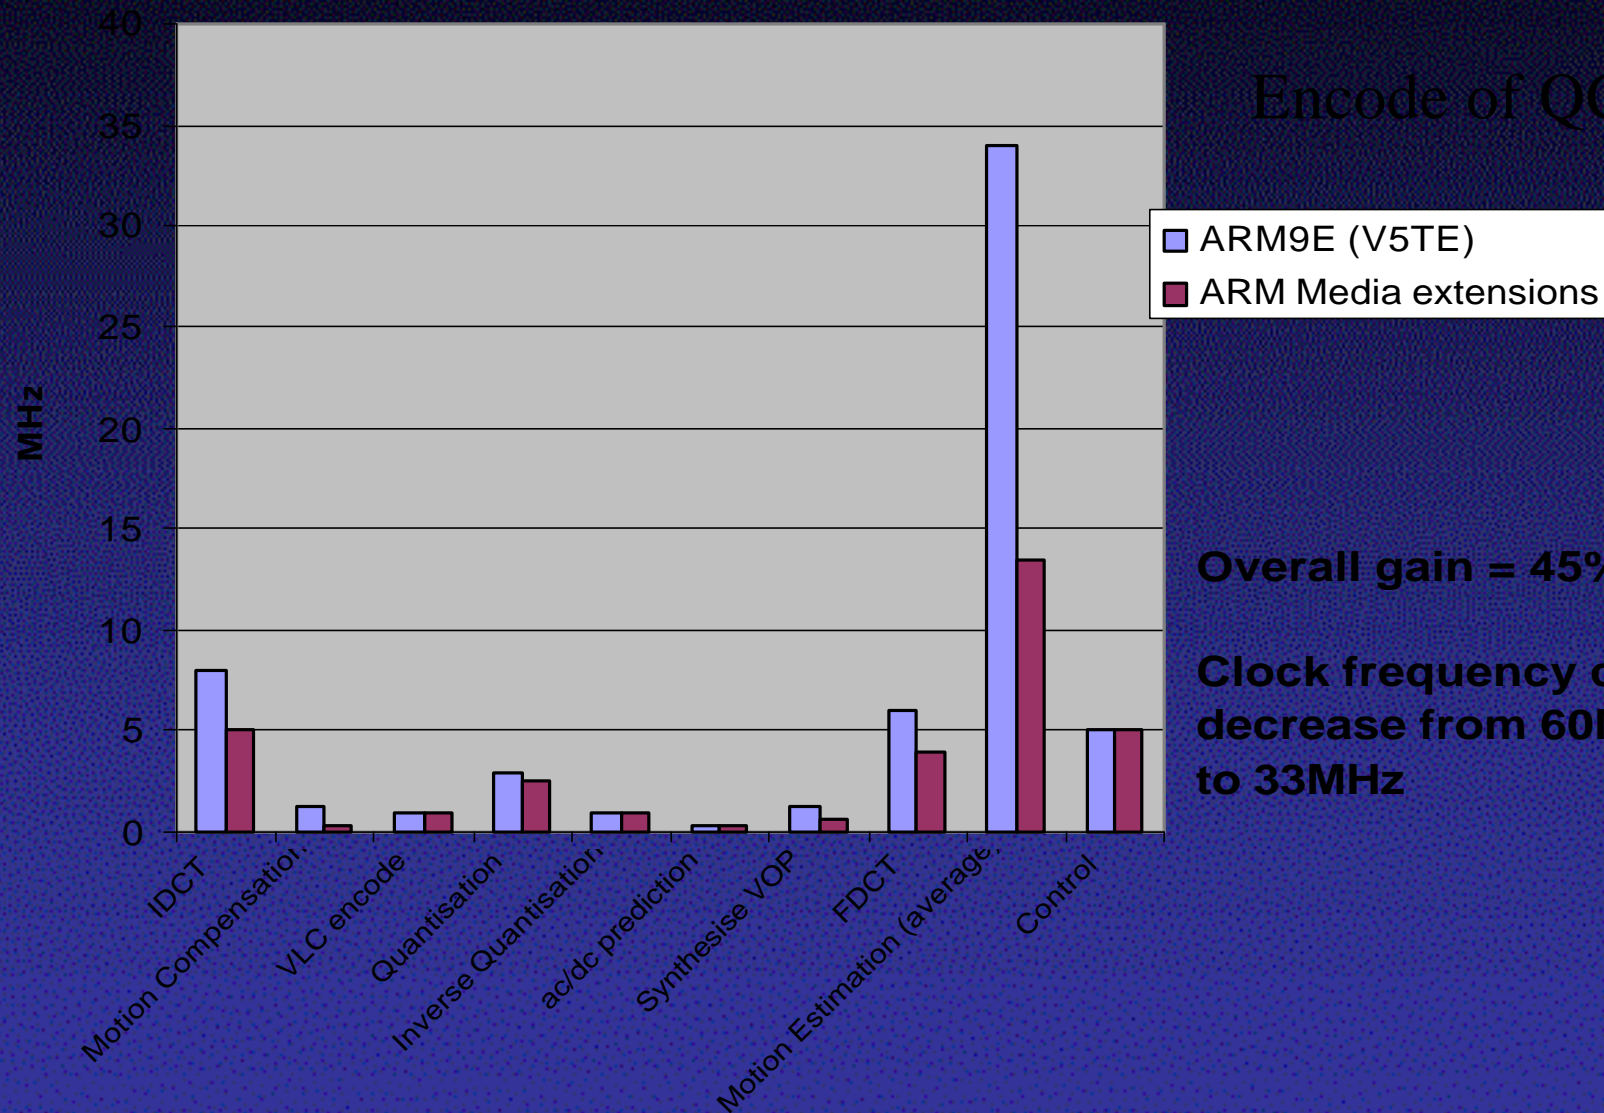
ADD8	Rp1, Rs1, Rb1	; path 1 metric = state 1 metric + metric for branch 1
ADD8	Rp2, Rs2, Rb2	; path 2 metric = state 2 metric + metric for branch 2
USUB8	Rt, Rp1, Rp2	; compare metrics - setting the SIMD flags
SEL8	Rd, Rp2, Rp1	; choose best (smallest) metric

- Total of 4-cycles for a single-cycle implementation

Total of 12-cycles for previous ARM ISA v5TE implementation

MPEG4 performance on ARM

Encode of QCIF at 15fps



Overall gain = 45%

Clock frequency can decrease from 60MHz to 33MHz

Dominated by motion estimation

- Underling operation is a sum of absolute differences

;Implementation using the SIMD operations

LDR	Rx, [pRx], #Offset	;get x(i,j)
LDR	Ry, [pRy], #Offset	;get y(i+m, j+n)
USUB8	Rt1, Rx, Ry	;Rt1 = x(i,j) - y(i+m,j+n)
USUB8	Rt2, Ry, Rx	;Rt2 = y(i+m,j+n) - x(i,j)
SEL8	Rd, Rt2, Rt1	;ABS(x(i,j) - y(i+m,j+n))

- Taking 5-cycles in a single-cycle implementation
 - Previously on v5TE cores, this takes around 15-cycles

Conclusions

- Up to 4x performance at virtually zero power increment
- Single core solutions for high performance platform processing requirements such as wireless and portable applications
- Fully backward compatible with the ARM roadmap
 - Value of existing 3rd party software support is maintained
 - No OS changes required for applications to use the new extensions