



Using Intrinsics to Exploit the Multimedia 8-way SIMD Instructions in SH-5

Speaker: Andy Sturges, Director of IP, SuperH, Inc.

**andy.sturges@superh.com
www.superh.com**

SuperH, Inc. Introduction



- ◆ Licenses RISC CPU cores and development tools
 - ◆ Extensive compiler developments within GCC

SIMD for multimedia data processing
Single Instruction Multiple Data

SH-4 family
32-bit RISC
Vector FPU

SH-5 family
64-bit RISC
VFPU + SIMD

SH-6 family
Superscalar
Announce 2004


SH-7 family
Multi-threading
Announce 2006

SuperH Multimedia
RISC CPU family



Building Applications with SIMD

- ◆ **SIMD instructions:**
 - ◆ Not well supported by standard industry compilers
 - ◆ Not automatically generated
 - ◆ Assembly language inserts
- ◆ **No type checking**
- ◆ **No portability**
- ◆ **Overhead caused by compiler**




**SIMD support required for
mainstream application use**



'Ideal' SIMD Compiler

- ◆ **Assembler level performance with no overhead**
- ◆ **Full language support for targeting SIMD**
- ◆ **Rapid optimization of existing 'C' code**
 - ◆ **Minimal changes to original code**
 - ◆ **Full benefit of compiler optimizations**
- ◆ **As portable as standard 'C' code**



GCC Vector extension solution

- ◆ **Ultimately: automatic vector generation from C**



What are Intrinsics?

Functions built-in to the compiler to describe instructions

- ◆ **SIMD instruction (8-way unsigned 8-bit add with saturation)**

`madds.ub source1, source2, result`

- ◆ **Intrinsic definition**

`unsigned int64 madds_ub (unsigned int64 source1, unsigned int64 source2)`

- ◆ **Knowledge of intrinsic via GCC instruction pattern**

`(define_insn "usaddv8qi3"[...match_operands...] " madds_ub %1,%2,%0"..)`

- ◆ **Vector definition**


`vector unsigned char madds_ub (vector unsigned char, vector unsigned char)`



GCC Vector Extensions

- ◆ Introduced in 1998 to GCC technology
 - ◆ Extends C language to describe vector objects
 - ◆ Initially supported ALTIVEC (PowerPC) SIMD architecture
 - ◆ Other architectures now supported
 - ◆ e.g. TI C6X, SuperH and MMX / x86

Vector C ideal for SIMD



- ◆ But still no automatic vectorizations!



GCC Vector Extensions Deliver

◆ Performance

- ◆ SIMD instructions fully described to the compiler as intrinsics
- ◆ Results in no run-time overhead in use of intrinsics

◆ Language

- ◆ Interface to SIMD instructions forces use of correct vector types
- ◆ Vector types extended through all of the C language
- ◆ Provides natural interface to SIMD instructions

◆ Portability

- ◆ SIMD programs can be ported across architectures with compiler support (e.g. detection of different vector widths)
- ◆ First portable language for SIMD programming
- ◆ SuperH investigating MMX to SH-5 translation
- ◆ Can define standard SIMD API



GCC Vector SH-5 SIMD Intrinsics

- ◆ **SH-5 fully implements the GCC vector extensions**
- ◆ **Intrinsic provided for each SIMD operation**
- ◆ **ABI defined functional interface to SIMD instructions**
- ◆ **Forces use of correctly typed operands to SIMD operations**
 - ◆ Important where SIMD instructions perform subtle type conversions
- ◆ **Fully type checked and optimized by compiler**
 - ◆ Compiler performs unrolling, in-lining and scheduling
 - ◆ Compiler takes care of register allocations and load/store
- ◆ **Returned to open source**
 - ◆ Other SIMD architectures can exploit technology



Detailed Example: MPEG4 Encode

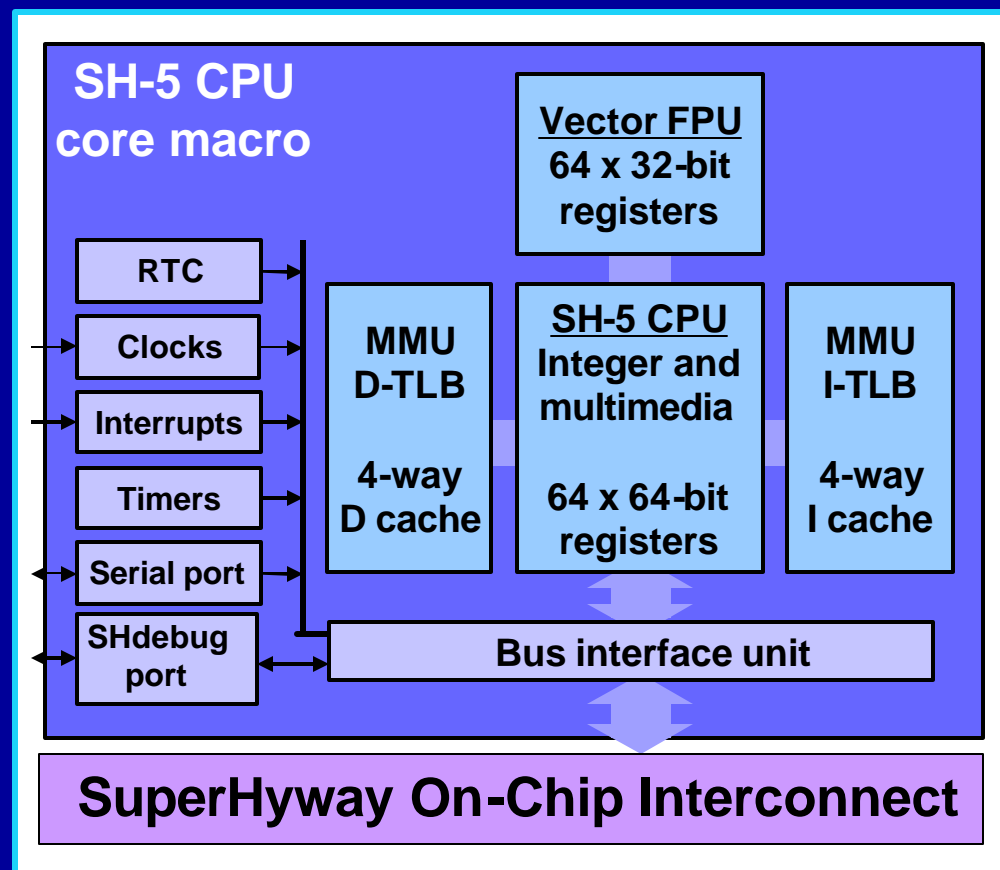
◆ Objective

- ◆ Using public source code demonstrate the ability to run MPEG4 encode in software on the SH-5 CPU core
- ◆ SuperH have used the 'Xvid' publicly available codec (see <http://www.xvid.org>)
- ◆ SuperH optimized the encoder using SHmedia SIMD intrinsics for the SH-5 64-bit CPU



SH-5: 64-bit Multimedia CPU

- ◆ **Integrated RISC and DSP**
 - ◆ 1.5 DMIPS / MHz
 - ◆ 7 MFLOPS / MHz
- ◆ **Exploits 64-bit registers for**
 - ◆ 2,4,8-way SIMD delivers
 - ◆ 24 MOPS / MHz
- ◆ **CPU core 3mm² in 0.13u**
 - ◆ Complete hard macro 12mm²
 - ◆ 400MHz





1st Embedded 64-bit SIMD CPU

- ◆ **SH-5 support for video and DSP data types:**
 - ◆ 8x8 / 4x16 / 2x32-bit, signed/unsigned/fractional, saturate/modulo
 - ◆ 64 high accuracy accumulators (64-bit general purpose registers)

- ◆ **SH-5 set of SIMD operators:**
 - ◆ Multiplies (four 16x16=32 per cycle): Mul, MulAcc, MulAdd, MulSub
 - ◆ Arithmetic: Add, Sub, Abs, Compare, Convert
 - ◆ Rearrangement: Shift, Shuffle, Permute, Bitwise Cmove, Extract
 - ◆ Motion estimation: Sum of Absolute Differences 9.6GOPS at 400MHz

- ◆ **Ideally suited to exploiting GCC vector intrinsics**



MPEG4: Motion Estimation C Code

Main inner loop of XVID MPEG4 motion estimation:

```
uint32_t sad8_c(const uint8_t * const cur,  const uint8_t
    * const ref, const uint32_t stride)
{
    uint32_t sad = 0, i, j;
    uint8_t const *ptr_cur = cur; *ptr_ref = ref;

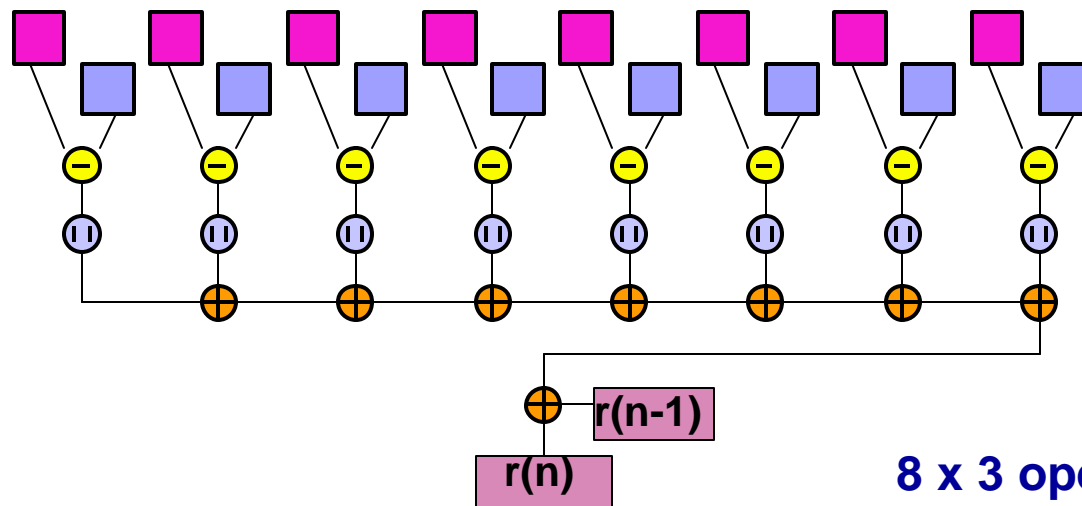
    for (j = 0; j < 8; j++) {
        for (i = 0; i < 8; i++)
            sad += ABS(*(ptr_cur + i) - *(ptr_ref + i));
        ptr_cur += stride; ptr_ref += stride;
    }

    return sad;
}
```



Motion Estimation Optimized

```
for (j = 0; j < 8; j++) {  
    vector char cur = *(vector char*)ptr_ref;  
    vector char ref =  
        (vector char)sh_media_unaligned_LD_Q((uint8_t*)ptr_ref,0);  
  
    sad = sh_media_MSAD_UBQ(cur, ref, sad);  
  
    ptr_cur +=stride; ptr_ref +=stride;  
}
```



Optimized using
SH-5 MSAD instruction

MSAD.UBQ Rm, Rn, Rw
Sum of absolute differences
of unsigned 8-bit

8 x 3 operation/cycle x 400MHz = 9.6 GOPS



Comparison of Compiler Output

◆ Original

- ◆ Generates 122 instructions
- ◆ Takes 976 cycles

```
...  
ld.ub    r41,0,r38  
ld.ub    r40,0,r39  
add.l    r19,r21,r19  
add.l    r17,r21,r17  
sub.l    r38,r39,r37  
add.l    r37,r63,r7  
sub      r63,r7,r8  
cmpgt    r63,r7,r9  
cmvne    r9,r8,r7  
...
```

} X 64

◆ Optimized

- ◆ Generates 52 instructions
- ◆ Takes 59 cycles



```
...  
or        r60,r61,r21  
ldhi.q    r43,0,r6  
ldlo.q    r20,0,r7  
addi.l    r8,7,r9  
ld.q      r20,0,r23  
msad.ubq  r22,r21,r36  
...
```

} X 8




MPEG-4 Encode Result

Unoptimized



SIMD Optimized
>80% improvement



SH-5 Performance	QCIF 64kbps, 15fps	CIF 384bps, 25fps	VGA 1mbps, 25fps
Unoptimized	97.9MHz	771.1MHz	n/a
Optimized	18.3MHz	146MHz	398MHz

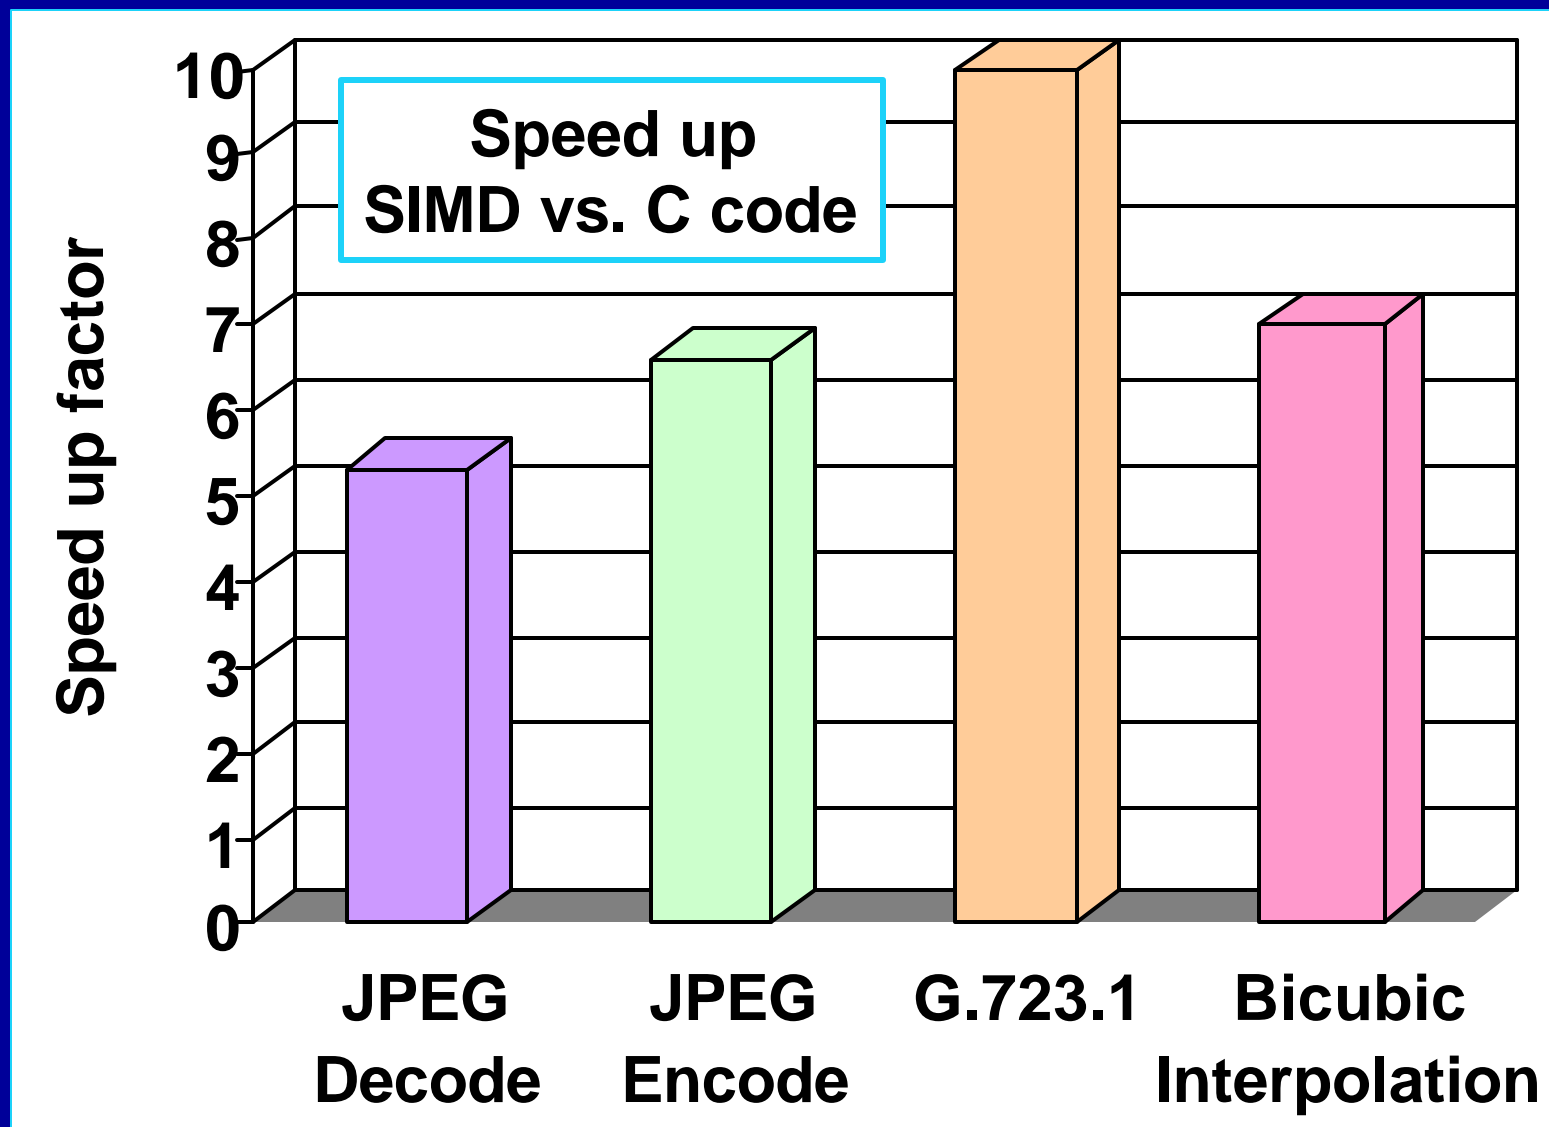


Benchmarking

- ◆ **CPU performance is combination of**
 - ◆ **Architecture and compiler**
- ◆ **EEMBC provides a mechanism for evaluating SIMD**
 - ◆ **‘Out of box’ and ‘full fury’**
- ◆ **SH-5 ‘full fury’ (optimized) benchmark will be published later this year using GCC vector intrinsics**



Examples of 64-bit SIMD Intrinsic



JPEG using EEMBC® benchmarks. To be certified in 2H 2003.



Optimizing SH-5 Cores

- ◆ **Future SH-5 targeting handheld multimedia devices**
 - ◆ Dynamic power: clock gating and low power datapath elements
 - ◆ Power management: dynamic frequency and voltage control
 - ◆ Static power: mixed Vth technology and 'unit' shutdown
 - ◆ Target: 0.30mW/MHz at 300MHz
- ◆ **Delivers very low power consumption for mobile multimedia**

Xvid MPEG-4 encode power requirements		
Profile	QCIF, 64kbs, 15fps	CIF, 384kbs, 25fps
Power	5.5mW	44mW



SH-6 Plans

◆ SH-6 will deliver even higher performance:

- ◆ With GCC Vector extensions SIMD code will be portable to SH-6
 - ◆ Used existing applications to optimize microarchitecture
- ◆ Super-scalar optimized for the embedded market
 - ◆ Dual issue 8-way SIMD
 - ◆ Deeper pipeline for higher clock frequency
- ◆ New SIMD instructions for specific applications, e.g.
 - ◆ H.264
 - ◆ JPEG 2000
 - ◆ Cryptography
 - ◆ Networking



Summary

- ◆ **GCC vector extensions enable effective programming model for SIMD**
 - ◆ This has been demonstrated on a variety of applications
 - ◆ Proven in EEMBC benchmarks
- ◆ **SH-5 8-way SIMD delivers high architecture performance**
 - ◆ Replaces dual CPU and DSP solutions
- ◆ **SH-6 can fully exploit all existing SH-5 optimized code:**
 - ◆ Full binary compatibility
 - ◆ Re-compile C code for maximum performance