

# SPARC64™ VIIIfx: Fujitsu's New Generation Octo Core Processor for PETA Scale computing



August 25, 2009

Takumi Maruyama

LSI Development Division  
Next Generation Technical Computing Unit  
Fujitsu Limited

# Fujitsu Processor Development

Features

HPC-ACE  
System On Chip

Hardware Barrier

Multi-core / Multi-thread

L2\$ on Die

Non-Blocking \$  
O-O-O Execution  
Super-Scalar

Single-chip CPU

Store Ahead  
Branch History  
Prefetch

\$ ECC  
Register/ALU Parity  
Instruction Retry  
\$ Dynamic Degradation  
RC/RT/History

High Performance  
Technology

High Reliability  
Technology

Tr=760M  
CMOS Cu / 45nm

SPARC64  
VIIIfx

HPC

UNIX

SPARC64™  
Processor

Mainframe

GS21

Tr=600M  
CMOS Cu  
65nm

SPARC64  
VII

Tr=540M  
CMOS Cu  
90nm

SPARC64  
VI

Tr=400M  
CMOS Cu  
90nm

SPARC64  
V+

Tr=190M  
CMOS Cu  
130nm

SPARC64  
V

Tr=30M  
CMOS Cu  
180nm / 150nm

SPARC64  
GP

GS21  
900

Tr=500M  
CMOS Cu  
90nm

Tr=190M  
CMOS Cu  
130nm

GS21  
600

Tr=46M  
CMOS Cu  
180nm

GS8900

GS8800B

Tr=30M  
CMOS Al  
250nm / 220nm

GS8800

Tr=10M  
CMOS Al  
350nm

GS8600

SPARC64  
II

SPARC64

~1995

1996  
~1997

1998  
~1999

2000  
~2003

2004  
~2007

2008~

SPARC64™ VIIIfx

# SPARC64™ VIIIfx Design Target

- ◆ Processor for Fujitsu's supercomputer for the peta-scale computing age, which realizes both high performance and low power
  - Unachievable goal with conventional processor design
    - More GF (Giga Flops)
    - More Efficiency
    - No more high frequency
  - ➔ Large ISA (Instruction Set Architecture) extension is required
- High Integration: SoC (System On Chip)
- High Reliability
- ◆ Reuse SPARC64™ VII design when applicable

Focus of this presentation

# HPC-ACE


(High Performance Computing - Arithmetic Computational Extensions)

- ◆ ISA (Instruction Set Architecture) of SPARC64Vllfx is:
  - Complies with
    - The SPARC-V9 standard
    - JPS (Joint Programmer's Specification): Extension to SPARC-V9
  - HPC-ACE: Fujitsu's unique ISA extension for HPC
    - Large register sets
    - SIMD (single instruction multiple data) instructions
    - etc
- ◆ You can download SPARC64™ Vllfx ISA documents from <http://jp.fujitsu.com/solutions/hpc/brochures/>
  - The SPARC® Architecture Manual Version 9
  - SPARC® Joint Programming Specification (JPS1): Commonality
  - SPARC64 Vllfx Extensions

# Large register sets 1/2

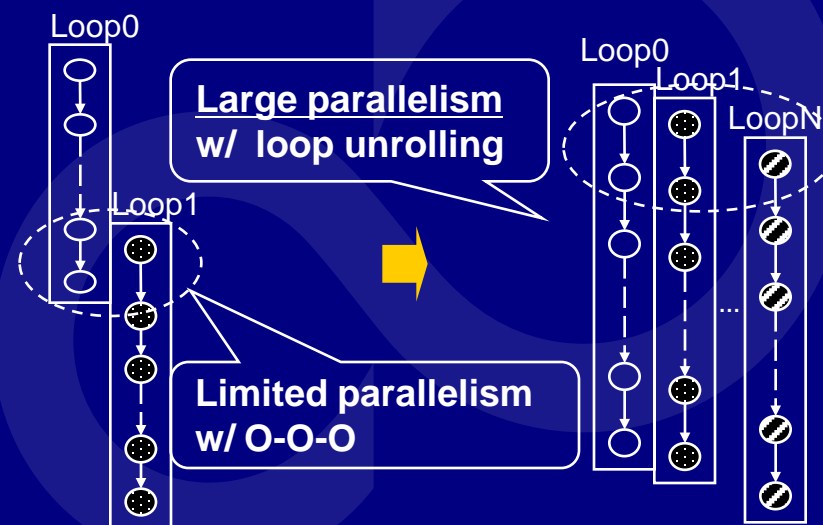
## ◆ Register enhancements from V9

- 32→64 INT registers
- 32→256 FP registers (DP)
  - Lower 32 registers are the same with SPARC-V9's.
  - FP registers are “flat”. Extended FP registers can be accessed with non-SIMD instruction as well.



## ◆ Why needed

- To extract more parallelism currently limited by the number of Arch registers.
- Reduce spill/fill overhead



# Large register sets 2/2


## ◆ Instruction format for 256 FP registers

- 8 bit x 4 (3 read +1 write) register number fields are necessary for FMA (Floating-point Multiply and Add) instruction.
- But SPARC-V9 instruction length is limited (32bits – fixed)


➔ Defined a new prefix instruction (SXAR) to specify upper-3bit of register numbers of the following two instructions.

## ◆ SXAR (Set XAR) instruction

- XAR: Extended Arithmetic Register
  - Set by the SXAR instruction
  - Valid bit is cleared once the corresponding subsequent instruction gets executed.



- SXAR1: set XAR for subsequent one instruction.
- SXAR2: set XAR for subsequent two instructions.



# SIMD

## ◆ SIMD FP operation

- 1 SIMD FP instruction executes two SP or DP floating-point operations


## ◆ SIMD Load/Store

- 1 SIMD load/store instruction accesses two contiguous SP or DP data in memory
- Data alignment requirement for SIMD-load (DP) is 8 byte rather than usual 16byte.
- Combine two independent FP operation into one with the Flat FP registers

## ◆ SXAR instruction

- Specifies SIMD operation of the subsequent two instructions.
- No new opcode is required for SIMD operation

➔ More SW optimization possibility with SIMD



### Conventional SIMD usage

SIMD-Load (Src1)  
SIMD-Load (Src2)  
SIMD-FP  
SIMD-Store (Dst)

### Combine two FP ops with SIMD

Load (Src1-A)  
Load (Src1-B)  
Load (Src2-A)  
Load (Src2-B)  
SIMD-FP  
Store (Dst-A)  
Store (Dst-B)

# FP Trigonometric Functions

## ◆ Instructions for fast Trigonometric Function calculation

### ◆ Taylor series approximation of sin (x)

$$\sim x - 1/3!x^3 + 1/5!x^5 - 1/7!x^7 + 1/9!x^9 - 1/11!x^{11} + 1/13!x^{13} - 1/15!x^{15}$$

$$= x (((((((0 - 1/15!)x^2 + 1/13!)x^2 - 1/11!)x^2 + 1/9!)x^2 - 1/7!)x^2 + 1/5!)x^2 - 1/3!)x^2 + 1)$$

### ◆ New instruction - ftrimaddd

■ Function:  $rs1 \times \text{abs}(rs2) + T[\text{index}] \rightarrow rd$

■ Usage:

# S = 0

ftrimaddd S, X<sup>2</sup>, 7, S # S \* X<sup>2</sup> - 1/15! → S

ftrimaddd S, X<sup>2</sup>, 6, S # S \* X<sup>2</sup> + 1/13! → S

ftrimaddd S, X<sup>2</sup>, 5, S # S \* X<sup>2</sup> - 1/11! → S

ftrimaddd S, X<sup>2</sup>, 4, S # S \* X<sup>2</sup> + 1/9! → S

ftrimaddd S, X<sup>2</sup>, 3, S # S \* X<sup>2</sup> - 1/7! → S

ftrimaddd S, X<sup>2</sup>, 2, S # S \* X<sup>2</sup> + 1/5! → S

ftrimaddd S, X<sup>2</sup>, 1, S # S \* X<sup>2</sup> - 1/3! → S

Fmuld S, X, S # S \* X → S




# Software controlled Cache

- ◆ Give SW to ability to control cache to optimize performance while keeping cache coherency
- ◆ Divide Cache into 2 groups (sectors)
  - SXAR instruction specifies the sector
    - sector 0: Instruction fetch / normal operand access (default)
    - sector 1: operand access explicitly specified by SXAR
  - Sector cache configuration register
    - Specifies the ratio of sector 0 and 1 at the same index

## ◆ HW Implementation

- Keep sector info of each cache line.
  - Select the way to be replaced to meet the sector ratio on cache misses
- SW specifies a sector depending on temporal and spatial locality of data




# Other HPC-ACE features

## ◆ Conditional operation

- For efficient execution of Loop with 'if'
- Conditional branch can be removed with the following conditional instructions.
  - FP Conditional Compare to Register
  - Move Selected FP-register on FP-register's condition
  - Store FP-register on FP-register's condition

→ Compiler can optimize the loop with SW pipeline

```
for (i=0; i<00; i++) {  
  if (A(i)>0)  
    FPop X(i)  
  else  
    FPop Y(i)  
}
```



## ◆ FP Reciprocal Approximation of Divide/Square-root

- Calculate reciprocal approximation with rounding error  $< 1/256$
- To achieve higher divide/square-root performance with pipelined operation

## ◆ FP Minimum and Maximum

Usage:

# %f2 / %f10 → %f0

**frcpad** %f10, %f6

fmuld %f2, %f6, %f2

fnmsubd %f6, %f10, 1.0, %f6

fmuld %f6, %f6, %f0

fmadd %f6, %f6, %f6, %f4

fmadd %f0, %f0, %f6, %f0

fmadd %f4, %f2, %f2, %f4

fmadd %f0, %f4, %f2, %f0

# Integrated Multicore Parallel Architecture

◆ SPARC64™ VIIIfx HW

- HPC-ACE
- Shared L2 cache to avoid false sharing
- Hardware barrier for fast inter-core synchronization


◆ Fujitsu's compiler technology

- Automatic parallelization

➔ More possibility of optimization:  
Parallelize the innermost loop


● Vector

```
DO J=1,N
V DO I=1,M
V A(I,J)=A(I,J+1)*B(I,J)
V END
END
```




● Conventional Scalar

```
P DO J=1,N
DO I=1,M
A(I,J)=A(I,J)*B(I,J)
END
P END
```




● SPARC64™ VIIIfx

```
DO J=1,N
P DO I=1,M
P A(I,J)=A(I,J+1)*B(I,J)
P END
END
```

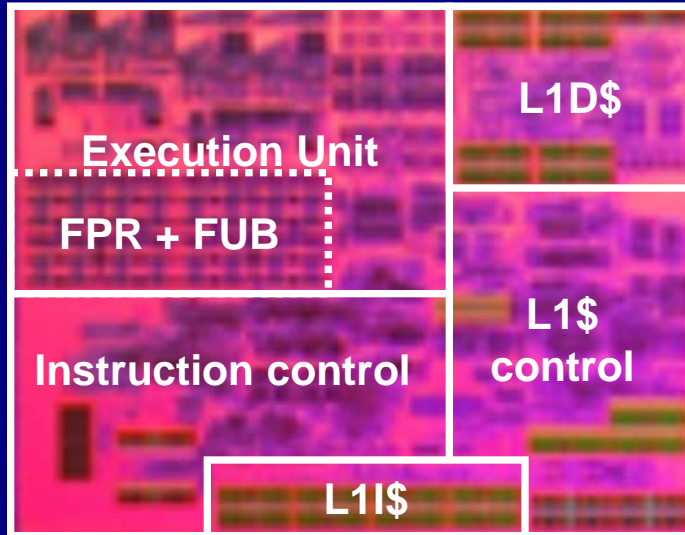


# SPARC64™ VIIIfx Chip Overview



- **Architecture Features**
  - 8 cores
  - Shared 5 MB L2\$
  - Embedded Memory Controller
  - 2 GHz
- **Fujitsu 45nm CMOS**
  - 22.7mm x 22.6mm
  - 760M transistors
  - 1271 signal pins
- **Performance (peak)**
  - 128GFlops
  - 64GB/s memory throughput
- **Power**
  - 58W (TYP, 30°C)
  - Water Cooling – Low leakage power and High reliability

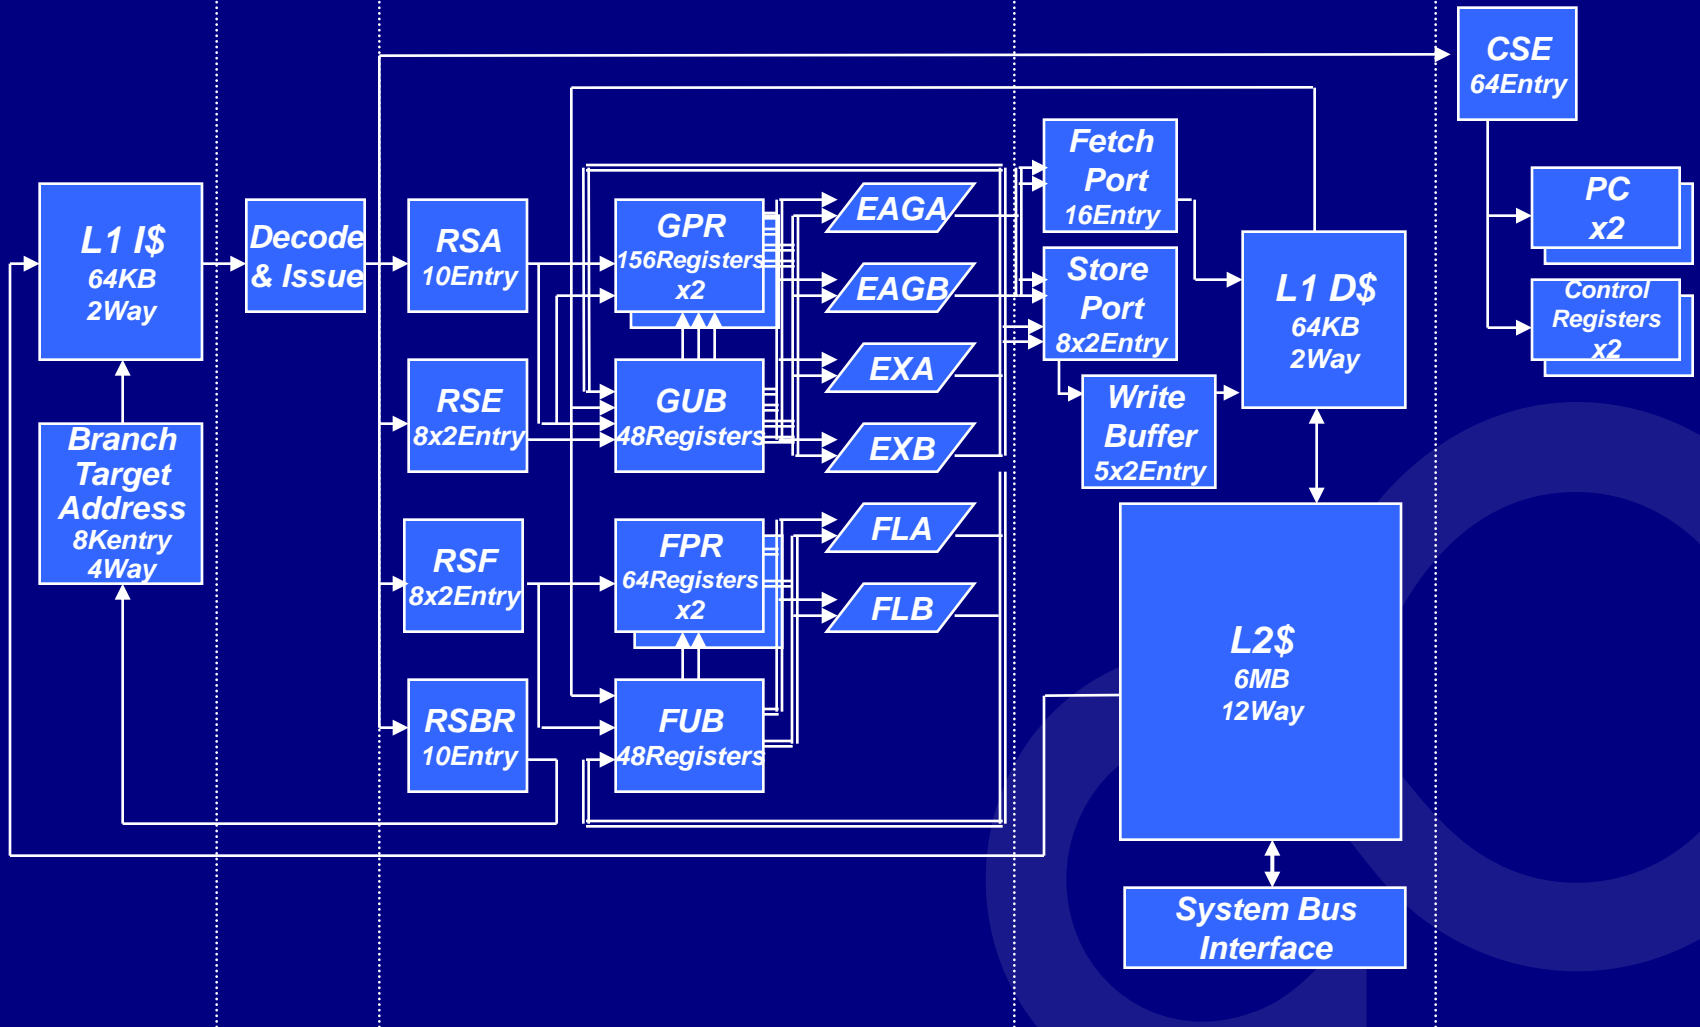
# SPARC64™ VIIIfx Core spec



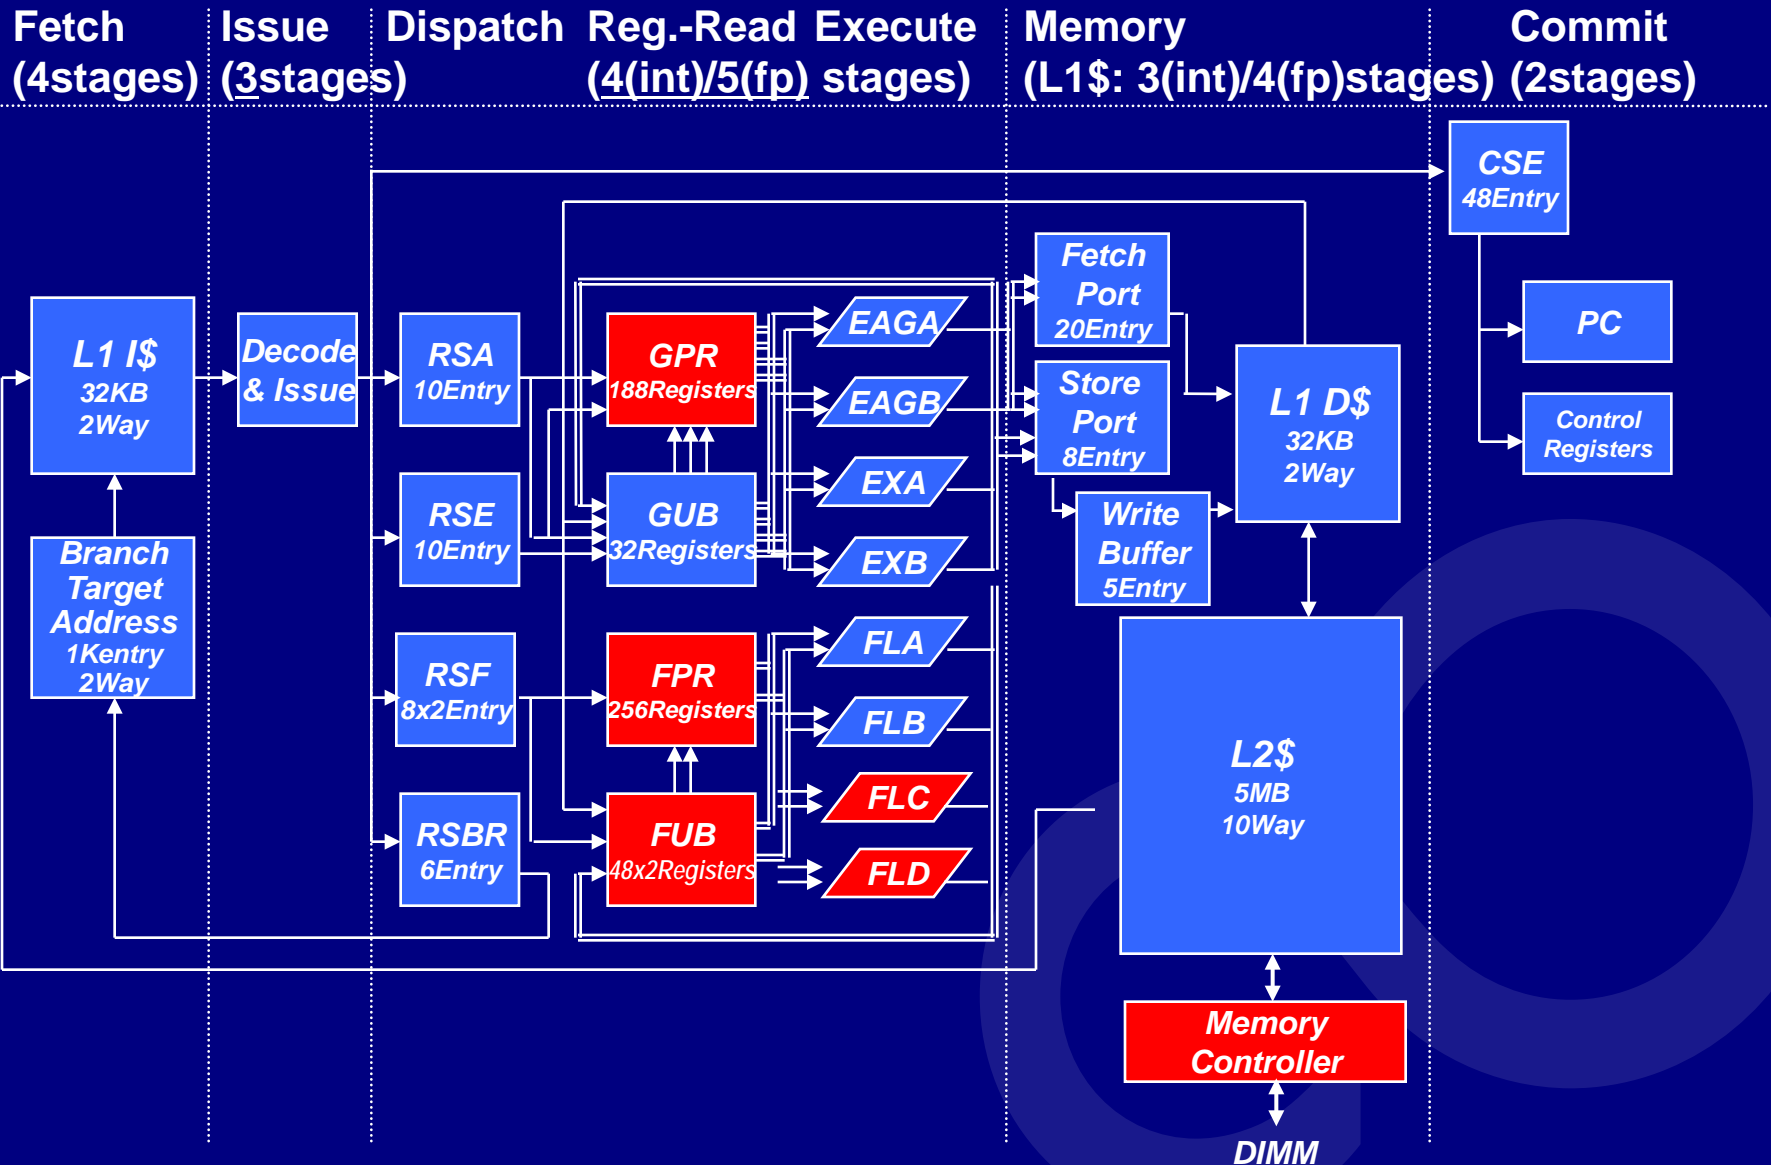
|                                     |   |   |
|-------------------------------------|---|---|
| <b>Instruction Set Architecture</b> | SPARC-V9/JPS + HPC-ACE                                |   |
| <b>#FP-operations per clock</b>     | 8 (= 4 Floating Multiply and Add)                     |   |
| <b>Execution units</b>              | <integer>   | <floating-point>  |
|                                     | ALU x2<br>SHIFT x2<br>MULT x1<br>DIVIDE x1<br>AGEN x2 | FMA x4 (2SIMD)<br>COMPARE x2<br><br>DIVIDE x2<br>VIS x1 |
| <b>#Registers</b>                   | 188 (GPR)<br>32 (GUB)                                 | 256 (FPR)<br>48 x2 (FUB)                                |
| <b>L1\$</b>                         | L1I\$ 32KB/2way<br>L1D\$ 32KB/2way                    |   |

# SPARC64™ VII Pipeline

Fetch (4stages)    Issue (2stages)    Dispatch (4stages)    Reg.-Read (4stages)    Execute (4stages)    Memory (L1\$: 3(int)/4(fp)stages)    Commit (2stages)



# SPARC64™ VIIIfx Pipeline







# Changes in the pipeline

## ◆ Issue stage

- Added “PD” stage to handle SXAR
- 6-instructions packed into 4 at “PD”
- An CSE entry is assigned at “D”
- The packed instructions have
  - Extended register fields
  - SIMD operation attributes
- 4 packed (=6 unpacked) instructions can be committed at the same time.

## ■ Execution Stage

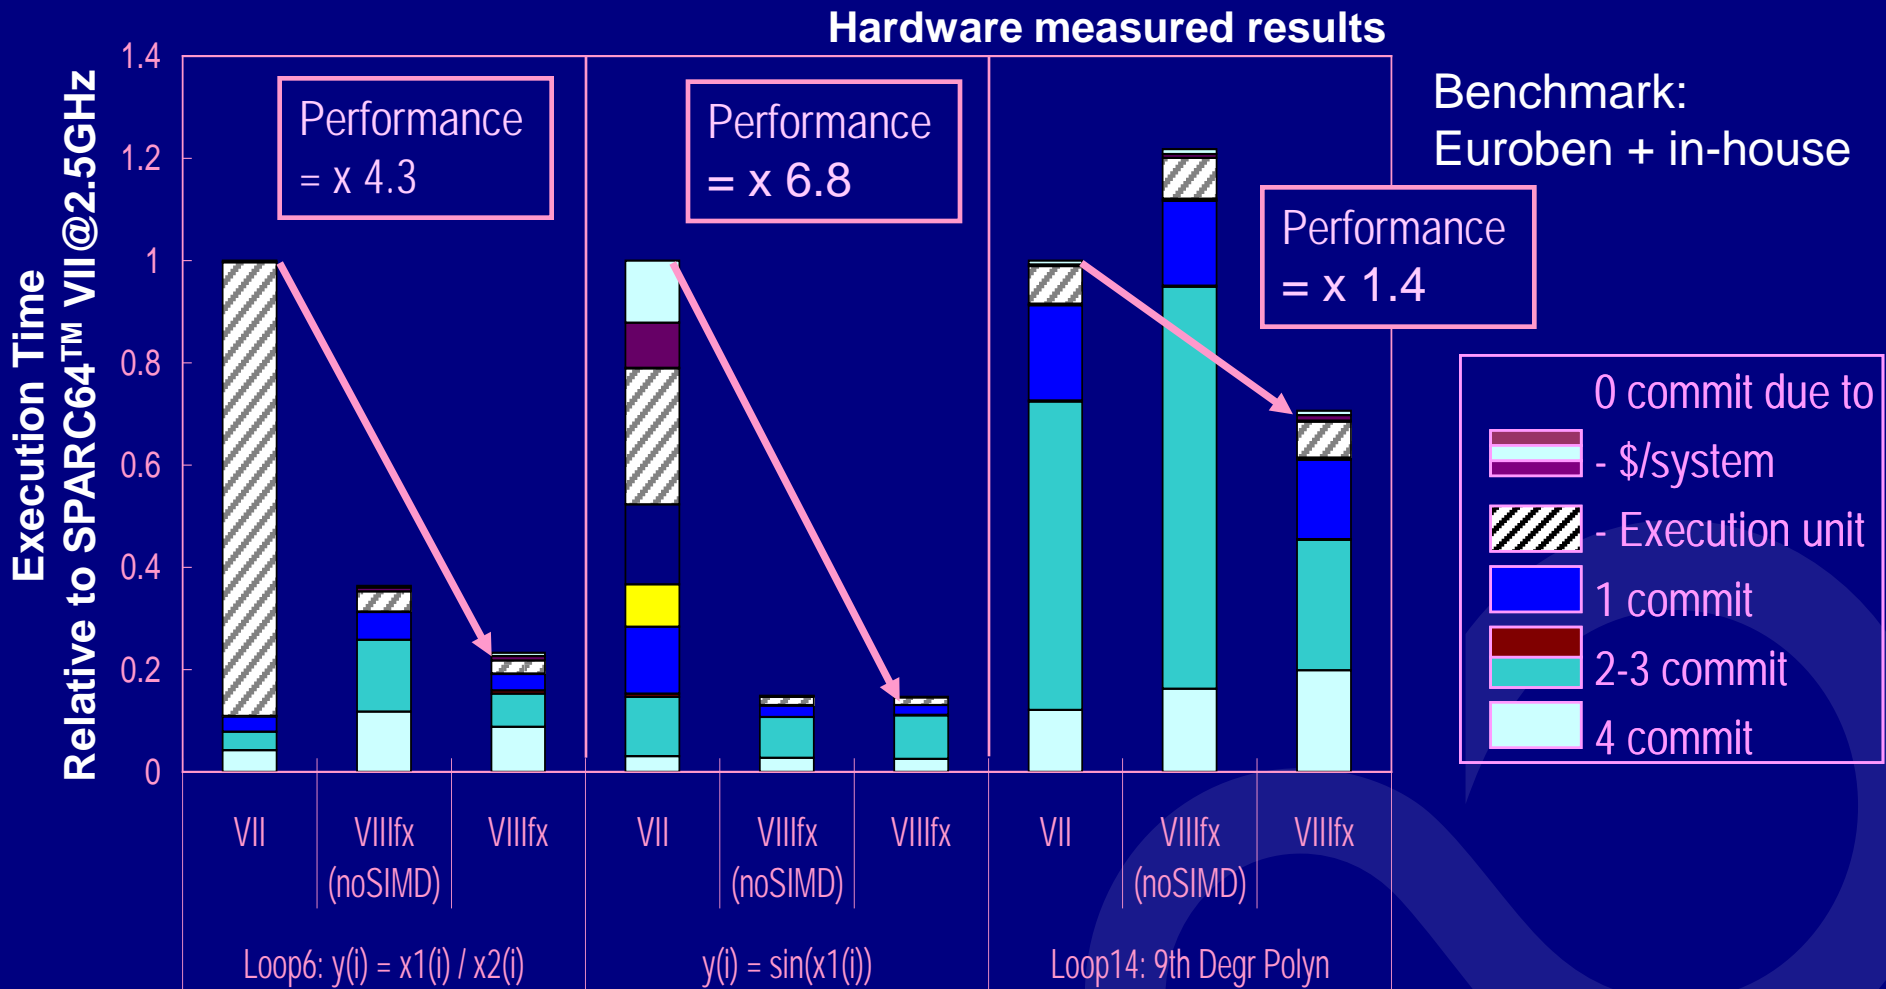
- “B” stage has been split into two for flat Floating-point Register access





# Performance Results – Core


## SPARC64™ VIIIfx@2.0GHz



➔ SPARC64™ VIIIfx realizes much higher core performance than SPARC64™ VII thanks to HPC-ACE despite its lower frequency

# Performance Results – Chip

## SPARC64™ VIIIfx@2.0GHz




- ➔ SPARC64™ VIIIfx shows about x 2.5 performance of SPARC64™ VII
  - Stall time due to the execution unit busy has been reduced dramatically.
- ➔ Expect x 3 performance with further compiler optimization

# SPARC64™ History and Future

Evolution rather than revolution

Peak Performance & Power



## ■ SPARC64™ V (1 core)

- RAS
- Single thread performance

## ■ SPARC64™ VI (2 core x 2 VMT)

- Throughput

## ■ SPARC64™ VII (4 core x 2 SMT)

- More Throughput
- High Performance Computing



## ■ SPARC64™ VIIIfx (8 core)

- High Performance Computing
- Low Power
- SoC

# SPARC64™ VIIIfx Summary

- ◆ SPARC64™ VIIIfx has been designed to be used for Fujitsu's supercomputer for the PETA-scale computing age.
- ◆ HPC-ACE instruction sets overcomes the limitation of conventional SPARC-V9 architecture.
- ◆ SPARC64™ VIIIfx has combined high performance and low power.
- ◆ SPARC64™ VIIIfx chip is up and running in the lab.
- ◆ Fujitsu will continue to develop SPARC64™ series to meet the needs of a new era.

# Abbreviations

- SPARC64™ VIIIfx
  - IB: Instruction Buffer
  - RSA: Reservation Station for Address generation
  - RSE: Reservation Station for Execution
  - RSF: Reservation Station for Floating-point
  - RSBR: Reservation Station for Branch
  - GUB: General Update Buffer
  - FUB: Floating point Update Buffer
  - GPR: General Purpose Register
  - FPR: Floating Point Register
  - CSE: Commit Stack Entry