

Vector Extensions to the MIPS-IV Instruction Set Architecture
(The V-IRAM Architecture Manual)
Revision 3.7.5

David Martin

March 4, 2000

Contents

1	Vector Architectural State	11
1.1	Overview	11
1.2	Vector Control Registers	11
1.3	Vector Scalar Registers	11
1.4	Vector Flag Registers	12
1.5	Number of Virtual Processors	12
1.6	Virtual Processor Width	12
2	Instruction Set	19
2.1	Open Questions	19
2.2	Data Types	19
2.3	Vector Operations	19
2.4	Flag Register Use	20
2.5	Memory Consistency	20
2.6	Exception Model	20
2.7	Speculative Execution	20
2.8	Assembly Language	20
2.9	Instructions	21
2.9.1	Vector Integer Arithmetic Instructions	21
2.9.2	Vector Logical Instructions	21
2.9.3	Vector Floating-Point Arithmetic Instructions	23
2.9.4	Vector Floating-Point Convert Instructions	23

2.9.5	Vector Fixed-Point Arithmetic Instructions	23
2.9.6	Vector Logical Flag Instructions	24
2.9.7	Flag Processing Instructions	24
2.9.8	Vector Processing Instructions	25
2.9.9	Vector Memory Instructions	25
2.9.10	Coprocessor Interface Instructions	26
2.9.11	Miscellaneous Vector Instructions	26
2.9.12	Miscellaneous Kernel Vector Instructions	26
A	Instruction Formats	45
B	Instruction Definitions	47
	CFC2 — Control From Coprocessor 2	48
	CTC2 — Control To Coprocessor 2	49
	MTC2 — Move To Coprocessor 2	50
	VABS — Vector Integer Absolute Value	51
	VABS.fmt — Vector Floating-Point Absolute Value	53
	VACOMMIT — Commit Speculative Arithmetic	56
	VADD — Signed Vector Integer Add	57
	VADD.U — Unsigned Vector Integer Add	59
	VADD.fmt — Vector Floating-Point Add	60
	VAND — Vector And	64
	VCEIL — Vector Floating-Point Ceiling	65
	VCIOTA — Vector Continuous Iota	70
	VCMP — Signed Vector Integer Compare	71
	VCMP.U — Unsigned Vector Integer Compare	73
	VCMP.fmt — Vector Floating-Point Compare	75
	VCOMPRESS — Vector Compress	79
	VCVT — Vector Convert	80
	VDIV — Signed Vector Integer Divide	91

VDIVU — Unsigned Vector Integer Divide	93
VDIV.fmt — Vector Floating-Point Divide	95
VEXPAND — Vector Expand	99
VEXT.SV — Signed Scalar-Vector Extract	100
VEXT.U.SV — Unsigned Scalar-Vector Extract	101
VEXT.VV — Vector-Vector Extract	102
VEXTHALF — Vector Extract Half	103
VHALF — Vector Half	104
VHALFDN — Vector Half Down	106
VHALFUP — Vector Half Up	108
VFAND — Vector Flag And	110
VFCLR — Vector Flag Clear	111
VFCLR8 — Vector Flag Clear 8	112
VFFF1 — Vector Flag Find First One	113
VFFL1 — Vector Flag Find Last One	114
VFINS — Scalar-Vector Insert	115
VFLD — Vector Flag Load	116
VFLOOR — Vector Floating-Point Floor	119
VFMF8 — Vector Flag Move From 8	125
VFMT8 — Vector Flag Move To 8	126
VFNOR — Vector Flag Nor	127
VFOR — Vector Flag Or	128
VFOR8 — Vector Flag Or 8	129
VFPOP — Vector Flag Population Count	130
VFLUSH — Flush Vector Unit	131
VFSET — Vector Flag Set	132
VFSETBF — Vector Flag Set Before First One	133
VFSETIF — Vector Flag Set Including First One	134
VFSETOF — Vector Flag Set Only First One	135

VFST — Vector Flag Store	136
VFXOR — Vector Flag Xor	138
VINS.SV — Scalar-Vector Insert	139
VINS.VV — Vector-Vector Insert	140
VIOTA — Vector Iota	142
VLD — Unit Stride Signed Vector Load	143
VLD.U — Unit Stride Unsigned Vector Load	146
VLDS — Variable Stride Signed Vector Load	148
VLDS.U — Variable Stride Unsigned Vector Load	151
VLDX — Indexed Signed Vector Load	153
VLDX.U — Indexed Unsigned Vector Load	155
VMADD.fmt — Vector Floating-Point Multiply Add	157
VMAX — Signed Vector Integer Maximum	161
VMAX.U — Unsigned Vector Integer Maximum	162
VMCOMMIT — Commit Speculative Arithmetic	163
VMCTS — Vector Move Control To Scalar	164
VMERGE — Vector Merge	165
VMIN — Signed Vector Integer Minimum	167
VMIN.U — Unsigned Vector Integer Minimum	168
VMOD — Signed Vector Integer Modulus	169
VMODU — Unsigned Vector Integer Modulus	171
VMSTC — Vector Move Scalar To Control	173
VMSUB.fmt — Vector Floating-Point Multiply Subtract	174
VMUL.fmt — Vector Floating-Point Multiply	178
VMULHI — Vector Integer Multiply High	182
VMULHI.U — Vector Integer Multiply High	183
VMULLO — Vector Integer Multiply Low	184
VNEG.fmt — Vector Floating-Point Negate	185
VNMADD.fmt — Vector Floating-Point Negative Multiply Add	189

VNMSUB.fmt — Vector Floating-Point Negative Multiply Subtract	193
VNOR — Vector Nor	197
VOR — Vector Or	198
VRECIP.fmt — Vector Floating-Point Reciprocal	199
VROUND — Vector Floating-Point Round	203
VRSQRT.fmt — Vector Floating-Point Reciprocal Square Root	209
VRSYNC — Vector Register Sync	213
VSADD — Signed Saturating Add	214
VSADD.U — Unsigned Saturating Add	216
VSAT — Signed Vector Saturate	218
VSAT.SU — Signed to Unsigned Vector Saturate	220
VSAT.U — Unsigned Vector Saturate	222
VSATVL — Saturate Vector Length	224
VSLL — Vector Shift Left Logical	225
VSLS — Signed Saturating Left Shift	227
VSLS.U — Unsigned Saturating Left Shift	229
VSQRT.fmt — Vector Floating-Point Square Root	231
VSRA — Vector Arithmetic Right Shift	235
VSRL — Vector Shift Right Logical	237
VSRR — Signed Shift Right And Round	239
VSRR.U — Unsigned Shift Right And Round	240
VSSUB — Signed Saturating Subtract	241
VSSUB.U — Signed Saturating Subtract	243
VST — Unit Stride Vector Store	245
VSTS — Variable Stride Vector Store	247
VSTX — Unordered Indexed Vector Store	249
VSTXO — Ordered Indexed Vector Store	251
VSUB — Signed Vector Integer Subtract	253
VSUB.U — Unsigned Vector Integer Subtract	255

VSUB.fmt — Vector Floating-Point Subtract	257
VSYNC — Vector Sync	261
VTLBP — Vector TLB Probe	262
VTLBR — Vector TLB Read	263
VTLBWI — Vector TLB Write Indexed	264
VTLBWR — Vector TLB Write Random	265
VTRUNC — Vector Floating-Point Truncate	266
VXLMADD — Signed Multiply Add Lower Halves	272
VXLMADD.U — Unsigned Multiply Add Lower Halves	275
VXLMSUB — Signed Multiply Subtract Lower Halves	278
VXLMSUB.U — Unsigned Multiply Subtract Lower Halves	281
VXMLMUL — Signed Multiply Lower Halves	284
VXMLMUL.U — Unsigned Multiply Lower Halves	286
VXOR — Vector Xor	288
VXUMADD — Signed Multiply Add Upper Halves	289
VXUMADD.U — Unsigned Multiply Add Upper Halves	292
VXUMSUB — Signed Multiply Subtract Upper Halves	295
VXUMSUB.U — Unsigned Multiply Subtract Upper Halves	298
VXUMUL — Signed Multiply Upper Halves	301
VXUMUL.U — Unsigned Multiply Upper Halves	303

List of Figures

1.1	Vector Architectural State	14
1.2	Data and Flag Register Element Mappings	15
1.3	Vector Control Registers	16
1.4	Vector Flag Register Contents	17
2.1	IEEE Floating-Point Comparisons	22
2.2	The Vector Convert Instruction	23
2.3	Fixed-Point Rounding Modes	24
2.4	Vector Instruction Set Summary	27
2.5	Instruction Qualifiers	28
2.6	Signed Vector Integer Arithmetic Instructions	29
2.7	Unsigned Vector Integer Arithmetic Instructions	30
2.8	Vector Logical Instructions	31
2.9	Vector Floating-Point Instructions	32
2.10	Vector Floating-Point Convert Instructions	33
2.11	Signed Vector Fixed-Point Instructions	34
2.12	Unsigned Fixed-Point Instructions	35
2.13	Vector Flag Logical Instructions	36
2.14	Vector Flag Processing Instructions	37
2.15	Vector Processing Instructions I	38
2.16	Vector Processing Instructions II	39
2.17	Vector Load Instructions	40
2.18	Vector Store Instructions	41

2.19 Coprocessor Interface Instructions	42
2.20 Miscellaneous Instructions	43
2.21 Miscellaneous Kernel-Mode Instructions	44

Chapter 1

Vector Architectural State

1.1 Overview

The vector architectural state is shown in Figure 1.1. A vector processor is a SIMD array of virtual processors (VP), where each VP has 32 data registers and 32 flag registers. The number of virtual processors is equal to the vector length (VL), which is never greater than the maximum vector length (MVL). MVL must be at least 16. A vector instruction performs the same operation on each VP in lockstep.

The VP width is set by the vector control register vc_{vpw} . MVL may be different for different VP widths, but the number of registers does not change. A register's contents are defined if and only if they are read and written with the same VP width.

Apart from the vector flag and data register files, there are two additional register files. The vector control register file is used in a variety of ways to control the operation of the vector unit (see Section 1.2). The vector scalar registers are used as source and destination registers for vector computations that require scalar sources and/or destinations.¹

1.2 Vector Control Registers

Figure 1.3 shows the contents of the 64 64-bit vector control registers. See the IRAM microarchitecture specification for a detailed description of these registers. Vector control register `vinc0` is fixed at zero.

1.3 Vector Scalar Registers

There are 32 64-bit vector scalar registers. All are general purpose registers. Vector scalar register 0 is fixed at zero.

¹A vector architecture that was not constrained by a coprocessor interface would probably use the MIPS general purpose integer registers and floating-point registers for this purpose.

1.4 Vector Flag Registers

Figure 1.4 shows the conventional usage of the flag registers. Either $\mathbf{vf}_{\text{mask}0}$ or $\mathbf{vf}_{\text{mask}1}$ is used as a mask for nearly all vector operations.

NOTE: More detail coming soon. May flags are used for exception processing. Many may be used as general purpose registers.

1.5 Number of Virtual Processors

The number of active virtual processors is given by the contents of the vector length register \mathbf{vc}_{v1} . The number of active VPs may vary anywhere between 0 and the implementation-defined maximum vector length, which is stored in the control register \mathbf{vc}_{mvl} . The architecture defines a minimum maximum vector length. This value is chosen to mediate two opposing pressures:

- MVL should be small to enable minimal implementations.
- MVL should be large so that more loops with a known number of iterations do not need to be strip-mined.

The minimum MVL is 16.

1.6 Virtual Processor Width

The width of each VP is given by the virtual processor width register \mathbf{vc}_{vpw} . Valid VP widths are 8, 16, 32, and 64 bits.² In order to make maximal use of available hardware resources, an implementation should increase MVL when VPW decreases. Figure 1.2 shows how MVL is doubled when VPW is halved in order to use all bits in the data register file for all VP widths. Moving to smaller VP widths causes the data elements to be fragmented into smaller pieces. Moving to larger VP widths causes the data elements to be catenated. Note that moving to smaller VP widths is always well defined. Moving to larger VP widths is defined for data elements only when the number of small elements is an exact multiple of the number of small elements per large element.

Note that the mapping of small VPs within larger VPs exposes the endian-ness of the implementation. When an implementation is big-endian, the mapping of small VPs within large VPs should also be big-endian. When an implementation is little-endian, the VP mapping should also be little-endian. For example, if 8-bit data is loaded into 64-bit VPs with 64-bit loads on a big-endian machine, the first 8-bit VP will contain the most-significant byte of the first 64-bit VP.

The mapping of flag register elements between different VP widths is more complex than for the data elements. For all VP widths, there are 32 1-bit flag registers per VP. Since MVL is larger for a smaller VPW, the number of bits of flag state is larger for a smaller VPW. This means that from an architectural point of view, the flag registers are different sizes for different VP widths. Because of this variability, the simple inclusion property that holds for the data registers does not work. Nevertheless, it is useful to define the state of the flag registers across VP width changes.

Moving to smaller VP widths causes flag elements to be replicated, and is always well defined. Moving to larger VP widths causes flag elements to be combined, under two restrictions. First, only identical flag values may be combined.

²The first implementation of this architecture – VIRAM-1 – will not implement 8-bit VPs.

An attempt to combine unlike flag values produces an undefined flag value. Second, the number of flags in the small VP width must be an exact multiple of the number of flags that combine to produce one flag in the large VP width.




Figure 1.1: Vector Architectural State

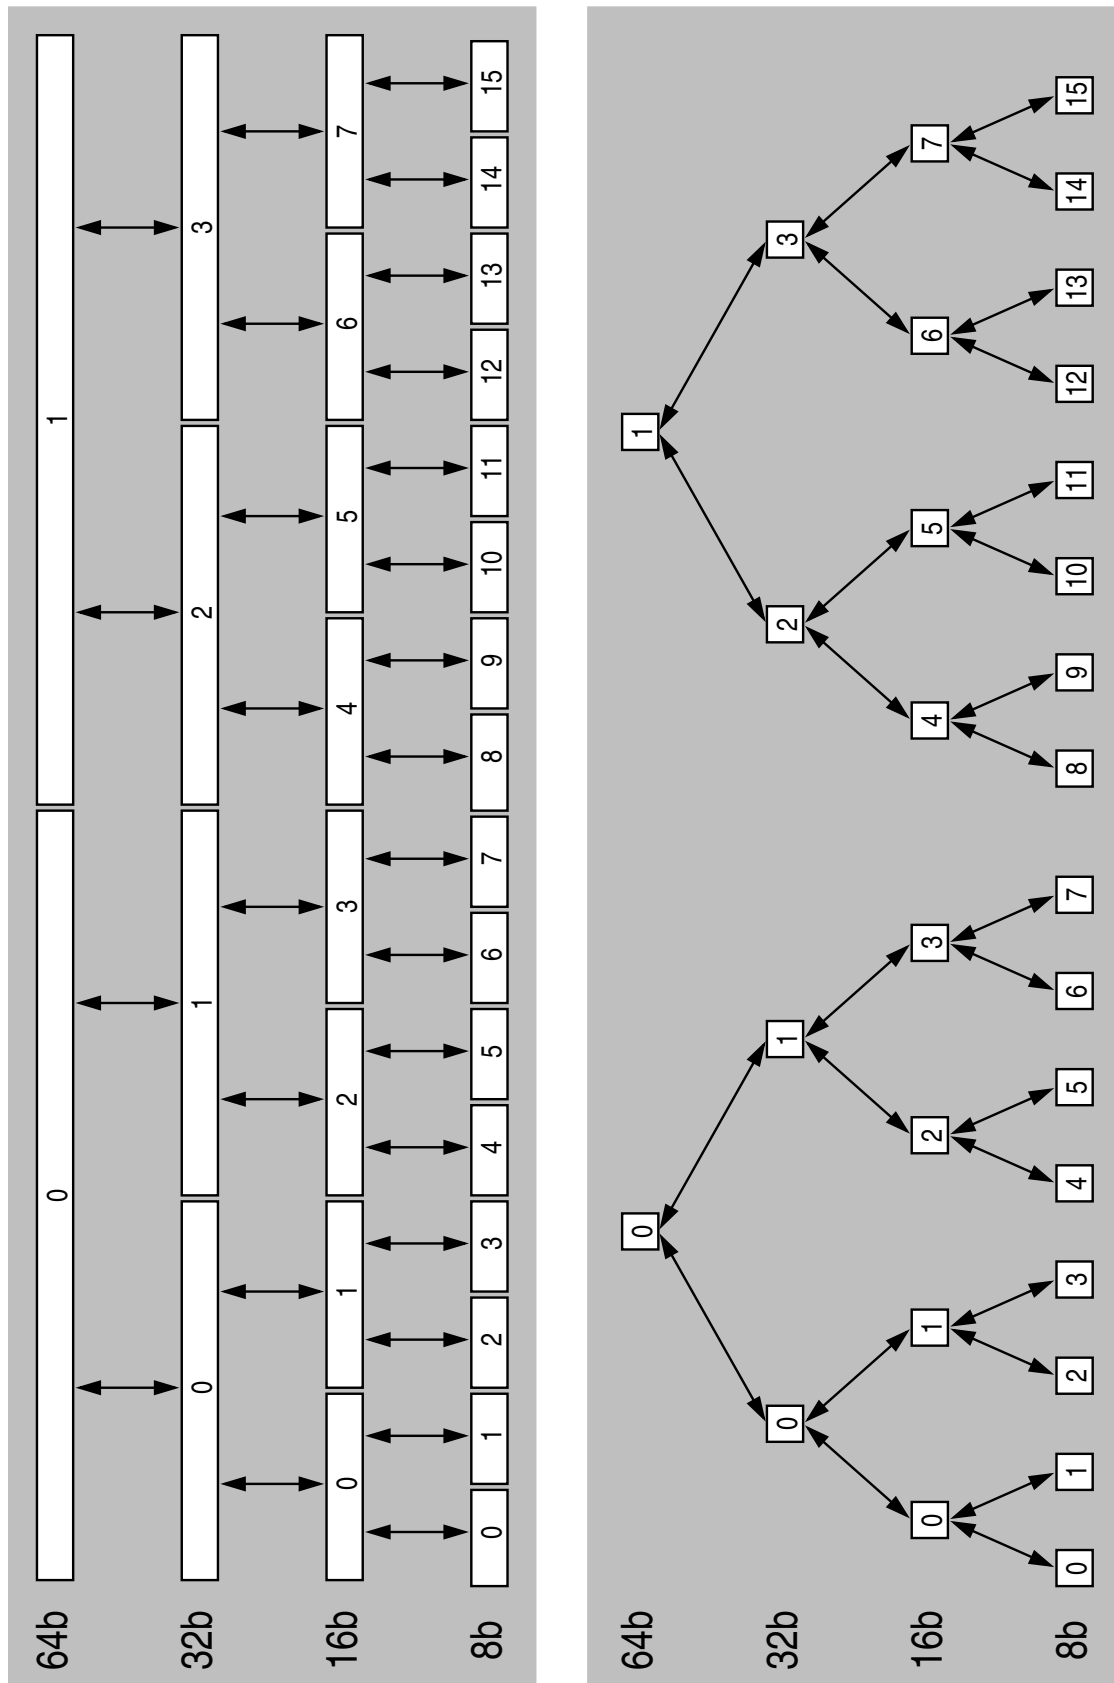


Figure 1.2: Data and Flag Register Element Mappings The upper portion of the figure shows the mapping of *data* elements when VP width changes. The lower portion of the figure shows the mapping of *flag* elements when VP width changes.

Hardware Name	Software Name	User Access	Kernel Access	Contents
\$vc0	v1	rw	rw	Vector length
\$vc1	vpw	rw	rw	Virtual processor width
\$vc2	vshamt	rw	rw	Shift amount
\$vc3	vindex	rw	rw	Element index
\$vc4	vcat	rw	rw	Assembler temporary
\$vc5	-	-	-	<i>Reserved</i>
\$vc6	vmode	rw	rw	Execution Mode
\$vc7	vregstatus	rw	rw	Vector register valid and dirty bits
\$vc8	vpsw	-	rw	Vector processor status word
\$vc9	vconfig	-	rw	I/O configuration
\$vc10	vcause	-	rw	Vector exception cause
\$vc11	vbadvaddr	-	rw	Bad virtual address
\$vc12	veiword	-	r	Instruction word for exceptions
\$vc13	vaiword	-	r	Instruction word for arithmetic exceptions
\$vc14	vwatch	-	rw	Physical watch address
\$vc15	vmemerror	-	rw	Memory error and status
\$vc16	vtlbindex	-	rw	Vector TLB index
\$vc17	vtlbbrandom	-	r	Vector TLB random index
\$vc18	vtlbentryhi	-	rw	Vector TLB entry hi
\$vc19	vtlbentrylo	-	rw	Vector TLB entry lo
\$vc20	vpagemask	-	rw	Vector TLB entry mask
\$vc21	vtlbwired	-	rw	Vector TLB wired entry index
\$vc22	vsafevl	r	rw	Safe vector length
\$vc23	vrev	r	r	Revision
\$vc24	mv1	r	r	Maximum vector length
\$vc25	logmv1	r	r	Base-2 logarithm of mv1
\$vc26	vcycle	r	rw	Cycle counter
\$vc27	-	-	-	<i>Reserved</i>
\$vc28	-	-	-	<i>Reserved</i>
\$vc29	-	-	-	<i>Reserved</i>
\$vc30	-	-	-	<i>Reserved</i>
\$vc31	-	-	-	<i>Reserved</i>
\$vc32	vbase0	rw	rw	Base Register 0
...
\$vc47	vbase15	rw	rw	Base Register 15
\$vc48	vinc0	r	r	Auto-Increment Register 0
\$vc49	vinc1	rw	rw	Auto-Increment Register 1
...
\$vc55	vinc7	rw	rw	Auto-Increment Register 7
\$vc56	vstride0	rw	rw	Stride Register 0
...
\$vc63	vstride7	rw	rw	Stride Register 7

Figure 1.3: Vector Control Registers

<i>Hardware Name</i>	<i>Software Name</i>	<i>Contents</i>
\$vf0	vfmask0	Primary mask
\$vf1	vfmask1	Alternate mask
\$vf2	vfgr0	General purpose
\$vf3	vfgr1	General purpose
\$vf4	vfgr2	General purpose
\$vf5	vfe0.L	Speculative load fault
\$vf6	vfe1.L	Speculative load fault
\$vf7	vfe2.L	Speculative load fault
\$vf8	vfe0.F	Integer overflow
\$vf9	vfe0.S	Integer saturate
\$vf10	vfe0.I	Floating-point inexact
\$vf11	vfe0.U	Floating-point underflow
\$vf12	vfe0.O	Floating-point overflow
\$vf13	vfe0.Z	Floating-point divide-by-zero
\$vf14	vfe0.V	Floating-point invalid
\$vf15	vfe0.E	Floating-point unimplemented
\$vf16	vfe1.F	Integer overflow
\$vf17	vfe1.S	Integer saturate
\$vf18	vfe1.I	Floating-point inexact
\$vf19	vfe1.U	Floating-point underflow
\$vf20	vfe1.O	Floating-point overflow
\$vf21	vfe1.Z	Floating-point divide-by-zero
\$vf22	vfe1.V	Floating-point invalid
\$vf23	vfe1.E	Floating-point unimplemented
\$vf24	vfe2.F	Integer overflow
\$vf25	vfe2.S	Integer saturate
\$vf26	vfe2.I	Floating-point inexact
\$vf27	vfe2.U	Floating-point underflow
\$vf28	vfe2.O	Floating-point overflow
\$vf29	vfe2.Z	Floating-point divide-by-zero
\$vf30	vfe2.V	Floating-point invalid
\$vf31	vfe2.E	Floating-point unimplemented

Figure 1.4: Vector Flag Register Contents

Chapter 2

Instruction Set

2.1 Open Questions

This section contains a list of random open questions about the instruction set architecture.

1. There will probably be some instructions relating to fast user-level message passing.
2. It is not clear how relaxed the vector memory consistency model should be. In particular, should a virtual processor see the order of its own memory references? This is called intra-VP consistency, or simply VP consistency in this document. What implementation could take advantage of relaxing this consistency guarantee?
3. Linkage conventions. MIPS registers have conventional uses for parameter passing, globals, etc. Such conventions will be defined soon for the vector, flag, and vector-scalar registers.

2.2 Data Types

The integer data width for an integer vector operation is always equal to the VP width. Integers loaded from memory are either sign- or zero-extended to the VP width.

Both 32-bit and 64-bit IEEE-754 standard floating-point formats are supported when the VP width is 64 bits. The 32-bit floating-point format is supported when the VP width is 32 bits. A 32-bit floating-point value is always stored in the low-order half of a 64-bit register. In this case, the upper 32 bits aren't touched by floating-point operations.

2.3 Vector Operations

Nearly all vector operations follow the same operational idiom. They are governed by a vector length and a vector mask. The vector length is taken from the vector length register. The vector mask is chosen from one of the two vector flag registers that are reserved for masking. Each virtual processor whose number is less than the vector length and whose mask bit is set performs the operation. If the vector length is zero, or if all VPs are masked out, then the instruction is a nop.

A VP takes its source operands from either a vector register or a scalar register. The destination is either a vector register or a scalar register, but is usually a vector register. Instructions whose source operands are both vectors are called *vector-vector* operations (qualified with $\cdot vv$). Instructions may also be *vector-scalar* (qualified with $\cdot vs$) if the second source is a scalar, or *scalar-vector* (qualified with $\cdot sv$), if the first source is a scalar. Non-commutative operations are provided in both scalar-vector and vector-scalar forms.

2.4 Flag Register Use

Each VP has 32 1-bit flag registers. The flag registers have two important uses: They are used to store mask vectors to provide predicated execution, and they are used to store exception bits.

A flag register reserved for exception processing is reserved by convention only. It may be used as a general purpose flag register in the absense of the relevant exception. This must be done with care, as the register is implicitly set by any instruction that may produce the relevant exception.

2.5 Memory Consistency

The memory consistency model may be best described as guaranteeing processor consistency for each virtual processor. Thus, each VP observes the order of its own memory operations, but it does not observe any guaranteed order of other VPs' memory operations. There are no ordering guarantees between scalar and vector instructions. Memory ordering must be explicitly enforced with vector sync instructions. These instructions control memory consistency between the scalar and vector unit and between virtual processors within the vector unit.

There is also a memory ordering guarantee between virtual processors of different widths. Operations between a narrow and a wide VP are ordered if the narrow VP is contained within the wide VP. For example, the memory operations of the 64-bit wide VP_0 , the 32-bit wide $VP_0 \cdots VP_1$, and the 16-bit wide $VP_1 \cdots VP_3$ are all ordered.

2.6 Exception Model

Coming soon...

2.7 Speculative Execution

Coming soon...

2.8 Assembly Language

Registers should be specified, wherever possible, using software names. The software names for vector flag registers are listed in Figure 1.4, and the software names for vector control registers are listed in Figure 1.3. There are currently no software names for vector or scalar data registers. The hardware names for the vector data registers are $\$vr0..\$vr31$. The hardware names for the scalar data registers are $\$vs0..\$vs31$.

2.9 Instructions

The vector instruction set is described in this document in three levels of detail. First, Figure 2.4 summarizes the entire instruction set in a single table. This table is divided into blocks of instructions of similar types. Each block is expanded into a more detailed table on the page named in the block's title. This collection of tables is the second level of detail, showing assembly formats and giving a short description of the instruction's operation.

The final level of detail is contained in Appendix B, where each instruction is defined in detail on a separate page. The first column of Figure 2.4 lists the pages in Appendix B on which the instructions are formally defined.

The assembly language format of an instruction is written with a shorthand notation, since the cross-product of legal qualifiers can produce many distinct opcodes. For example, the format of the vector floating-point add instruction in Figure 2.9 on Page 32 is shown as:

$$\text{vadd} \left\{ \begin{array}{l} .s \\ .d \end{array} \right\} \left\{ \begin{array}{l} .vv[.1] \\ .sv[.1] \end{array} \right\} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}}$$

The brackets [] denote an optional qualifier. The braces {} contain all choices for a required qualifier. Note that in this case, the assembly language register specifiers are contained within a set of braces so that they match the .vv/ .vs qualifier. Thus, the expanded list of formats for this instruction is:

```
vadd.s.vv vr_dest, vr_src1, vr_src2
vadd.s.vv.1 vr_dest, vr_src1, vr_src2
vadd.s.sv vr_dest, vs_src1, vr_src2
vadd.s.sv.1 vr_dest, vs_src1, vr_src2
vadd.d.vv vr_dest, vr_src1, vr_src2
vadd.d.vv.1 vr_dest, vr_src1, vr_src2
vadd.d.sv vr_dest, vs_src1, vr_src2
vadd.d.sv.1 vr_dest, vs_src1, vr_src2
```

2.9.1 Vector Integer Arithmetic Instructions

The vector integer arithmetic instructions are summarized in Figure 2.6 and Figure 2.7. All integer operations operate on VP-width operands and produce VP-width results.

Both signed and unsigned versions of some operations are provided. In most cases, this is because the computation is different depending on signedness (e.g. divide). However, both signed and unsigned add and subtract are provided, even though the computations are identical. The difference here is that signed add and subtract can overflow, while unsigned add and subtract never overflow under any circumstances.

2.9.2 Vector Logical Instructions

Figure 2.8 summarizes the vector logical operations. The usual set of logical operations is provided.

Predicate	Predicate true if...				Invalid if Unordered
	Unordered	Less Than	Equal	Greater Than	
(F)					
* LT		✓			✓
* EQ			✓		
* GT				✓	✓
* LE		✓	✓		✓
* GE			✓	✓	✓
LG		✓		✓	✓
LEG		✓	✓	✓	✓
† UN	✓				
ULT	✓	✓			
UEQ	✓		✓		
UGT	✓			✓	
ULE	✓	✓	✓		
UGE	✓		✓	✓	
(ULG)	✓	✓		✓	
(ULEG)	✓	✓	✓	✓	
(T)	✓	✓	✓	✓	
‡ NLT	✓		✓	✓	✓
* NEQ	✓	✓		✓	
‡ NGT	✓	✓	✓		✓
‡ NLE	✓			✓	✓
‡ NGE	✓	✓			✓
NLG	✓		✓		✓
NLEG	✓				✓
‡ NUN		✓	✓	✓	
NULT			✓	✓	
NUEQ		✓		✓	
NUGT		✓	✓		
NULE				✓	
NUGE		✓			
(NULG)			✓		
(NULEG)					

Figure 2.1: IEEE Floating-Point Comparisons

The full cross-product of the independent conditions *unordered*, *less than*, *equal*, and *greater than* produces 32 predicates, but only 26 are distinct. The 6 predicates in parentheses are either degenerate (F, ULEG, T, NULEG evaluate to either true or false) or redundant (ULG is identical to NEQ, and NULG is identical to EQ).

The 6 predicates marked with * are “obligatory in any [IEEE 754-1985] conforming implementation.” The predicate marked with † is “strongly recommended.” An implementation should provide “a means of logically negating predicates.” Negating all required and recommended predicates adds the 5 predicates marked with ‡.

Two operands are unordered if at least one is a NaN. The invalid operation exception is raised when unordered operands are compared using a predicate involving either *less than* or *greater than*, and not involving *unordered*. A predicate and its logical negation have the same exception behavior.

From Format		To Format			
		W	L	S	D
W				†‡	
L				‡	‡
S	†‡	‡			‡
D	‡	‡	‡		

Figure 2.2: The Vector Convert Instruction Empty locations represent conversions not provided by the `vcvt` instruction either because they are redundant (e.g. `vcvt.d.w` can be accomplished by `vcvt.d.l`) or because they are not floating-point conversions (e.g. `vcvt.w.l`). Conversions marked with † are valid for 32-bit VPs. Conversions marked with ‡ are valid for 64-bit VPs.

2.9.3 Vector Floating-Point Arithmetic Instructions

Figure 2.9 shows the support for floating-point arithmetic operations. The support conforms to the IEEE 754-1985 standard for binary floating-point arithmetic. The 32-bit single format and the 64-bit double format are supported for all operations.

Figure 2.1 shows all possible floating-point comparison predicates and those that should be provided by an IEEE 754-1985 compliant implementation. This architecture provides the 7 required and recommended predicates. The EQ, NEQ, LT, LE, and UN predicates are provided directly. The GT and GE predicates may be synthesized by swapping operands and using either LT or LT, respectively.

2.9.4 Vector Floating-Point Convert Instructions

Figure 2.10 shows the support for floating-point arithmetic operations. The support conforms to the IEEE 754-1985 standard for binary floating-point arithmetic. The `vcvt.a.b` instruction uses the default floating-point rounding mode to convert *to* format *a*, *from* format *b*. The set of valid conversions is shown in Figure 2.2. Floating-point values may be converted to to signed integer format using an explicit rounding mode with the `vtrunc`, `vround`, `vceil`, and `vfloor` instructions.

2.9.5 Vector Fixed-Point Arithmetic Instructions

The fixed-point instructions are chosen to provide DSP functionality in a general purpose computing device. Many of the features of a DSP architecture are already provided by this vector architecture or are not needed at all:

- *Fixed ↔ Floating-Point Conversion* instructions already exist.
- *Multiple Loads and Stores Per Cycle* may be provided by an implementation with multiple memory units.
- *Auto Increment / Decrement* is built into strided vector memory operations.
- *Circular Addressing* is not needed. The cost of coding circular addressing is amortized over long vector operations.
- *Bit-Reverse Addressing* is not needed, as alternate FFT algorithms may be used.

Name	Action
Truncate	$x.y \dots \Rightarrow x$
Round Up	$x.0y \dots \Rightarrow x$ $x.1y \dots \Rightarrow x + 1$
Round to Nearest Even	if $xy.1$ [exactly] then $x0.1 \Rightarrow x0$ $x1.1 \Rightarrow x1 + 1$ else $x.0y \dots \Rightarrow x$ $x.1y \dots \Rightarrow x + 1$ endif
Jam	$xy.z \dots \Rightarrow xw$ where $w = y \parallel z \parallel \dots$

Figure 2.3: Fixed-Point Rounding Modes

Some remaining DSP features need special architectural support, namely rounding, saturation, and special multiply instructions. Figure 2.11 and Figure 2.12 list the fixed-point instructions. Many of the operations involve shifting off low-order bits. For these instructions, this shift is accompanied by a rounding step. The rounding mode is one of the four modes described in Figure 2.3.

Since multiply-add is such a frequent operation in fixed-point codes, this operation (along with multiply-subtract) is provided as a basic operation, though it can be synthesized from more basic instructions. Note that fixed-point multiply instructions are different from normal integer multiply low or multiply high. Multiply low and high multiply their entire n -bit operands, and select either the lower or upper half of the $2n$ -bit result. In contrast, the fixed-point multiply instructions multiply either the lower or upper $n/2$ -bit halves of the n -bit operands to produce an n -bit result.

2.9.6 Vector Logical Flag Instructions

Figure 2.13 shows the simple flag processing instructions. This includes the standard logical operations plus instructions to set and clear registers. The scalar-vector flag instructions take their scalar operand from a vector scalar register. This register is not treated as a bit vector; the operand is false if the register is zero, and true otherwise. The `vfmf8` and `vfmt8` move blocks of 8 flag registers to and from a vector register. These instructions exist in order to save and restore the flag register set efficiently for context switching. None of these instructions operate under mask, but they all operate under vector length.

2.9.7 Flag Processing Instructions

Figure 2.14 lists the more complex flag processing instructions. Some of these instructions – `vfpop`, `vfffl1`, `vffl1` – place information about a flag register in a scalar register. Others – `vfsetbf`, `vfsetif`, `vfsetof` – perform complex operations directly on the flag registers. The remaining instructions – `viota` and `vciota` – are used to generate index vectors in a vector register from a flag register.

2.9.8 Vector Processing Instructions

Figure 2.15 and Figure 2.16 list the more complex vector processing instructions. All of these instructions are either not masked, or use the vector mask in an unusual manner.

Many of these instructions replicate functionality in order to gain performance. The `vcompress` and `vexpand` instructions can perform any operation that can be done with `vins.vv`, `vext.vv`, `vhalfup`, `vhalfdn`, and `vexthalf`. Their generality comes at a performance cost, so more restricted instructions are provided for specific applications. The `vhalfup` and `vhalfdn` can be used to perform butterfly permutations. The `vexthalf` instruction is designed specifically to accelerate reductions.

2.9.9 Vector Memory Instructions

The instructions to transfer data between vector registers and memory are listed in Figure 2.17 and Figure 2.18. There are two categories of vector memory instructions: *strided* and *indexed*. A strided load takes a base address and a signed stride, and loads a vector of values starting at the base address, where each element is separated by the stride amount. The base address is taken from the vector control register `vbase`, and the stride is taken from the vector control register `vstride`. *Stride is in units of elements, not bytes*. A special case is *unit stride*, which loads a contiguous vector from memory. Unit stride loads and stores are provided as separate instructions because they are the most common type of strided access.

Note that both negative and zero strides are legal. Zero stride is the most obvious case where the order of operations in a strided access can be observed. In fact, virtual memory can cause the order of operations of non-zero stride stores to be observed as well. If a strided store straddles a page boundary, then the indirection provided by virtual memory can create physical address aliases that are not virtual address aliases. Zero-stride is still a special case, since the address aliases existed in the virtual address space, while non-zero-stride aliases only exist in the physical address space. However, all strided stores are treated equally with respect to memory ordering: *All strided stores occur in VP order*.

An indexed memory operation takes a scalar base address, and a vector of byte (not element) offsets. A scalar-vector add is performed to yield a vector of addresses. The elements are loaded / stored using these addresses. Since aliases can exist, it is important to specify the order in which operations occur. The ordered indexed store (`vstxo`) stores elements in element order. The `vstx` instruction store elements in an undefined order.

Every vector memory operation specifies a data width in the opcode. This is the width of data in memory. This width must be less than or equal to the VP width. If it is less than the VP width, then the value is zero-extended for unsigned loads and sign-extended for signed loads. Floating-point values should be loaded with unsigned loads.

The instructions used to transfer flag registers directly to and from memory are also shown in Figure 2.17 and Figure 2.18. The vector flag store instruction stores a single flag register as a packed bit vector in memory, where the bytes are stored in a little-endian manner (i.e. the first 8 bits of the bit vector are at a lower address than the next 8 bits). It should be possible to strip-mine a loop in order to create a contiguous bit vector in memory. The bit vector always begins at the beginning of a half-word, but may end in the middle of another half-word. If a flag store does not have a multiple of 16 bits to store, then it may pad with zeros in order to store a multiple of 16 bits. Note that MINMVL is also 16, and this is not a coincidence. It is advantageous to an implementation to pad and align to larger boundaries, but the boundary should not be more than MINMVL bits.

2.9.10 Coprocessor Interface Instructions

The standard MIPS coprocessor move instructions are used in the vector unit to move data between the MIPS general purpose registers and the vector unit's control and scalar registers. The vector unit occupies coprocessor 2, so `cfc2`, `mtc2`, and `ctc2` are used. `mfc2` is not used because it causes problems with exceptions and the coprocessor interface in VIRAM-1.

2.9.11 Miscellaneous Vector Instructions

Figure 2.20 lists miscellaneous user-mode vector instructions that do not fit other categories. This includes an instruction to aid strip-mining (`vsatv1`), instructions to move data between the vector unit's control and scalar register files, instructions to commit speculative operations, and instructions to enforce memory ordering.

2.9.12 Miscellaneous Kernel Vector Instructions

Figure 2.21 lists various instructions that may be executed only in kernel mode. These instructions are used to manipulate the TLB and to flush implementation-specific state when the vector unit is frozen.

Page	Mnemonic	Operation	Page	Mnemonic	Operation
Integer Arithmetic (29,30)			Flag Logical (36)		
51	vabs	Absolute Value	110	vfand	And
57,59	vadd	Add	128	vfor	Or
253,255	vsub	Subtract	138	vf xor	Xor
184	vmul lo	Multiply Low	127	vf nor	Nor
182,183	vmul hi	Multiply High	111	vfclr	Clear
91,93	vdiv	Divide	132	vfset	Set
169,171	vmod	Modulus	Flag Processing (37)		
235	vsra	Arithmetic Right Shift	142	viota	Iota
71,73	vcmp	Compare	70	vciota	Continuous Iota
167,168	vmin	Minimum	130	vfpop	Population Count
161,162	vmax	Maximum	113	vfff1	Find First One
Logical (31)			114	vffl1	Find Last One
64	vand	And	133	vfsetbf	Set Before First One
198	vor	Or	134	vfsetif	Set Including First One
288	vxor	Xor	135	vfsetof	Set Only First One
197	vnor	Nor	126	vfmt8	Move To 8
225	vsll	Logical Left Shift	125	vfmf8	Move From 8
237	vsrl	Logical Right Shift	112	vfclr8	Clear 8
Floating-Point Arithmetic (32,33)			129	vfor8	Or 8
53	vabs	Absolute Value	Vector Processing (38,39)		
60	vadd	Add	139,140	vins	Vector Insert
257	vsub	Subtract	102,101,100	vext	Vector Extract
178	vmul	Multiply	79	vcompress	Compress
157	vmadd	Multiply Add	99	vexpand	Expand
174	vmsub	Multiply Subtract	165	vmerge	Merge
189	vmadd	Negative Multiply Add	115	vfits	Scalar Insert
193	vmsub	Negative Multiply Subtract	103	vexthalf	Extract Half
95	vdiv	Divide	104	vhalf	Half
231	vsqrt	Square Root	108	vhalfup	Half Up
209	vrsqrt	Reciprocal Square Root	106	vhalfdn	Half Down
199	vrecip	Reciprocal	Miscellaneous User (43)		
185	vneg	Negate	164	vmcts	Move Control To Scalar
75	vcmp	Compare	173	vmstc	Move Scalar To Control
80	vcvt	Convert	224	vsatvl	Saturate Vector Length
266	vtrunc	Truncate	56	vacomm it	Commit Speculative Arithmetic
203	vround	Round	163	vmcomm it	Commit Speculative Arithmetic
65	vceil	Ceiling	261	vsync	Sync
119	vfloor	Floor	213	vr sync	Register Sync
Fixed-Point Arithmetic (34,35)			Miscellaneous Kernel (44)		
218,222,220	vsat	Saturate	262	vtlbp	TLB Probe
214,216	vsadd	Saturating Add	263	vtlbr	TLB Read
241,243	vssub	Saturating Subtract	264	vtlbwi	TLB Write Indexed
239,240	vsrr	Shift Right And Round	265	vtlbwr	TLB Write Random
227,229	vslls	Saturating Left Shift	131	vflush	Flush Vector Unit
301,303	vxumul	Multiply Upper Halves	Load (40)		
284,286	vxlmul	Multiply Lower Halves	116	vfld	Flag Load
289,292	vxumadd	Multiply Add Upper Halves	143,146	vld	Unit Stride Load
295,298	vxumsub	Multiply Subtract Upper Halves	148,151	vlds	Variable Stride Load
272,275	vxlmadd	Multiply Add Lower Halves	153,155	vldx	Indexed Load
278,281	vxlmsub	Multiply Subtract Lower Halves	Store (41)		
Coprocessor Interface (42)			136	vfst	Flag Store
48	cfc2	Control From Cop2	245	vst	Unit Stride Store
50	mtc2	Move To Cop2	247	vsts	Variable Stride Store
49	ctc2	Control To Cop2	249	vstx	Unordered Indexed Store
			251	vstxo	Ordered Indexed Store

Figure 2.4: Vector Instruction Set Summary Each sub-block within this table expands into the more detailed table(s) on the page(s) specified in the sub-block title. The instructions are defined in full on the page(s) named in the first column.

<i>Qualifier</i>	<i>Meaning</i>	<i>Notes</i>
<i>op.vv</i> <i>op.vs</i> <i>op.sv</i>	Vector-Vector Vector-Scalar Scalar-Vector	Vector arithmetic instructions may take up to one source operand from a scalar register. A vector-vector operation takes two vector source operands; a vector-scalar operation takes its second operand from the scalar register file; a scalar-vector operation takes its first operand from the scalar register file.
<i>op.b</i> <i>op.h</i> <i>op.w</i> <i>op.l</i>	1B Byte 2B Halfword 4B Word 8B Longword	The saturate instruction, floating-point convert instructions, and all vector memory instructions need to specify the width of integer data.
<i>op.u</i>	Unsigned	Where sensible, integer arithmetic operations are provided in both signed and unsigned versions. The unsigned versions are identified by applying this qualifier.
<i>op.s</i> <i>op.d</i>	4B Single FP Format 8B Double FP Format	Floating-point instructions need to specify data type.
<i>op.l</i>	Use vf_{mask1} as the mask.	By default, the vector mask is taken from vf_{mask0} . This qualifier enables vf_{mask1} as vector mask.
<i>op.lt</i> <i>op.eq</i> <i>op.le</i> <i>op.neq</i> <i>op.un</i>	Less Than Equal Greater Than Less Than Or Equal Greater Than Or Equal Not Equal Unordered	These are the predicates used by the integer and floating-point compare instructions.

Figure 2.5: Instruction Qualifiers

Page	Operation	Assembly	Summary
51	Absolute Value	$\text{vabs}[\cdot.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP writes into vr_{dest} the absolute value of vr_{src} .
57	Add	$\text{vadd} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer sum of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
253	Subtract	$\text{vsub} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{vs}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer subtraction of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
184	Multiply Low	$\text{vmullo} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the lower half of the full-VP-width integer product of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
182	Multiply High	$\text{vmulhi} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the upper half of the full-VP-width signed integer product of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
91	Divide	$\text{vdiv} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{vs}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer quotient of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
169	Modulus	$\text{vmod} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{vs}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer modulus of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
235	Arithmetic Right Shift	$\text{vsra} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{vs}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ right-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is sign-extended.
71	Compare	$\text{vcmp} \begin{cases} .\text{eq} \\ .\text{neq} \end{cases} \begin{cases} .\text{vv}[\cdot.1] & \text{vf}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vf}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$ $\text{vcmp} \begin{cases} .\text{lt} \\ .\text{le} \end{cases} \begin{cases} .\text{vv}[\cdot.1] & \text{vf}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vf}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{vs}[\cdot.1] & \text{vf}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vf_{dest} the signed integer comparison of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.
167	Minimum	$\text{vmin} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the lesser of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as signed integers.
161	Maximum	$\text{vmax} \begin{cases} .\text{vv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .\text{sv}[\cdot.1] & \text{vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the greater of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as signed integers.

Figure 2.6: Signed Vector Integer Arithmetic Instructions

Page	Operation	Assembly	Summary
59	Add	$\text{vadd.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the unsigned integer sum of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
255	Subtract	$\text{vsub.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the unsigned integer subtraction of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
183	Multiply High	$\text{vmulhi.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the upper half of the full-VP-width unsigned integer product of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
93	Divide	$\text{vdiv.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the unsigned integer quotient of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
171	Modulus	$\text{vmod.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the unsigned integer modulus of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.
73	Compare	$\text{vcmp.u} \begin{cases} .lt \\ .le \end{cases} \begin{cases} .vv[.1] & \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vf_{dest} the unsigned integer comparison of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.
168	Minimum	$\text{vmin.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the lesser of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as unsigned integers.
162	Maximum	$\text{vmax.u} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the greater of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as unsigned integers.

Figure 2.7: Unsigned Vector Integer Arithmetic Instructions

Page	Operation	Assembly	Summary
64	And	$\text{vand} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each VP writes into vr_{dest} the bit-wise logical <i>and</i> of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
198	Or	$\text{vor} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each VP writes into vr_{dest} the bit-wise logical <i>or</i> of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
288	Xor	$\text{vxor} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each VP writes into vr_{dest} the bit-wise logical <i>xor</i> of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
197	Nor	$\text{vnor} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$	Each VP writes into vr_{dest} the bit-wise logical <i>nor</i> of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .
225	Logical Left Shift	$\text{vsl} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{vs}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ left-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is zero-filled.
237	Logical Right Shift	$\text{vsr} \begin{cases} .\text{vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .\text{vs}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$	Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ right-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is zero-extended.

Figure 2.8: Vector Logical Instructions

Page	Operation	Assembly	Summary
53	Absolute Value	$\text{vabs}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places the floating-point absolute value of vr_{src} into vr_{dest} .
60	Add	$\text{vadd}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP places the floating-point sum of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and vr_{src2} into vr_{dest} .
257	Subtract	$\text{vsub}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \\ .vs [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{matrix}$	Each unmasked VP places the floating-point difference of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and $\text{vr}_{\text{src2}}/\text{vf}_{\text{src2}}$ into vr_{dest} , where at least one source is a vector.
178	Multiply	$\text{vmul}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP places the floating-point product of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and vr_{src2} into vr_{dest} .
157	Multiply Add	$\text{vmadd}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP adds the floating-point product of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and vr_{src2} into vr_{dest} .
174	Multiply Subtract	$\text{vmsub}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP subtracts vr_{dest} from the floating-point product of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$.
189	Negative Multiply Add	$\text{vnmadd}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP adds the floating-point product of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and vr_{src2} into vr_{dest} , and negates the result.
193	Negative Multiply Subtract	$\text{vnmsub}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP subtracts vr_{dest} from the floating-point product of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$, and negates the result.
95	Divide	$\text{vdiv}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \\ .vs [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{matrix}$	Each unmasked VP places the floating-point quotient of $\text{vr}_{\text{src1}}/\text{vs}_{\text{src1}}$ and $\text{vr}_{\text{src2}}/\text{vf}_{\text{src2}}$ into vr_{dest} , where at least one source is a vector.
231	Square Root	$\text{vsqrt}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the floating-point square root of vr_{src} .
209	Reciprocal Square Root	$\text{vrsqrt}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the floating-point reciprocal square root of vr_{src} .
199	Reciprocal	$\text{vrecip}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the floating-point reciprocal of vr_{src} .
185	Negate	$\text{vneg}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the floating-point square root of vr_{src} .
75	Compare	$\text{vcmp}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .f \\ .sf \\ .un \\ .ngle \\ .eq \\ .seq \\ .ueq \\ .ngl \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$ $\text{vcmp}\left\{\begin{smallmatrix} .s \\ .d \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .olt \\ .lt \\ .ult \\ .nge \\ .ole \\ .le \\ .ule \\ .ngt \end{smallmatrix}\right\} \left\{\begin{smallmatrix} .vv [.1] \\ .sv [.1] \end{smallmatrix}\right\} \begin{matrix} \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix}$	Each unmasked VP writes into vf_{dest} the floating-point comparison of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.

Figure 2.9: Vector Floating-Point Instructions

Page	Operation	Assembly	Summary
80	Convert	$\text{vcvt.s.} \left\{ \begin{smallmatrix} .d \\ .w \\ .l \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$ $\text{vcvt.d.} \left\{ \begin{smallmatrix} .s \\ .l \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$ $\text{vcvt.w.} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$ $\text{vcvt.l.} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of converting vr_{src} from one integer or floating-point format to another.
266	Truncate	$\text{vtrunc} \left\{ \begin{smallmatrix} .w \\ .l \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of converting vr_{src} from a floating-point format to a signed integer format, using the <i>truncate</i> rounding mode.
203	Round	$\text{vround} \left\{ \begin{smallmatrix} .w \\ .l \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of converting vr_{src} from a floating-point format to a signed integer format, using the <i>round</i> rounding mode.
65	Ceiling	$\text{vceil} \left\{ \begin{smallmatrix} .w \\ .l \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of converting vr_{src} from a floating-point format to a signed integer format, using the <i>ceiling</i> rounding mode.
119	Floor	$\text{vfloor} \left\{ \begin{smallmatrix} .w \\ .l \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of converting vr_{src} from a floating-point format to a signed integer format, using the <i>floor</i> rounding mode.

Figure 2.10: Vector Floating-Point Convert Instructions

Page	Operation	Assembly	Summary
218	Saturate	$\text{vsat} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of saturating vr_{src} to a signed integer narrower than the VP width. The result is sign-extended to the VP width.
220	Saturate	$\text{vsat.su} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of saturating vr_{src} from a signed VP width value to an unsigned value that is as wide or narrower than the VP width. The result is zero-extended to the VP width.
214	Saturating Add	$\text{vsadd} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP writes into vr_{dest} the signed integer sum of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . The sum saturates to the VP width instead of overflowing.
241	Saturating Subtract	$\text{vssub} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .vs[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP writes into vr_{dest} the signed integer subtraction of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The difference saturates to the VP width instead of overflowing.
239	Shift Right And Round	$\text{vsrr} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP writes into vr_{dest} the right arithmetic shift of vr_{src} . The result is rounded as per the fixed-point rounding mode. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
227	Saturating Left Shift	$\text{vsls} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src}}$	Each unmasked VP writes into vr_{dest} the signed saturating left shift of vr_{src} . The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
301	Multiply Upper Halves	$\text{vxumul} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after an arithmetic right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
284	Multiply Lower Halves	$\text{vxlmul} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after an arithmetic right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
289	Multiply Add Upper Halves	$\text{vxumadd} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is added into vr_{dest} after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
295	Multiply Subtract Upper Halves	$\text{vxumsub} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is subtracted from vr_{dest} after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
272	Multiply Add Lower Halves	$\text{vxlmadd} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is added into vr_{dest} after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
278	Multiply Subtract Lower Halves	$\text{vxlmsub} \left\{ \begin{matrix} .vv[.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ .sv[.1] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \end{matrix} \right\}$	Each unmasked VP computes the signed integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is subtracted from vr_{dest} after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

Figure 2.11: Signed Vector Fixed-Point Instructions

Page	Operation	Assembly	Summary
222	Saturate	$\text{vsat.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP places into vr_{dest} the result of saturating vr_{src} to an unsigned integer narrower than the VP width. The result is zero-extended to the VP width.
216	Saturating Add	$\text{vsadd.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP writes into vr_{dest} the unsigned integer sum of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . The sum saturates to the VP width instead of overflowing.
243	Saturating Subtract	$\text{vssub.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP writes into vr_{dest} the unsigned integer subtraction of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The difference saturates to zero instead of underflowing.
240	Shift Right And Round	$\text{vsrr.u} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP writes into vr_{dest} the right logical shift of vr_{src} . The result is rounded as per the fixed-point rounding mode. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
229	Saturating Left Shift	$\text{vsls.u} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$	Each unmasked VP writes into vr_{dest} the unsigned saturating left shift of vr_{src} . The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
303	Multiply Upper Halves	$\text{vxumul.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after a logical right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
286	Multiply Lower Halves	$\text{vxlmul.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after a logical right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
292	Multiply Add Upper Halves	$\text{vxumadd.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is added into vr_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
298	Multiply Subtract Upper Halves	$\text{vxumsub.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is subtracted from vr_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
275	Multiply Add Lower Halves	$\text{vxlmadd.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is added into vr_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.
281	Multiply Subtract Lower Halves	$\text{vxlmsub.u} \left\{ \begin{matrix} .vv[.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{matrix} \right.$	Each unmasked VP computes the unsigned integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is subtracted from vr_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

Figure 2.12: Unsigned Fixed-Point Instructions

Page	Operation	Assembly	Summary
110	And	$\text{vfind} \begin{cases} .vv \text{ vfd}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .sv \text{ vfd}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$	Each VP writes into vfd_{dest} the logical <i>and</i> of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.
128	Or	$\text{vfor} \begin{cases} .vv \text{ vfd}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .sv \text{ vfd}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$	Each VP writes into vfd_{dest} the logical <i>or</i> of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.
138	Xor	$\text{vfxor} \begin{cases} .vv \text{ vfd}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .sv \text{ vfd}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$	Each VP writes into vfd_{dest} the logical <i>xor</i> of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.
127	Nor	$\text{vfnor} \begin{cases} .vv \text{ vfd}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .sv \text{ vfd}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$	Each VP writes into vfd_{dest} the logical <i>nor</i> of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.
111	Clear	$\text{vfclr } \text{vfd}_{\text{dest}}$	Each VP writes zero into vfd_{dest} . This instruction is not masked.
132	Set	$\text{vfset } \text{vfd}_{\text{dest}}$	Each VP writes one into vfd_{dest} . This instruction is not masked, but is subject to vector length.

Figure 2.13: Vector Flag Logical Instructions

Page	Operation	Assembly	Summary
142	Iota	<code>viota vr_dest, vf_src</code>	The list of VPs with bits set in <code>vf_src</code> is placed in <code>vr_dest</code> . This is the compressed index vector of <code>vf_src</code> . This instruction is not masked.
70	Continuous Iota	<code>vciota vr_dest, vf_src</code>	The continuous index vector of <code>vf_src</code> is placed in register <code>vr_dest</code> . The continuous index vector's value at <code>vp_i</code> is equal to the population count of flag values for <code>vp₀ ... vp_{i-1}</code> . This instruction is not masked.
130	Population Count	<code>vfpop vs_dest, vf_src</code>	The population count of <code>vf_src</code> is placed in <code>vs_dest</code> . This instruction is not masked.
113	Find First One	<code>vfffl vs_dest, vf_src</code>	The location of the first set bit of <code>vf_src</code> is placed in <code>vs_dest</code> . This instruction is not masked. If there is no set bit in <code>vf_src</code> , then the vector length is placed in <code>vs_dest</code> .
114	Find Last One	<code>vffll vs_dest, vf_src</code>	The location of the last set bit of <code>vf_src</code> is placed in <code>vs_dest</code> . This instruction is not masked. If there is no set bit in <code>vf_src</code> , then the vector length is placed in <code>vs_dest</code> .
133	Set Before First One	<code>vfsetbf vf_dest, vf_src</code>	Register <code>vf_dest</code> is filled with ones up to and not including the first set bit in <code>vf_src</code> . Remaining positions in <code>vf_dest</code> are cleared. If <code>vf_src</code> contains no set bits, <code>vf_dest</code> is set to all ones. This instruction is not masked.
134	Set Including First One	<code>vfsetif vf_dest, vf_src</code>	Register <code>vf_dest</code> is filled with ones up to and including the first set bit in <code>vf_src</code> . Remaining positions in <code>vf_dest</code> are cleared. If <code>vf_src</code> contains no set bits, <code>vf_dest</code> is set to all ones. This instruction is not masked.
135	Set Only First One	<code>vfsetof vf_dest, vf_src</code>	Register <code>vf_dest</code> is filled with zeros except for the position of the first set bit in <code>vf_src</code> . If <code>vf_src</code> contains no set bits, <code>vf_dest</code> is set to all zeros. This instruction is not masked.
126	Move To 8	<code>vfmt8 vf_dest, vr_src</code>	Registers <code>vf_dest ... vf_dest+7</code> are set with the contents of <code>vr_src</code> . This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.
125	Move From 8	<code>vfmf8 vr_dest, vf_src</code>	Registers <code>vf_src ... vf_src+7</code> are copied into <code>vr_dest</code> . This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.
112	Clear 8	<code>vfclr8 vf_dest</code>	Each VP writes zero into <code>vf_dest ... vf_dest+7</code> . This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.
129	Or 8	<code>vfor8 vf_dest, vf_src</code>	Each VP performs an OR-write of <code>vf_src ... vf_src+7</code> into <code>vf_dest ... vf_dest+7</code> . This instruction is not masked, but is subject to vector length. Each flag register specifier must be a multiple of 8.

Figure 2.14: Vector Flag Processing Instructions

Page	Operation	Assembly	Summary
140	Vector Insert	<code>vins.vv vr_dest, vr_src2</code>	The leading portion of <code>vr_src</code> is inserted into <code>vr_dest</code> . <code>vr_dest</code> must be different than <code>vr_src</code> . Leading and trailing entries of <code>vr_dest</code> are not touched. The lower <code>vc_logmv1</code> bits of vector control register <code>vc_index</code> specifies the starting position in <code>vr_dest</code> . The vector length specifies the number of elements to transfer. This instruction is not masked.
139	Scalar Insert	<code>vins.sv vr_dest, vs_src</code>	The contents of <code>vs_src</code> are written into <code>vr_dest</code> at position <code>vc_index</code> . The lower <code>vc_logmv1</code> bits of <code>vc_index</code> are used. This instruction is not masked and does not use vector length.
102	Vector Extract	<code>vext.vv vr_dest, vr_src</code>	A portion of <code>vr_src</code> is extracted into the front of <code>vr_dest</code> . <code>vr_dest</code> must be different than <code>vr_src</code> . Trailing entries of <code>vr_dest</code> are not touched. The lower <code>vc_logmv1</code> bits of vector control register <code>vc_index</code> specifies the starting position in <code>vr_src</code> . The vector length specifies the number of elements to transfer. This instruction is not masked.
100	Scalar Extract	<code>vext.sv vs_dest, vr_src2</code>	Element <code>vc_index</code> of <code>vr_src</code> is written into <code>vs_dest</code> . The lower <code>vc_logmv1</code> bits of <code>vc_index</code> are used. The value is sign-extended. This instruction is not masked and does not use vector length.
101	Scalar Extract	<code>vext.u.sv vs_dest, vr_src</code>	Element <code>vc_index</code> of <code>vr_src</code> is written into <code>vs_dest</code> . The lower <code>vc_logmv1</code> bits of <code>vc_index</code> are used. The value is zero-extended. This instruction is not masked and does not use vector length.
79	Compress	<code>vcompress[.1] vr_dest, vr_src</code>	All unmasked elements of <code>vr_src</code> are catenated to form a vector whose length is the population count of the mask (subject to vector length). The result is placed at the front of <code>vr_dest</code> , leaving trailing elements untouched. <code>vr_dest</code> must be different than <code>vr_src</code> .
99	Expand	<code>vexpand[.1] vr_dest, vr_src</code>	The first n elements of <code>vr_src</code> are written into the unmasked positions of <code>vr_dest</code> , where n is the population count of the mask (subject to vector length). Masked positions in <code>vr_dest</code> are not touched. <code>vr_dest</code> must be different than <code>vr_src</code> .
165	Merge	$\text{vmerge} \begin{cases} .vv[.1] & \text{vr_dest}, \text{vr_src1}, \text{vr_src2} \\ .sv[.1] & \text{vr_dest}, \text{vs_src1}, \text{vr_src2} \\ .vs[.1] & \text{vr_dest}, \text{vr_src1}, \text{vs_src2} \end{cases}$	Each VP copies into <code>vr_dest</code> either <code>vs_src1/vr_src1</code> if the mask is 0, or <code>vs_src2/vr_src2</code> if the mask is 1. At least one source is a vector. Scalar sources are truncated to the virtual processor width.
115	Scalar Insert	<code>vfins vf_dest, vs_src</code>	The boolean value of <code>vs_src</code> is written into <code>vf_dest</code> at position <code>vc_index</code> . The lower <code>vc_logmv1</code> bits of <code>vc_index</code> are used. This instruction is not masked and does not use vector length.

Figure 2.15: Vector Processing Instructions I

Page	Operation	Assembly	Summary
103	Extract Half	<code>vexthalf vr_dest, vr_src</code>	The top half of <code>vr_src</code> is extracted into the front of <code>vr_dest</code> . <code>vr_dest</code> must be different than <code>vr_src</code> . Trailing entries of <code>vr_dest</code> are not touched. The vector length specifies the length of <code>vr_src</code> . This instruction is not masked. The vector length must be a power of two, and must be at least two.
104	Half	<code>vhalf vr_dest, vr_src</code>	The top half of <code>vr_src</code> is extracted into the front of <code>vr_dest</code> . <code>vr_dest</code> must be different than <code>vr_src</code> . Trailing entries of <code>vr_dest</code> are not touched. The vector length specifies the length of <code>vr_src</code> . This instruction is not masked. The vector length must be a power of two, and must be at least two. After execution, the vector length register is divided by two.
108	Half Up	<code>vhalfup vr_dest, vr_src</code>	This instruction does one half of one step of a butterfly permutation. Paired with <code>vhalfdn</code> , it can be used to perform an in-register butterfly permutation. <code>vr_dest</code> must be different than <code>vr_src</code> . <code>vc_vindex</code> specifies the log of both the length and the separation of the sub-vectors that are moved from <code>vr_src</code> to <code>vr_dest</code> . The sub-vectors are shifted up within each $2 * 2^{vc_vindex}$ block by 2^{vc_vindex} elements. This writes half of the elements within each block in <code>vr_dest</code> ; the remaining elements are not touched. This instruction is not masked. <code>vc_vl</code> must be at least 2 and must be a power of 2. 2^{vc_vindex} must be less than <code>vc_vl</code> .
106	Half Down	<code>vhalfdn vr_dest, vr_src</code>	This instruction does one half of one step of a butterfly permutation. Paired with <code>vhalfup</code> , it can be used to perform an in-register butterfly permutation. <code>vr_dest</code> must be different than <code>vr_src</code> . <code>vc_vindex</code> specifies the log of both the length and the separation of the sub-vectors that are moved from <code>vr_src</code> to <code>vr_dest</code> . The sub-vectors are shifted down within each $2 * 2^{vc_vindex}$ block by 2^{vc_vindex} elements. This writes half of the elements within each block in <code>vr_dest</code> ; the remaining elements are not touched. This instruction is not masked. <code>vc_vl</code> must be at least 2 and must be a power of 2. 2^{vc_vindex} must be less than <code>vc_vl</code> .

Figure 2.16: Vector Processing Instructions II

Page	Operation	Assembly	Summary
116	Flag Load	<code>vfld vf_{dest}[, vbase[, vinc]]</code>	The VPs perform a contiguous vector flag load into <code>vf_{dest}</code> . <code>vf_{dest}</code> should be different than the speculative load fault flag register if the instruction is speculative. The base address is given by <code>vbase</code> (default is <code>vbase₀</code>), and must be 16-bit aligned. The bytes are loaded in little-endian order. The signed increment in <code>vinc</code> (default is <code>vinc₀</code>) is added to <code>vbase</code> as a side-effect. This instruction is not masked.
143	Unit Stride Load	$\text{vld} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}[, \text{vbase}[, \text{vinc}]]$	The VPs perform a contiguous vector load into <code>vr_{dest}</code> . The base address is given by <code>vbase</code> (default is <code>vbase₀</code>), and must be aligned to the width of the data in memory. The signed increment in <code>vinc</code> (default is <code>vinc₀</code>) is added to <code>vbase</code> as a side-effect. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.
146	Unit Stride Load	$\text{vld.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}[, \text{vbase}[, \text{vinc}]]$	Operation is identical to <code>vld</code> , except that loaded values are zero-extended to the virtual processor width instead of sign-extended.
148	Variable Stride Load	$\text{vlds} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}[, \text{vbase}[, \text{vstride}[, \text{vinc}]]]$	The VPs perform a strided vector load into <code>vr_{dest}</code> . The base address is given by <code>vbase</code> (default is <code>vbase₀</code>), and must be aligned to the width of the data in memory. The <i>signed</i> stride is given by <code>vstride</code> (default is <code>vstride₀</code>). The stride is in terms of elements, not in terms of bytes. The signed increment in <code>vinc</code> (default is <code>vinc₀</code>) is added to <code>vbase</code> as a side-effect. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.
151	Variable Stride Load	$\text{vlds.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}[, \text{vbase}[, \text{vstride}[, \text{vinc}]]]$	Operation is identical to <code>vlds</code> , except that loaded values are zero-extended to the virtual processor width instead of sign-extended.
153	Indexed Load	$\text{vldx} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{offsets}}[, \text{vbase}]$	The VPs perform an indexed vector load into <code>vr_{dest}</code> . The base address is given by <code>vbase</code> (default is <code>vbase₀</code>). The <i>signed</i> offsets are given by <code>vr_{offsets}</code> . The offsets are in units of bytes, not in units of elements. The effective addresses must be aligned to the width of the data in memory. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.
155	Indexed Load	$\text{vldx.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{offsets}}[, \text{vbase}]$	Operation is identical to <code>vldx</code> , except that loaded values are zero-extended to the virtual processor width instead of sign-extended.

Figure 2.17: Vector Load Instructions

Page	Operation	Assembly	Summary
136	Flag Store	<code>vfst $\mathbf{vf}_{src}[, \mathbf{vbase}[, \mathbf{vinc}]]$</code>	The VPs perform a contiguous vector flag store of \mathbf{vf}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be 16-bit aligned. A multiple of 16 bits is written, padding with zeros as necessary. The bytes are stored in little-endian order. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. This instruction is not masked.
245	Unit Stride Store	<code>$\mathbf{vst} \left\{ \begin{smallmatrix} .b \\ .h \\ .w \\ .l \end{smallmatrix} \right\} [.1] \mathbf{vr}_{src}[, \mathbf{vbase}[, \mathbf{vinc}]]$</code>	The VPs perform a contiguous vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The register value is truncated from the VP width to the memory width. The VPs access memory in order.
247	Variable Stride Store	<code>$\mathbf{vsts} \left\{ \begin{smallmatrix} .b \\ .h \\ .w \\ .l \end{smallmatrix} \right\} [.1] \mathbf{vr}_{src}[, \mathbf{vbase}[, \mathbf{vstride}[, \mathbf{vinc}]]]$</code>	The VPs perform a contiguous vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The <i>signed</i> stride is given by $\mathbf{vstride}$ (default is $\mathbf{vstride}_0$). The stride is in terms of elements, not in terms of bytes. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The register value is truncated from the VP width to the memory width. The VPs access memory in order.
249	Unordered Indexed Store	<code>$\mathbf{vstx} \left\{ \begin{smallmatrix} .b \\ .h \\ .w \\ .l \end{smallmatrix} \right\} [.1] \mathbf{vr}_{src}, \mathbf{vr}_{offsets}[, \mathbf{vbase}]$</code>	The VPs perform an indexed vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0). The <i>signed</i> offsets are given by $\mathbf{vr}_{offsets}$. The offsets are in units of bytes, not in units of elements. The effective addresses must be aligned to the width of the data in memory. The register value is truncated from the VP width to the memory width. The stores may be performed in any order.
251	Ordered Indexed Store	<code>$\mathbf{vstxo} \left\{ \begin{smallmatrix} .b \\ .h \\ .w \\ .l \end{smallmatrix} \right\} [.1] \mathbf{vr}_{src}, \mathbf{vr}_{offsets}[, \mathbf{vbase}]$</code>	Operation is identical to \mathbf{vstx} , except that the VPs access memory in order.

Figure 2.18: Vector Store Instructions

<i>Page</i>	<i>Operation</i>	<i>Assembly</i>	<i>Summary</i>
48	Control From Cop2	<code>cfc2 r_{dest}, vc_{src}</code>	<code>vc_{src}</code> is copied into register <code>r_{dest}</code> .
50	Move To Cop2	<code>mtc2 r_{src}, vs_{dest}</code>	<code>r_{src}</code> is copied into register <code>vs_{dest}</code> .
49	Control To Cop2	<code>ctc2 r_{src}, vc_{dest}</code>	<code>r_{src}</code> is copied into register <code>vc_{dest}</code> . Writing <code>vc_{vpz}</code> changes <code>vc_{mv1}</code> , <code>vc_{logmv1}</code> , and <code>vc_{vsafev1}</code> as a side-effect.

Figure 2.19: Coprocessor Interface Instructions

Page	Operation	Assembly	Summary
164	Move Scalar To Control	<code>vmstc vc_{dest}, vs_{src}</code>	Register <code>vs_{src}</code> is copied to <code>vc_{dest}</code> . Writing <code>vc_{vpw}</code> changes <code>vc_{mv1}</code> , <code>vc_{logmv1}</code> , and <code>vc_{vsafev1}</code> as a side-effect. This instruction may write only <code>vc_{v1}</code> , <code>vc_{vpw}</code> , <code>vc_{vshamt}</code> , <code>vc_{vindex}</code> , and <code>vc_{vcat}</code> .
173	Move Control To Scalar	<code>vmcts vs_{dest}, vc_{src}</code>	Register <code>vc_{src}</code> is copied to <code>vs_{dest}</code> .
224	Saturate Vector Length	<code>vsatv1</code>	The vector length register is saturated to the maximum vector length.
56	Commit Speculative Arithmetic	<code>vacommit[.exact][.1]</code>	Any pending vector arithmetic exceptions generated by speculative operations are raised. This instruction operates under mask and vector length. The exceptions are precise if either <code>.exact</code> is specified, or if <code>vc_{vmode}</code> specifies exact exceptions.
163	Commit Speculative Arithmetic	<code>vmcommit[.exact][.1]</code>	Any pending vector memory exceptions generated by speculative operations are raised. This instruction operates under mask and vector length. The exceptions are precise if either <code>.exact</code> is specified, or if <code>vc_{vmode}</code> specifies exact exceptions.
261	Sync	$vsync \begin{Bmatrix} .sav \\ .vas \\ .vav \\ .vp \end{Bmatrix} \begin{Bmatrix} \phi \\ .raw \\ .war \\ .waw \end{Bmatrix}$	Enforces scalar-after-vector, vector-after-scalar, vector-after-vector, or intra-vp memory ordering. If none of <code>raw/war/waw</code> are specified, then all are in effect. All relevant memory references preceeding the sync must appear to execute before all relevant memory references following the sync. Order is defined as program order.
213	Register Sync	$vrsync \begin{Bmatrix} .sav \\ .vav \\ .vp \end{Bmatrix} \begin{Bmatrix} \phi \\ .raw \\ .war \\ .waw \end{Bmatrix} vr_{dest}$	Enforces scalar-after-vector, vector-after-scalar, vector-after-vector, or intra-vp memory ordering. If none of <code>raw/war/waw</code> are specified, then all are in effect. The most recent memory reference to <code>vr_{dest}</code> preceeding the sync must appear to execute before all relevant memory references following the sync. Order is defined as program order.

Figure 2.20: Miscellaneous Instructions


<i>Page</i>	<i>Operation</i>	<i>Assembly</i>	<i>Summary</i>
262	TLB Probe	vtlbp	This instruction is not yet implemented.
263	TLB Read	vtlbr	This instruction is not yet implemented.
264	TLB Write Indexed	vtlbwi	This instruction is not yet implemented.
265	TLB Write Random	vtlbwr	This instruction is not yet implemented.
131	Flush Vector Unit	vflush	The vector unit must be frozen. All in-flight instructions are flushed. The register files are not touched.

Figure 2.21: Miscellaneous Kernel-Mode Instructions

Appendix A


Instruction Formats

The vector register format (VR-Type) specifies three registers and a 10-bit opcode in the coprocessor-2 opcode space. The 10-bit vector opcode is split into two pieces:




Note that single-source VR-type instructions use field `src1` to specify the source register.



There are two instruction formats that are used by the coprocessor interface instructions to move data between the scalar and vector units:



There are several instruction formats used by the instructions that move data between the vector scalar and vector control register files:




Each class of vector memory access – unit-stride, constant-stride, and indexed – has its own instruction format:



NOTE: The vsync instruction uses some special formats, to be defined soon...

Appendix B

Instruction Definitions

CFC2**Control From Coprocessor 2****Assembly**

```
cfc2 rdest, vCsrc
```

Operation

```
if (!user_readable (VC[src])) { /* assume user mode for now */
    Raise vIUI;
}
x = pdinst->get_src_vc (src);
GPR[dest] = x;
```

Description

vC_{src} is copied into register r_{dest}.

Exceptions


vIUI Illegal Use of Instruction Exception

Notes

None.

Control To Coprocessor 2

CTC2



Assembly

```
ctc2 rsrc, vcdest
```

Operation

```
if (!user_writable (VC[dest])) { /* assume user mode for now */
    Raise vIUI;
}
x = pdinst->get_src_gpr (src);
if (dest == VcVpw) {
    x &= 0x3; /* force vpw to a legal value */
    VC[vpw] = x;
    VC[mvl] = MVL[x];
    VC[logmvl] = LOGMVL[x];
} else {
    VC[dest] = x;
}
```

Description

r_{src} is copied into register vc_{dest} . Writing vc_{vpw} changes vc_{mvl} , vc_{logmvl} , and $vc_{vsafevl}$ as a side-effect.

Exceptions


vInt	Vector Integer Exception
vFPE	Vector Floating-Point Exception
vIUI	Illegal Use of Instruction Exception

Notes

None.

MTC2

Move To Coprocessor 2



Assembly

```
mtc2  $r_{src}$ ,  $vs_{dest}$ 
```

Operation

```
x = pdinst->get_src_gpr (src);
VS[dest] = x;
```

Description

r_{src} is copied into register vs_{dest} .

Exceptions


None.

Notes

None.

Vector Integer Absolute Value

VABS



Assembly

```
vabs[.1] vr_dest, vr_src
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
raise_F = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = abs (x);
        if (abs_overflow (x, z)) {
            f = true;
            if (!speculative && F_enabled) {
                raise_F = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_F][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_F][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_F][vp] = f;
}
if (raise_F) {
    Raise vAri;
}

```

Description

Each unmasked VP writes into `vr_dest` the absolute value of `vr_src`.

VABS(cont.)**Vector Integer Absolute Value****Exceptions**


vIVL Invalid Vector Length Exception

Notes

This operation is only defined for signed operands.

Vector Floating-Point Absolute Value

VABS.fmt



Assembly

$$\text{vabs} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = abs_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VABS.fmt(cont.)

Vector Floating-Point Absolute Value

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = abs_d(x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Absolute Value

VABS.fmt(cont.)

Description

Each unmasked VP places the floating-point absolute value of \mathbf{vr}_{src} into \mathbf{vr}_{dest} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

VACOMMIT

Commit Speculative Arithmetic



Assembly

```
vacommit[.exact][.1]
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
F = S = I = U = O = Z = V = E = false;
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        F = F || (VF[vfe_F][vp] && F_enabled);
        S = S || (VF[vfe_S][vp] && S_enabled);
        I = I || (VF[vfe_I][vp] && I_enabled);
        U = U || (VF[vfe_U][vp] && U_enabled);
        O = O || (VF[vfe_O][vp] && O_enabled);
        Z = Z || (VF[vfe_Z][vp] && Z_enabled);
        V = V || (VF[vfe_V][vp] && V_enabled);
        E = E || VF[vfe_E][vp];
    }
}
if (F || S || I || U || O || Z || V || E) {
    Raise vAri;
}

```

Description

Any pending vector arithmetic exceptions generated by speculative operations are raised. This instruction operates under mask and vector length. The exceptions are precise if either `.exact` is specified, or if `vcvmode` specifies exact exceptions.

Exceptions


None.

Notes

None.

Signed Vector Integer Add

VADD



Assembly

```
vadd { .vv[.1] vr_dest, vr_src1, vr_src2
      .sv[.1] vr_dest, vs_src1, vr_src2
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
raise_F = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x + y;
        if (add_overflow (x, y, z)) {
            f = true;
            if (!speculative && F_enabled) {
                raise_F = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_F][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_F][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_F][vp] = f;
}
if (raise_F) {
    Raise vAri;
}
```

VADD(cont.)**Signed Vector Integer Add****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the signed integer sum of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Unsigned Vector Integer Add

VADD.U



Assembly

```
vadd.u { .vv[.1] vr_dest, vr_src1, vr_src2
        .sv[.1] vr_dest, vs_src1, vr_src2
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x + y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
```

Description

Each unmasked VP writes into `vr_dest` the unsigned integer sum of `vs_src1/vr_src1` and `vr_src2`.

Exceptions


`vIVL` Invalid Vector Length Exception

Notes

The only difference between `vadd` and `vadd.u` is that `vadd` can overflow, while `vadd.u` cannot.

VADD.fmt

Vector Floating-Point Add



Assembly

$$\text{vadd} \begin{cases} .s \\ .d \end{cases} \begin{cases} .vv[.1] \\ .sv[.1] \end{cases} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vadd} \begin{cases} .s \\ .d \end{cases} \begin{cases} .vv[.1] \\ .sv[.1] \end{cases} \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}}$$

Vector Floating-Point Add

VADD.fmt(cont.)

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = add_s (x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VADD.fmt(cont.)

Vector Floating-Point Add

Operation (Double Precision)

```

    if (vl > mvl) {
        Raise vIVL;
    }
    if (vpw < 3) {
        Raise vIUI;
    }
    raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
    for (vp = 0; vp < mvl; vp++) {
        I = U = O = Z = V = E = false;
        if (vp < vl && VF[mask][vp]) {
            if (vv) {
                x = VR[src1][vp];
                y = VR[src2][vp];
            } else {
                x = VS[src1];
                y = VR[src2][vp];
            }
            z = add_d(x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
            if (!speculative) {
                write = true;
                if (I && I_enabled) { raise_I = true; write = false; }
                if (U && U_enabled) { raise_U = true; write = false; }
                if (O && O_enabled) { raise_O = true; write = false; }
                if (Z && Z_enabled) { raise_Z = true; write = false; }
                if (V && V_enabled) { raise_V = true; write = false; }
                if (E) { raise_E = true; write = false; }
                if (!write) {
                    z = VR[dest][vp];
                }
            }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }

```

Vector Floating-Point Add**VADD.fmt(cont.)****Description**


Each unmasked VP places the floating-point sum of $\mathbf{vr}_{src1}/\mathbf{vs}_{src1}$ and \mathbf{vr}_{src2} into \mathbf{vr}_{dest} .

Exceptions

\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

Vector And



Assembly

$$\text{vand} \begin{cases} \text{.VV[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.SV[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x & y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
}

```

Description

Each VP writes into `vrdest` the bit-wise logical *and* of `vssrc1`/`vrsrc1` and `vrsrc2`.

Exceptions


vIVL	Invalid Vector Length Exception
------	---------------------------------

Notes

None.

Vector Floating-Point Ceiling

VCEIL



Assembly

$$\text{vceil} \left\{ \begin{smallmatrix} .w \\ .l \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} .s \\ .d \end{smallmatrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Operation (Single to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = ceil_s_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCEIL(cont.)

Vector Floating-Point Ceiling

Operation (Double to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = ceil_d_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Ceiling

VCEIL(cont.)

Operation (Single to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = ceil_s_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCEIL(cont.)

Vector Floating-Point Ceiling

Operation (Double to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = ceil_d_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Ceiling

VCEIL(cont.)

Description

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of converting \mathbf{vr}_{src} from a floating-point format to a signed integer format, using the *ceiling* rounding mode.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

VCIOTA

Vector Continuous Iota



Assembly

```
vciota vrdest, vfsrc
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
count = 0;
for (vp = 0; vp < vl; vp++) {
    VR[dest][vp] = count;
    if (VF[src][vp]) {
        count++;
    }
}
for (vp = vl; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

The continuous index vector of $\mathbf{vf}_{\text{src1}}$ is placed in register $\mathbf{vr}_{\text{dest}}$. The continuous index vector's value at \mathbf{vp}_i is equal to the population count of flag values for $\mathbf{vp}_0 \cdots \mathbf{vp}_{i-1}$. This instruction is not masked.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

The continuous and compressed index vectors of a flag register of all 1's are the same.

Signed Vector Integer Compare

VCMP



Assembly

$$\text{vcmp} \begin{cases} \text{.eq} \\ \text{.neq} \end{cases} \begin{cases} \text{.vv[.1]} \\ \text{.sv[.1]} \end{cases} \begin{cases} \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

$$\text{vcmp} \begin{cases} \text{.lt} \\ \text{.le} \end{cases} \begin{cases} \text{.vv[.1]} \\ \text{.sv[.1]} \\ \text{.vs[.1]} \end{cases} \begin{cases} \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        b = x OP y;
    } else {
        b = VF[dest][vp];
    }
    VF[dest][vp] = b;
}

```

Description

Each unmasked VP writes into vf_{dest} the signed integer comparison of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.

VCMP(cont.)

Signed Vector Integer Compare

Exceptions


vIVL Invalid Vector Length Exception

Notes

Commutative operations are not provided in `.vs` form.

Unsigned Vector Integer Compare

VCMP.U



Assembly

$$\text{vcmp.u} \begin{cases} .lt \\ .le \end{cases} \begin{cases} .vv[.1] & \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vf}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        b = x OP y;
    } else {
        b = VF[dest][vp];
    }
    VF[dest][vp] = b;
}

```

Description

Each unmasked VP writes into vf_{dest} the unsigned integer comparison of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.

VCMP.U(cont.)**Unsigned Vector Integer Compare****Exceptions**


`vIVL` Invalid Vector Length Exception

Notes

Commutative operations are not provided in `.vs` form. The `.eq` and `.neq` are provided only as signed comparisons, since those operations are independent of signedness.

Vector Floating-Point Compare

VCMP.fmt



Assembly

$$\begin{aligned} \text{vcmp} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} & \left\{ \begin{matrix} .ngle \\ .eq \\ .seq \\ .ueq \\ .ngl \end{matrix} \right\} & \left\{ \begin{matrix} .vv[.1] \text{ } \mathbf{vf}_{\text{dest}} , \mathbf{vr}_{\text{src1}} , \mathbf{vr}_{\text{src2}} \\ .sv[.1] \text{ } \mathbf{vf}_{\text{dest}} , \mathbf{vs}_{\text{src1}} , \mathbf{vr}_{\text{src2}} \end{matrix} \right. \\ \\ \text{vcmp} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} & \left\{ \begin{matrix} .olt \\ .lt \\ .ult \\ .nge \\ .ole \\ .le \\ .ule \\ .ngt \end{matrix} \right\} & \left\{ \begin{matrix} .vv[.1] \text{ } \mathbf{vf}_{\text{dest}} , \mathbf{vr}_{\text{src1}} , \mathbf{vr}_{\text{src2}} \\ .sv[.1] \text{ } \mathbf{vf}_{\text{dest}} , \mathbf{vs}_{\text{src1}} , \mathbf{vr}_{\text{src2}} \end{matrix} \right. \end{aligned}$$

VCMP.fmt(cont.)

Vector Floating-Point Compare

Operation (Single Precision)

```

    if (vl > mvl) {
        Raise vIVL;
    }
    if (vpw < 2) {
        Raise vIUI;
    }
    raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
    for (vp = 0; vp < mvl; vp++) {
        I = U = O = Z = V = E = false;
        if (vp < vl && VF[mask][vp]) {
            if (vv) {
                x = VR[src1][vp];
                y = VR[src2][vp];
            } else if (sv) {
                x = VS[src1];
                y = VR[src2][vp];
            } else {
                x = VR[src1][vp];
                y = VS[src2];
            }
            b = cmp_s_<op> (x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
            if (!speculative) {
                write = true;
                if (I && I_enabled) { raise_I = true; write = false; }
                if (U && U_enabled) { raise_U = true; write = false; }
                if (O && O_enabled) { raise_O = true; write = false; }
                if (Z && Z_enabled) { raise_Z = true; write = false; }
                if (V && V_enabled) { raise_V = true; write = false; }
                if (E) { raise_E = true; write = false; }
                if (!write) {
                    b = VF[dest][vp];
                }
            }
        } else {
            b = VF[dest][vp];
        }
        VF[dest][vp] = b;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }

```

Vector Floating-Point Compare

VCMP.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        b = cmp_d_<op> (x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                b = VF[dest][vp];
            }
        }
    } else {
        b = VF[dest][vp];
    }
    VF[dest][vp] = b;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCMP.fmt(cont.)**Vector Floating-Point Compare****Description**

Each unmasked VP writes into `vfdest` the floating-point comparison of `vssrc1/vrsrc1` and `vssrc2/vrsrc2`, where at least one source is a vector. A true predicate yields 1; a false predicate yields 0.

Exceptions


<code>vIVL</code>	Invalid Vector Length Exception
<code>vIUI</code>	Illegal Use of Instruction Exception

Notes

Commutative operations are not provided in `.vs` form.

Vector Compress

VCOMPRESS



Assembly

```
vcompress[.1] vr_dest, vr_src
```

Operation

```
if (vl > mvl) {
    Raise vIUI;
}
if (src == dest) {
    Raise vIUI;
}
count = 0;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        VR[dest][count++] = x;
    }
}
for (vp = count; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

All unmasked elements of \mathbf{vr}_{src} are catenated to form a vector whose length is the population count of the mask (subject to vector length). The result is placed at the front of \mathbf{vr}_{dest} , leaving trailing elements untouched. \mathbf{vr}_{dest} must be different than \mathbf{vr}_{src} .

Exceptions


vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

Compress is the opposite of expand. The vector length applies to \mathbf{vr}_{src} and the mask.

VCVT

Vector Convert



Assembly

$$\begin{aligned}
 &\text{vcvt.s.} \left\{ \begin{matrix} .d \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{vr}_{\text{src}} \\
 &\text{vcvt.d.} \left\{ \begin{matrix} .s \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{vr}_{\text{src}} \\
 &\text{vcvt.w.} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{vr}_{\text{src}} \\
 &\text{vcvt.l.} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{vr}_{\text{src}}
 \end{aligned}$$

Vector Convert

VCVT(cont.)

Operation (Double to Single)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_d_to_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCVT(cont.)

Vector Convert

Operation (Word to Single)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_w_to_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Convert

VCVT(cont.)

Operation (Long to Single)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_l_to_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCVT(cont.)

Vector Convert

Operation (Single to Double)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_s_to_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Convert

VCVT(cont.)

Operation (Long to Double)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_l_to_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCVT(cont.)

Vector Convert

Operation (Single to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_s_to_w (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Convert

VCVT(cont.)

Operation (Double to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_d_to_w (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCVT(cont.)

Vector Convert

Operation (Single to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_s_to_l (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```


Vector Convert

VCVT(cont.)

Operation (Double to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = cvt_d_to_l (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VCVT(cont.)

Vector Convert

Description

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of converting \mathbf{vr}_{src} from one integer or floating-point format to another.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

The difference between `vcvt` and the other floating-point conversion instructions (`vtrunc`, `vfloor`, `vround`, and `vceil`) is that `vcvt` implicitly uses the default rounding mode (specified in vector control register $\mathbf{vc}_{\text{mode}}$), where the other instructions use explicit rounding modes.

Signed Vector Integer Divide

VDIV



Assembly

$$\text{vdiv} \begin{cases} \text{.vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        if (y == 0) {
            z = VR[dest][vp];
        } else {
            z = x / y;
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the signed integer quotient of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.

VDIV(cont.)

Signed Vector Integer Divide

Exceptions


$vIVL$ Invalid Vector Length Exception

Notes

The result is undefined if the divisor is zero.

Unsigned Vector Integer Divide

VDIVU



Assembly

$$\text{vdiv.u} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        if (y == 0) {
            z = VR[dest][vp];
        } else {
            z = x / y;
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the unsigned integer quotient of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.

VDIVU(cont.)

Unsigned Vector Integer Divide

Exceptions


$vIVL$ Invalid Vector Length Exception

Notes

The result is undefined if the divisor is zero.

Vector Floating-Point Divide

VDIV.fmt



Assembly

$$\text{vdiv} \begin{cases} .s \\ .d \end{cases} \begin{cases} .vv[.1] \\ .sv[.1] \\ .vs[.1] \end{cases} \begin{cases} \text{vr}_{\text{dest}} , & \text{vr}_{\text{src1}} , & \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}} , & \text{vs}_{\text{src1}} , & \text{vr}_{\text{src2}} \\ \text{vr}_{\text{dest}} , & \text{vr}_{\text{src1}} , & \text{vs}_{\text{src2}} \end{cases}$$

VDIV.fmt(cont.)

Vector Floating-Point Divide

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = div_s(x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }
}

```

Vector Floating-Point Divide

VDIV.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = div_d (x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }
}

```

VDIV.fmt(cont.)

Vector Floating-Point Divide

Description

Each unmasked VP places the floating-point quotient of $\mathbf{vr}_{src1}/\mathbf{vs}_{src1}$ and $\mathbf{vr}_{src2}/\mathbf{vf}_{src2}$ into \mathbf{vr}_{dest} , where at least one source is a vector.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Expand

VEXPAND



Assembly

```
vexpand[.1] vr_dest, vr_src
```

Operation

```

if (vl > mvl) {
    Raise vIUI;
}
if (src == dest) {
    Raise vIUI;
}
count = 0;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][count++];
    } else {
        x = VR[dest][vp];
    }
    VR[dest][vp] = x;
}

```

Description

The first n elements of \mathbf{vr}_{src} are written into the unmasked positions of $\mathbf{vr}_{\text{dest}}$, where n is the population count of the mask (subject to vector length). Masked positions in $\mathbf{vr}_{\text{dest}}$ are not touched. $\mathbf{vr}_{\text{dest}}$ must be different than \mathbf{vr}_{src} .

Exceptions


vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

Expand is the opposite of compress. The vector length applies to $\mathbf{vr}_{\text{dest}}$ and the mask.

VEXT.SV

Signed Scalar-Vector Extract



Assembly

```
vext.sv vsdest, vrsrc2
```

Operation

```
vindex &= mvl - 1; /* force vindex into valid range */
x = VR[src][vindex];
VS[dest] = x;
```

Description

Element vc_{vindex} of vr_{src} is written into vs_{dest} . The lower vc_{logmvl} bits of vc_{vindex} are used. The value is sign-extended. This instruction is not masked and does not use vector length.

Exceptions


None.

Notes

None.

Unsigned Scalar-Vector Extract

VEXT.U.SV



Assembly

```
vext.u.sv vsdest, vrsrc
```

Operation

```
vindex &= mvl - 1; /* force vindex into valid range */
x = VR[src][vindex];
VS[dest] = x;
```

Description

Element vc_{vindex} of vr_{src} is written into vs_{dest} . The lower vc_{logmvl} bits of vc_{vindex} are used. The value is zero-extended. This instruction is not masked and does not use vector length.

Exceptions


None.

Notes

Floating-point extraction should be done with this instruction.

VEXT.VV

Vector-Vector Extract



Assembly

```
vext.vv vrdest, vrsrc
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
vindex &= mvl - 1; /* force vindex into valid range */
for (vp = 0; vp < vl && vp + vindex < mvl; vp++) {
    x = VR[src][vp + vindex];
    VR[dest][vp] = x;
}
for (; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}

```

Description

A portion of vr_{src} is extracted into the front of vr_{dest} . vr_{dest} must be different than vr_{src} . Trailing entries of vr_{dest} are not touched. The lower $\text{vc}_{\log\text{mvl}}$ bits of vector control register $\text{vc}_{\text{vindex}}$ specifies the starting position in vr_{src} . The vector length specifies the number of elements to transfer. This instruction is not masked.

Exceptions


vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

If $\text{vindex} + \text{vl} \geq \text{mvl}$, then the only elements copied are those whose source vp number is legal. Vector-vector extract is the opposite of vector-vector insert.

Vector Extract Half

VEXTHALF



Assembly

```
vexthalf vrdest, vrsrc
```

Operation

```
if (vl > mvl || vl == 1 || !is_pow2 (vl)) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
for (vp = 0; vp < vl / 2; vp++) {
    x = VR[src][vl / 2 + vp];
    VR[dest][vp] = x;
}
for (vp = vl / 2; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

The top half of vr_{src} is extracted into the front of vr_{dest} . vr_{dest} must be different than vr_{src} . Trailing entries of vr_{dest} are not touched. The vector length specifies the length of vr_{src} . This instruction is not masked. The vector length must be a power of two, and must be at least two.

Exceptions


vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

This instruction is a special-case of the `vext.vv` instruction. This instruction may be used to implement reductions efficiently.

VHALF

Vector Half



Assembly

```
vhalf vrdest, vrsrc
```

Operation

```
if (vl > mvl || vl == 1 || !is_pow2 (vl)) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
for (vp = 0; vp < vl / 2; vp++) {
    x = VR[src][vl / 2 + vp];
    VR[dest][vp] = x;
}
for (vp = vl / 2; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
VC[vc_vl] = vl / 2;
```

Description


The top half of vr_{src} is extracted into the front of vr_{dest} . vr_{dest} must be different than vr_{src} . Trailing entries of vr_{dest} are not touched. The vector length specifies the length of vr_{src} . This instruction is not masked. The vector length must be a power of two, and must be at least two. After execution, the vector length register is divided by two.

Vector Half**VHALF(cont.)****Exceptions**

vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

This instruction is identical to the `vexthalf` instruction, except that this instruction also halves the vector length register. This instruction may be used to implement reductions efficiently.

VHALFDN**Vector Half Down****Assembly**

```
vhalfdn vrdest, vrsrc
```

Operation

```
if (vl > mvl || vl == 1 || !is_pow2 (vl)) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
L = 1 << vindex;
if (L >= vl) {
    vl = 0; /* operation not defined */
}
for (vp = 0; vp < vl; vp += 2 * L) {
    for (i = 0; i < L; i++) {
        x = VR[src][vp + L + i];
        VR[dest][vp + i] = x;
        x = VR[dest][vp + L + i];
        VR[dest][vp + L + i] = x;
    }
}
for (vp = vl; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Vector Half Down

VHALFDN(cont.)

Description

This instruction does one half of one step of a butterfly permutation. Paired with `vhalfdn`, it can be used to perform an in-register butterfly permutation. `vrdest` must be different than `vrsrc`. `vcindex` specifies the log of both the length and the separation of the sub-vectors that are moved from `vrsrc` to `vrdest`. The sub-vectors are shifted down within each $2 * 2^{vc_{index}}$ block by $2^{vc_{index}}$ elements. This writes half of the elements within each block in `vrdest`; the remaining elements are not touched. This instruction is not masked. `vcv1` must be at least 2 and must be a power of 2. $2^{vc_{index}}$ must be less than `vcv1`.

Exceptions


<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vIVL</code>	Invalid Vector Length Exception

Notes

This instruction may be used to implement multiple reductions efficiently.

VHALFUP

Vector Half Up



Assembly

vhalfup **vr**_{dest}, **vr**_{src}

Operation

```

if (vl > mvl || vl == 1 || !is_pow2 (vl)) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
L = 1 << vindex;
if (L >= vl) {
    vl = 0; /* operation not defined */
}
for (vp = 0; vp < vl; vp += 2 * L) {
    for (i = 0; i < L; i++) {
        x = VR[src][vp + i];
        VR[dest][vp + L + i] = x;
        x = VR[dest][vp + i];
        VR[dest][vp + i] = x;
    }
}
for (vp = vl; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}

```

Vector Half Up

VHALFUP(cont.)

Description

This instruction does one half of one step of a butterfly permutation. Paired with `vhalfdn`, it can be used to perform an in-register butterfly permutation. `vrdest` must be different than `vrsrc`. `vcvindex` specifies the log of both the length and the separation of the sub-vectors that are moved from `vrsrc` to `vrdest`. The sub-vectors are shifted up within each $2 * 2^{vc_{vindex}}$ block by $2^{vc_{vindex}}$ elements. This writes half of the elements within each block in `vrdest`; the remaining elements are not touched. This instruction is not masked. `vcv1` must be at least 2 and must be a power of 2. $2^{vc_{vindex}}$ must be less than `vcv1`.

Exceptions


<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vIVL</code>	Invalid Vector Length Exception

Notes

If `vcvindex` is `vcv1/2`, this instruction is identical to the `vexthalf` instruction. This instruction may be used to implement multiple reductions efficiently.

VFAND

Vector Flag And



Assembly

```

vfand {
    .VV vf_dest, vf_src1, vf_src2
    .SV vf_dest, vs_src1, vf_src2
}

```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        if (vv) {
            a = VF[src1][vp];
            b = VF[src2][vp];
        } else {
            a = (VS[src1] != 0);
            b = VF[src2][vp];
        }
        c = a && b;
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}

```

Description

Each VP writes into `vf_dest` the logical *and* of `vs_src1/vf_src1` and `vf_src2`. This instruction is not masked.

Exceptions


`vIVL` Invalid Vector Length Exception

Notes

A scalar source used as a boolean operand is false when zero and true when non-zero.

Vector Flag Clear

VFCLR



Assembly

```
vfclr vf_dest
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        c = false;
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}
```

Description

Each VP writes zero into `vf_dest`. This instruction is not masked.

Exceptions


`vIVL` Invalid Vector Length Exception

Notes

None.

VFCLR8

Vector Flag Clear 8



Assembly

```
vfclr8 vfdest
```

Operation

```
assert ((dest % 8) == 0);
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    for (i = 0; i < 8; i++) {
        VF[dest + i][vp] = false;
    }
}
for (vp = vl; vp < mvl; vp++) {
    for (i = 0; i < 8; i++) {
        c = VF[dest + i][vp];
        VF[dest + i][vp] = c;
    }
}
```

Description

Each VP writes zero into $\text{vf}_{\text{dest}} \cdots \text{vf}_{\text{dest}+7}$. This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.

Exceptions


vIVL	Invalid Vector Length Exception
vIUI	Illegal Use of Instruction Exception

Notes

None.

Vector Flag Find First One

VFFF1



Assembly

```
vfffl vsdest, vfsrc
```

Operation

```
count = 0;
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        break;
    }
}
VS[dest] = vp;
```

Description

The location of the first set bit of **vf**_{src} is placed in **vs**_{dest}. This instruction is not masked. If there is no set bit in **vf**_{src}, then the vector length is placed in **vs**_{dest}.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VFFL1

Vector Flag Find Last One



Assembly

```
vffl1 vsdest, vfsrc
```

Operation

```
last = vl;
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        last = vp;
    }
}
VS[dest] = last;
```

Description

The location of the last set bit of **vf**_{src} is placed in **vs**_{dest}. This instruction is not masked. If there is no set bit in **vf**_{src}, then the vector length is placed in **vs**_{dest}.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Scalar-Vector Insert

VFINS



Assembly

```
vfins vfdest, vssrc
```

Operation

```
vindex &= mvl - 1; /* force vindex into valid range */
for (vp = 0; vp < vindex && vp < mvl; vp++) {
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
b = (VS[src] != 0);
VF[dest][vindex] = b;
for (vp = vindex + 1; vp < mvl; vp++) {
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
```

Description

The boolean value of vs_{src} is written into vf_{dest} at position vc_{vindex} . The lower $vc_{\log mvl}$ bits of vc_{vindex} are used. This instruction is not masked and does not use vector length.

Exceptions


None.

Notes

A scalar source used as a boolean operand is false when zero and true when non-zero.

VFLD

Vector Flag Load



Assembly

```
vfld vf_dest[, vbase[, vinc]]
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
/* load 16-bit chunks; not 8-bit chunks! */
for (vp = 0; vp < vl; vp += 16) {
    va = base + vp / 8;
    if ((va & 0x1) != 0) { /* Address must be 16-bit aligned */
        fault[vp] = vAdEL;
    } else {
        fault[vp] = translate (READ, 2, va, &pa[vp]);
    }
    if (!speculative && fault[vp] != Exc_none) {
        Raise fault[vp];
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (speculative) {
        a = VF[vfe_L][vp];
        VF[vfe_L][vp] = a;
    }
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
for (vp = 0; vp < vl; vp += 16) {
    if (fault[vp] != Exc_none) {
        assert (speculative);
        for (i = 0; i < 16; i++) {
            VF[vfe_L][vp + i] = true;
        }
    } else {
        x = unsigned_load (pa[vp], 2);
        for (i = 0; i < 16 && vp + i < vl; i++) {
            loc = i < 8 ? i + 8 : i - 8; /* bytes are little-endian */
            b = (x >> loc) & 0x1;
            VF[dest][vp + i] = b;
        }
    }
}
VC[vc_base] = base + inc;

```

Description

The VPs perform a contiguous vector flag load into $\mathbf{vf}_{\text{dest}}$. $\mathbf{vf}_{\text{dest}}$ should be different than the speculative load fault flag register if the instruction is speculative. The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be 16-bit aligned. The bytes are loaded in little-endian order. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. This instruction is not masked.

VFLD(cont.)**Vector Flag Load****Exceptions**


vIVL	Invalid Vector Length Exception
vIUI	Illegal Use of Instruction Exception
vAdEL	Vector Address Error (Load)
vMod	Vector TLB Modification Exception
vTLBL	Vector TLB Exception (Load)
vDBEA	Vector Bus Error Exception (Address)
vDBED	Vector Bus Error Exception (Data)
vWatch	Watch Address Exception

Notes

Exactly VL bits are loaded into the destination register. All addresses are checked and translated before any values are loaded. A load fault causes an exception only if the instruction is non-speculative. A speculative faulting load does not load any value.

Vector Floating-Point Floor

VFLOOR



Assembly

$$\text{vfloor} \left\{ \begin{smallmatrix} \cdot \text{w} \\ \cdot \text{l} \end{smallmatrix} \right\} \left\{ \begin{smallmatrix} \cdot \text{s} \\ \cdot \text{d} \end{smallmatrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VFLOOR(cont.)**Vector Floating-Point Floor****Operation (Single to Word)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = floor_s_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Floor

VFLOOR(cont.)

Operation (Double to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = floor_d_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VFLOOR(cont.)**Vector Floating-Point Floor****Operation (Single to Long)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = floor_s_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Floor

VFLOOR(cont.)

Operation (Double to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = floor_d_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VFLOOR(cont.)**Vector Floating-Point Floor****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of converting \mathbf{vr}_{src} from a floating-point format to a signed integer format, using the *floor* rounding mode.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Flag Move From 8

VFMF8



Assembly

```
vfmf8 vrdest, vfsrc
```

Operation

```
assert ((src % 8) == 0);
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    x = 0;
    for (i = 0; i < 8; i++) {
        if (VF[src+i][vp]) {
            x |= 1 << i;
        }
    }
    VR[dest][vp] = x;
}
for (vp = vl; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

Registers $\text{vf}_{\text{src}} \cdots \text{vf}_{\text{src}+7}$ are copied into vr_{dest} . This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.

Exceptions


vIVL	Invalid Vector Length Exception
vIUI	Illegal Use of Instruction Exception

Notes

The mapping of flag registers into vr_{dest} is not defined, but must be consistent with the `vfmt8` instruction. The result of this instruction is only meaningful to the `vfmt8` instruction.

VFMT8

Vector Flag Move To 8



Assembly

```
vfmt8 vfdest, vrsrc
```

Operation

```
assert ((dest % 8) == 0);
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    x = VR[src][vp];
    for (i = 0; i < 8; i++) {
        b = (x >> i) & 0x1;
        VF[dest + i][vp] = b;
    }
}
for (vp = vl; vp < mvl; vp++) {
    for (i = 0; i < 8; i++) {
        b = VF[dest + i][vp];
        VF[dest + i][vp] = b;
    }
}
```

Description

Registers **vf**_{dest} . . . **vf**_{dest+7} are set with the contents of **vr**_{src}. This instruction is not masked, but is subject to vector length. The flag register specifier must be a multiple of 8.

Exceptions


vIVL	Invalid Vector Length Exception
vIUI	Illegal Use of Instruction Exception

Notes

The mapping of flag registers in **vr**_{src} is not defined, but is consistent with the **vfmf8** instruction. Register **vr**_{src} must contain the result of a previous **vfmf8** instruction.

Vector Flag Nor

VFNOR



Assembly

$$\text{vfnor} \begin{cases} .VV & \text{vf}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .SV & \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        if (vv) {
            a = VF[src1][vp];
            b = VF[src2][vp];
        } else {
            a = (VS[src1] != 0);
            b = VF[src2][vp];
        }
        c = !(a || b);
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}

```

Description

Each VP writes into vf_{dest} the logical *nor* of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

A scalar source used as a boolean operand is false when zero and true when non-zero.

VFOR

Vector Flag Or



Assembly

$$\text{vfor} \begin{cases} .VV & \text{vf}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .SV & \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        if (vv) {
            a = VF[src1][vp];
            b = VF[src2][vp];
        } else {
            a = (VS[src1] != 0);
            b = VF[src2][vp];
        }
        c = a || b;
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}

```

Description

Each VP writes into vf_{dest} the logical *or* of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

A scalar source used as a boolean operand is false when zero and true when non-zero.

Vector Flag Or 8

VFOR8



Assembly

```
vfor8 vfdest, vfsrc
```

Operation

```
assert ((src % 8) == 0);
assert ((dest % 8) == 0);
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    for (i = 0; i < 8; i++) {
        a = VF[dest + i][vp];
        b = VF[src + i][vp];
        VF[dest + i][vp] = a || b;
    }
}
for (vp = vl; vp < mvl; vp++) {
    for (i = 0; i < 8; i++) {
        c = VF[dest + i][vp];
        VF[dest + i][vp] = c;
    }
}
```

Description

Each VP performs an OR-write of $\mathbf{vf}_{\text{src}} \cdots \mathbf{vf}_{\text{src}+7}$ into $\mathbf{vf}_{\text{dest}} \cdots \mathbf{vf}_{\text{dest}+7}$. This instruction is not masked, but is subject to vector length. Each flag register specifier must be a multiple of 8.

Exceptions


vIVL	Invalid Vector Length Exception
vIUI	Illegal Use of Instruction Exception

Notes

None.

VFPOP

Vector Flag Population Count



Assembly

```
vfpop  $\mathbf{vs}_{\text{dest}}$ ,  $\mathbf{vf}_{\text{src}}$ 
```

Operation

```
count = 0;
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        count++;
    }
}
VS[dest] = count;
```

Description

The population count of \mathbf{vf}_{src} is placed in $\mathbf{vs}_{\text{dest}}$. This instruction is not masked.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Flush Vector Unit

VFLUSH



Assembly

vflush

Operation

```
fprintf (stderr, "vflush not yet implementedn");
```

Description

The vector unit must be frozen. All in-flight instructions are flushed. The register files are not touched.

Exceptions


None.

Notes

None.

VFSET

Vector Flag Set



Assembly

vfset **vf**_{dest}

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        c = true;
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}
```

Description

Each VP writes one into **vf**_{dest}. This instruction is not masked, but is subject to vector length.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Vector Flag Set Before First One

VFSETBF



Assembly

```
vfsetbf vfdest, vfsrc
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        break;
    }
    VF[dest][vp] = true;
}
for (; vp < vl; vp++) {
    VF[dest][vp] = false;
}
for (vp = vl; vp < mvl; vp++) {
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
```

Description

Register **vf_{dest}** is filled with ones up to and not including the first set bit in **vf_{src}**. Remaining positions in **vf_{dest}** are cleared. If **vf_{src}** contains no set bits, **vf_{dest}** is set to all ones. This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VFSETIF

Vector Flag Set Including First One



Assembly

```
vfsetif vfdest, vfsrc
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    b = VF[src][vp];
    VF[dest][vp] = true;
    if (b) {
        break;
    }
}
for (++vp; vp < vl; vp++) {
    VF[dest][vp] = false;
}
for (vp = vl; vp < mvl; vp++) {
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
```

Description

Register **vf**_{dest} is filled with ones up to and including the first set bit in **vf**_{src}. Remaining positions in **vf**_{dest} are cleared. If **vf**_{src} contains no set bits, **vf**_{dest} is set to all ones. This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Vector Flag Set Only First One

VFSETOF



Assembly

```
vfsetof vfdest, vfsrc
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        break;
    }
    VF[dest][vp] = false;
}
if (vp < vl) {
    VF[dest][vp] = true;
}
for (++vp; vp < vl; vp++) {
    VF[dest][vp] = false;
}
for (vp = vl; vp < mvl; vp++) {
    b = VF[dest][vp];
    VF[dest][vp] = b;
}
```

Description

Register **vf**_{dest} is filled with zeros except for the position of the first set bit in **vf**_{src}. If **vf**_{src} contains no set bits, **vf**_{dest} is set to all zeros. This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VFST

Vector Flag Store



Assembly

```
vfst vfsrc[, vbase[, vinc]]
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < vl; vp += 16) {
    va = base + vp / 8;
    if ((va & 0x1) != 0) { /* Address must be 16-bit aligned */
        fault[vp] = vAdEL;
    } else {
        fault[vp] = translate (WRITE, 2, va, &pa[vp]);
    }
    if (fault[vp] != Exc_none) {
        Raise fault[vp];
    }
}
for (vp = 0; vp < vl; vp += 16) {
    addr = pa[vp];
    x = 0;
    for (i = 0; i < 16; i++) {
        loc = i < 8 ? i + 8 : i - 8; /* bytes are little-endian */
        b = vp + i > vl ? false : VF[src][vp + i];
        x |= b << loc;
    }
    store (addr, 2, x);
}
VC[vc_base] = base + inc;

```

Description

The VPs perform a contiguous vector flag store of `vfsrc`. The base address is given by `vbase` (default is `vbase0`), and must be 16-bit aligned. A multiple of 16 bits is written, padding with zeros as necessary. The bytes are stored in little-endian order. The signed increment in `vinc` (default is `vinc0`) is added to `vbase` as a side-effect. This instruction is not masked.

Vector Flag Store**VFST(cont.)****Exceptions**


vIVL	Invalid Vector Length Exception
vAdES	Vector Address Error (Store)
vMod	Vector TLB Modification Exception
vTLBS	Vector TLB Exception (Store)
vDBEA	Vector Bus Error Exception (Address)
vDBED	Vector Bus Error Exception (Data)
vWatch	Watch Address Exception

Notes

All addresses are checked and translated before any values are stored.

VFXOR

Vector Flag Xor



Assembly

$$\text{vfxor} \begin{cases} .VV & \text{vf}_{\text{dest}}, \text{vf}_{\text{src1}}, \text{vf}_{\text{src2}} \\ .SV & \text{vf}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vf}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        if (vv) {
            a = VF[src1][vp];
            b = VF[src2][vp];
        } else {
            a = (VS[src1] != 0);
            b = VF[src2][vp];
        }
        c = a ^ b;
    } else {
        c = VF[dest][vp];
    }
    VF[dest][vp] = c;
}

```

Description

Each VP writes into vf_{dest} the logical *xor* of $\text{vs}_{\text{src1}}/\text{vf}_{\text{src1}}$ and vf_{src2} . This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

A scalar source used as a boolean operand is false when zero and true when non-zero.

Scalar-Vector Insert

VINS.SV



Assembly

```
vins.sv vrdest, vssrc
```

Operation

```
vindex &= mvl - 1; /* force vindex into valid range */
for (vp = 0; vp < vindex; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
x = VS[src];
VR[dest][vindex] = x;
for (vp = vindex + 1; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

The contents of vs_{src} are written into vr_{dest} at position vc_{vindex} . The lower $vc_{\log mvl}$ bits of vc_{vindex} are used. This instruction is not masked and does not use vector length.

Exceptions


None.

Notes

The scalar source is truncated to the virtual processor width.

VINS.VV

Vector-Vector Insert



Assembly

```
vins.vv vr_dest, vr_src2
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
if (src == dest) {
    Raise vIUI;
}
vindex &= mvl - 1; /* force vindex into valid range */
for (vp = 0; vp < vindex; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
for (vp = vindex; vp < vindex + vl && vp < mvl; vp++) {
    x = VR[src][vp - vindex];
    VR[dest][vp] = x;
}
for (vp = vindex + vl; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}

```

Description

The leading portion of \mathbf{vr}_{src} is inserted into $\mathbf{vr}_{\text{dest}}$. $\mathbf{vr}_{\text{dest}}$ must be different than \mathbf{vr}_{src} . Leading and trailing entries of $\mathbf{vr}_{\text{dest}}$ are not touched. The lower $\mathbf{vc}_{\log mvl}$ bits of vector control register $\mathbf{vc}_{\text{vindex}}$ specifies the starting position in $\mathbf{vr}_{\text{dest}}$. The vector length specifies the number of elements to transfer. This instruction is not masked.

Exceptions


vIUI	Illegal Use of Instruction Exception
vIVL	Invalid Vector Length Exception

Notes

If $\text{vindex} + \text{vl} \geq \text{mvl}$, then the only elements copied are those whose destination vp number is legal.
Vector-vector insert is the opposite of vector-vector extract.

VIOTA

Vector Iota



Assembly

```
viota vrdest, vfsrc
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
count = 0;
for (vp = 0; vp < vl; vp++) {
    if (VF[src][vp]) {
        VR[dest][count++] = vp;
    }
}
for (vp = count; vp < mvl; vp++) {
    x = VR[dest][vp];
    VR[dest][vp] = x;
}
```

Description

The list of VPs with bits set in \mathbf{vf}_{src} is placed in $\mathbf{vr}_{\text{dest}}$. This is the compressed index vector of \mathbf{vf}_{src} . This instruction is not masked.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Unit Stride Signed Vector Load

VLD



Assembly

$$\text{vld} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}[, \text{vbase}[, \text{vinc}]]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = signed_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}
VC[vc_base] = base + inc;

```

Unit Stride Signed Vector Load

VLD(cont.)

Description

The VPs perform a contiguous vector load into \mathbf{vr}_{dest} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdEL}	Vector Address Error (Load)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBL}	Vector TLB Exception (Load)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

All addresses are checked and translated before any values are loaded. A load fault causes an exception only if the instruction is non-speculative. A speculative faulting load does not load any value. The only difference between \mathbf{vld} and $\mathbf{vld.u}$ is that \mathbf{vld} sign-extends loaded data, and $\mathbf{vld.u}$ zero-extends loaded data.

VLD.U

Unit Stride Unsigned Vector Load



Assembly

$$\text{vld.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, [\text{vbase}, \text{vinc}]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = unsigned_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}
VC[vc_base] = base + inc;

```

Unit Stride Unsigned Vector Load**VLD.U(cont.)****Description**

Operation is identical to `vld`, except that loaded values are zero-extended to the virtual processor width instead of sign-extended.

Exceptions


<code>vIVL</code>	Invalid Vector Length Exception
<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vAdEL</code>	Vector Address Error (Load)
<code>vMod</code>	Vector TLB Modification Exception
<code>vTLBL</code>	Vector TLB Exception (Load)
<code>vDBEA</code>	Vector Bus Error Exception (Address)
<code>vDBED</code>	Vector Bus Error Exception (Data)
<code>vWatch</code>	Watch Address Exception

Notes

`vld.u.l` is not provided, as it would be functionally identical to `vld.l`.

VLDS

Variable Stride Signed Vector Load



Assembly

$$\text{vlds} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}[, \text{vbase}[, \text{vstride}[, \text{vinc}]]]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * stride * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = signed_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}
VC[vc_base] = base + inc;

```

VLDS(cont.)

Variable Stride Signed Vector Load

Description

The VPs perform a strided vector load into \mathbf{vr}_{dest} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The *signed* stride is given by $\mathbf{vstride}$ (default is $\mathbf{vstride}_0$). The stride is in terms of elements, not in terms of bytes. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdEL}	Vector Address Error (Load)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBL}	Vector TLB Exception (Load)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

Stride may be negative or zero. All addresses are checked and translated before any values are loaded. A load fault causes an exception only if the instruction is non-speculative. A speculative faulting load does not load any value. The only difference between \mathbf{vlds} and $\mathbf{vlds.u}$ is that \mathbf{vlds} sign-extends loaded data, and $\mathbf{vlds.u}$ zero-extends loaded data.

Variable Stride Unsigned Vector Load

VLDS.U



Assembly

$$\text{vlds.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}[, \text{vbase}[, \text{vstride}[, \text{vinc}]]]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * stride * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = unsigned_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}
VC[vc_base] = base + inc;

```

VLDS.U(cont.)**Variable Stride Unsigned Vector Load****Description**

Operation is identical to `vlds`, except that loaded values are zero-extended to the virtual processor width instead of sign-extended.

Exceptions


<code>vIVL</code>	Invalid Vector Length Exception
<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vAdEL</code>	Vector Address Error (Load)
<code>vMod</code>	Vector TLB Modification Exception
<code>vTLBL</code>	Vector TLB Exception (Load)
<code>vDBEA</code>	Vector Bus Error Exception (Address)
<code>vDBED</code>	Vector Bus Error Exception (Data)
<code>vWatch</code>	Watch Address Exception

Notes

`vlds.u.l` is not provided, as it would be functionally identical to `vlds.l`.

Indexed Signed Vector Load

VLDX



Assembly

$$\text{vldx} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{offsets}}, \text{vbase}$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + VR[offsets][vp];
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = signed_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}

```

VLDX(cont.)

Indexed Signed Vector Load

Description

The VPs perform an indexed vector load into $\mathbf{vr}_{\text{dest}}$. The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0). The *signed* offsets are given by $\mathbf{vr}_{\text{offsets}}$. The offsets are in units of bytes, not in units of elements. The effective addresses must be aligned to the width of the data in memory. The width of each element in memory is given by the opcode. The loaded value is sign-extended to the virtual processor width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdEL}	Vector Address Error (Load)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBL}	Vector TLB Exception (Load)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

All addresses are checked and translated before any values are loaded. A load fault causes an exception only if the instruction is non-speculative. A speculative faulting load does not load any value. The only difference between \mathbf{vldx} and $\mathbf{vldx.u}$ is that \mathbf{vldx} sign-extends loaded data, and $\mathbf{vldx.u}$ zero-extends loaded data.

Indexed Unsigned Vector Load

VLDX.U



Assembly

$$\text{vldx.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{offsets}}[, \text{vbase}]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + VR[offsets][vp];
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (READ, nbytes, va, &pa[vp]);
        }
        if (!speculative && fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (fault[vp] != Exc_none) {
            if (speculative) b = true;
            x = VR[dest][vp];
        } else {
            if (speculative) b = VF[vfe_L][vp];
            x = unsigned_load (pa[vp], nbytes);
        }
    } else {
        if (speculative) b = VF[vfe_L][vp];
        x = VR[dest][vp];
    }
    if (speculative) VF[vfe_L][vp] = b;
    VR[dest][vp] = x;
}

```

VLDX.U(cont.)**Indexed Unsigned Vector Load****Description**

Operation is identical to `vldx`, except that loaded values are zero-extended to the virtual processor width instead of sign-extended.

Exceptions


<code>vIVL</code>	Invalid Vector Length Exception
<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vAdEL</code>	Vector Address Error (Load)
<code>vMod</code>	Vector TLB Modification Exception
<code>vTLBL</code>	Vector TLB Exception (Load)
<code>vDBEA</code>	Vector Bus Error Exception (Address)
<code>vDBED</code>	Vector Bus Error Exception (Data)
<code>vWatch</code>	Watch Address Exception

Notes

`vldx.u.l` is not provided, as it would be functionally identical to `vldx.l`.

Vector Floating-Point Multiply Add

VMADD.fmt



Assembly

$$\text{vmadd} \begin{Bmatrix} .s \\ .d \end{Bmatrix} \begin{Bmatrix} .vv[.1] \\ .sv[.1] \end{Bmatrix} \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}}$$

VMADD.fmt(cont.)**Vector Floating-Point Multiply Add****Operation (Single Precision)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = mul_add_s (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Multiply Add

VMADD.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = mul_add_d (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VMADD.fmt(cont.)**Vector Floating-Point Multiply Add****Description**

Each unmasked VP adds the floating-point product of $\mathbf{vr}_{\text{src1}}/\mathbf{vs}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$ into $\mathbf{vr}_{\text{dest}}$.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Signed Vector Integer Maximum

VMAX



Assembly

$$\text{vmax} \begin{cases} \text{.VV}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.SV}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x > y ? x : y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the greater of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as signed integers.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VMAX.U

Unsigned Vector Integer Maximum



Assembly

$$\text{vmax.u} \begin{cases} \text{.vv}[\cdot,1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot,1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x > y ? x : y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the greater of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as unsigned integers.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Commit Speculative Arithmetic

VMCOMMIT



Assembly

```
vmcommit[.exact][.1]
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
L = false;
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        L = L || VF[vfe_L][vp];
    }
}
if (L) {
    Raise vSLE;
}
```

Description

Any pending vector memory exceptions generated by speculative operations are raised. This instruction operates under mask and vector length. The exceptions are precise if either `.exact` is specified, or if `vcvmode` specifies exact exceptions.

Exceptions


None.

Notes

None.

VMCTS

Vector Move Control To Scalar



Assembly

```
vmcts  $vs_{dest}$ ,  $vc_{src}$ 
```

Operation

```
if (!reg_user_readable (src)) { /* assume user mode for now */
    Raise vIUI;
}
x = pdinst->get_src_vc (src);
VS[dest] = x;
```

Description

Register vc_{src} is copied to vs_{dest} .

Exceptions


$vIUI$	Illegal Use of Instruction Exception
--------	--------------------------------------

Notes

None.

Vector Merge

VMERGE



Assembly

$$\text{vmerge} \begin{cases} \text{.VV[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.SV[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.VS[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        if (VF[mask][vp]) {
            z = y;
        } else {
            z = x;
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each VP copies into vr_{dest} either $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ if the mask is 0, or $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$ if the mask is 1. At least one source is a vector. Scalar sources are truncated to the virtual processor width.

VMERGE(cont.)**Vector Merge****Exceptions**


$vIVL$ Invalid Vector Length Exception

Notes

This instruction can be used to efficiently replicate a scalar value into a vector register.

Signed Vector Integer Minimum

VMIN



Assembly

$$\text{vmin} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x < y ? x : y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the lesser of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as signed integers.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VMIN.U

Unsigned Vector Integer Minimum



Assembly

$$\text{vmin.u} \begin{cases} \text{.VV}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.SV}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x < y ? x : y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the lesser of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} , treated as unsigned integers.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Signed Vector Integer Modulus

VMOD



Assembly

$$\text{vmod} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        if (y == 0) {
            z = VR[dest][vp];
        } else {
            z = x % y;
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the signed integer modulus of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.

VMOD(cont.)**Signed Vector Integer Modulus****Exceptions**


\mathbf{vIVL} Invalid Vector Length Exception

Notes

The result is undefined if the divisor is zero.

Unsigned Vector Integer Modulus

VMODU



Assembly

$$\text{vmodu} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        if (y == 0) {
            z = VR[dest][vp];
        } else {
            z = x % y;
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the unsigned integer modulus of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector.

VMODU(cont.)

Unsigned Vector Integer Modulus

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

The result is undefined if the divisor is zero.

Vector Move Scalar To Control

VMSTC



Assembly

```
vmstc vCdest, vSsrc
```

Operation

```
if (dest != VcVl && dest != VcVpw && dest != VcShamt &&
    dest != VcIndex && dest != VcAt) {
    Raise vIUI;
}
x = VS[src];
if (dest == VcVpw) {
    assert (x <= 3);
    VC[vpw] = x;
    VC[mvl] = MVL[x];
    VC[logmvl] = LOGMVL[x];
} else {
    VC[dest] = x;
}
```

Description

Register `vssrc` is copied to `vcdest`. Writing `vcvpw` changes `vcmvl`, `vclogmvl`, and `vcvsafevl` as a side-effect. This instruction may write only `vcvl`, `vcvpw`, `vcvshamt`, `vcvindex`, and `vcvcat`.

Exceptions


vInt	Vector Integer Exception
vFPE	Vector Floating-Point Exception
vIUI	Illegal Use of Instruction Exception

Notes

None.

VMSUB.fmt

Vector Floating-Point Multiply Subtract



Assembly

$$\text{vmsub} \begin{Bmatrix} .s \\ .d \end{Bmatrix} \begin{Bmatrix} .vv[.1] \\ .sv[.1] \end{Bmatrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}}$$

Vector Floating-Point Multiply Subtract

VMSUB.fmt(cont.)

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = mul_sub_s (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VMSUB.fmt(cont.)

Vector Floating-Point Multiply Subtract

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = mul_sub_d (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Multiply Subtract**VMSUB.fmt(cont.)****Description**

Each unmasked VP subtracts $\mathbf{vr}_{\text{dest}}$ from the floating-point product of $\mathbf{vr}_{\text{src1}}/\mathbf{vs}_{\text{src1}}$.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

VMUL.fmt

Vector Floating-Point Multiply



Assembly

$$\text{vmul} \begin{cases} .s \\ .d \end{cases} \begin{cases} .vv[.1] \\ .sv[.1] \end{cases} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}}$$

Vector Floating-Point Multiply

VMUL.fmt(cont.)

Operation (Single Precision)

```

    if (vl > mvl) {
        Raise vIVL;
    }
    if (vpw < 2) {
        Raise vIUI;
    }
    raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
    for (vp = 0; vp < mvl; vp++) {
        I = U = O = Z = V = E = false;
        if (vp < vl && VF[mask][vp]) {
            if (vv) {
                x = VR[src1][vp];
                y = VR[src2][vp];
            } else {
                x = VS[src1];
                y = VR[src2][vp];
            }
            z = mul_s(x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
            if (!speculative) {
                write = true;
                if (I && I_enabled) { raise_I = true; write = false; }
                if (U && U_enabled) { raise_U = true; write = false; }
                if (O && O_enabled) { raise_O = true; write = false; }
                if (Z && Z_enabled) { raise_Z = true; write = false; }
                if (V && V_enabled) { raise_V = true; write = false; }
                if (E) { raise_E = true; write = false; }
                if (!write) {
                    z = VR[dest][vp];
                }
            }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }

```

VMUL.fmt(cont.)

Vector Floating-Point Multiply

Operation (Double Precision)

```

    if (vl > mvl) {
        Raise vIVL;
    }
    if (vpw < 3) {
        Raise vIUI;
    }
    raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
    for (vp = 0; vp < mvl; vp++) {
        I = U = O = Z = V = E = false;
        if (vp < vl && VF[mask][vp]) {
            if (vv) {
                x = VR[src1][vp];
                y = VR[src2][vp];
            } else {
                x = VS[src1];
                y = VR[src2][vp];
            }
            z = mul_d(x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
            if (!speculative) {
                write = true;
                if (I && I_enabled) { raise_I = true; write = false; }
                if (U && U_enabled) { raise_U = true; write = false; }
                if (O && O_enabled) { raise_O = true; write = false; }
                if (Z && Z_enabled) { raise_Z = true; write = false; }
                if (V && V_enabled) { raise_V = true; write = false; }
                if (E) { raise_E = true; write = false; }
                if (!write) {
                    z = VR[dest][vp];
                }
            }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }

```

Vector Floating-Point Multiply**VMUL.fmt(cont.)****Description**

Each unmasked VP places the floating-point product of $\mathbf{vr}_{src1}/\mathbf{vs}_{src1}$ and \mathbf{vr}_{src2} into \mathbf{vr}_{dest} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

VMULHI

Vector Integer Multiply High



Assembly

$$\text{vmulhi} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = signed_multiply_high (x, y);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the upper half of the full-VP-width signed integer product of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Vector Integer Multiply High

VMULHI.U



Assembly

```
vmulhi.u {
    .vv[.1] vr_dest, vr_src1, vr_src2
    .sv[.1] vr_dest, vs_src1, vr_src2
}
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = unsigned_multiply_high (x, y);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
```

Description

Each unmasked VP writes into `vr_dest` the upper half of the full-VP-width unsigned integer product of `vs_src1/vr_src1` and `vr_src2`.

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VMULLO

Vector Integer Multiply Low



Assembly

$$\text{vmullo} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x * y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the lower half of the full-VP-width integer product of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .

Exceptions


vIVL Invalid Vector Length Exception

Notes

There are not separate signed and unsigned multiply low instructions, because signed and unsigned multiply low are identical operations.

Vector Floating-Point Negate

VNEG.fmt



Assembly

$$\text{vneg} \left\{ \begin{matrix} \cdot \text{s} \\ \cdot \text{d} \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VNEG.fmt(cont.)

Vector Floating-Point Negate

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = neg_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Negate

VNEG.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = neg_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VNEG.fmt(cont.)**Vector Floating-Point Negate****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the floating-point square root of \mathbf{vr}_{src} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Floating-Point Negative Multiply Add

VNMADD.fmt



Assembly

$$\text{vnmadd} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} \left\{ \begin{matrix} .vv[.1] \\ .sv[.1] \end{matrix} \right\} \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}}$$

VNMADD.fmt(cont.)

Vector Floating-Point Negative Multiply Add

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = neg_mul_add_s (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Negative Multiply Add

VNMADD.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = neg_mul_add_d (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VNMADD.fmt(cont.)**Vector Floating-Point Negative Multiply Add****Description**

Each unmasked VP adds the floating-point product of $\mathbf{vr}_{\text{src1}}/\mathbf{vs}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$ into $\mathbf{vr}_{\text{dest}}$, and negates the result.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Floating-Point Negative Multiply Subtract

VNMSUB.fmt



Assembly

$$\text{vnmsub} \begin{Bmatrix} .s \\ .d \end{Bmatrix} \begin{Bmatrix} .vv[.1] \\ .sv[.1] \end{Bmatrix} \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}}$$

VNMSUB.fmt(cont.)

Vector Floating-Point Negative Multiply Subtract

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = neg_mul_sub_s (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Negative Multiply Subtract

VNMSUB.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = VR[src3][vp];
        z = neg_mul_sub_d (x, y, z, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VNMSUB.fmt(cont.) **Vector Floating-Point Negative Multiply Subtract****Description**

Each unmasked VP subtracts $\mathbf{vr}_{\text{dest}}$ from the floating-point product of $\mathbf{vr}_{\text{src1}}/\mathbf{vs}_{\text{src1}}$, and negates the result.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Nor

VNOR



Assembly

$$\text{vnor} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = ~(x | y);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each VP writes into vr_{dest} the bit-wise logical *nor* of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

VOR

Vector Or



Assembly

$$\text{vor} \begin{cases} \text{.vv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x | y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each VP writes into vr_{dest} the bit-wise logical *or* of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Vector Floating-Point Reciprocal

VRECIP.fmt



Assembly

$$\text{vrecip} \left\{ \begin{matrix} \cdot s \\ \cdot d \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VRECIP.fmt(cont.)

Vector Floating-Point Reciprocal

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = recip_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```


Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = recip_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VRECIP.fmt(cont.)**Vector Floating-Point Reciprocal****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the floating-point reciprocal of \mathbf{vr}_{src} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Floating-Point Round

VROUND



Assembly

$$\text{vround} \left\{ \begin{matrix} .w \\ .l \end{matrix} \right\} \left\{ \begin{matrix} .s \\ .d \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VROUND(cont.)**Vector Floating-Point Round****Operation (Single to Word)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = round_s_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Round

VROUND(cont.)

Operation (Double to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = round_d_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VROUND(cont.)**Vector Floating-Point Round****Operation (Single to Long)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = round_s_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Round

VROUND(cont.)

Operation (Double to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = round_d_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VROUND(cont.)**Vector Floating-Point Round****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of converting \mathbf{vr}_{src} from a floating-point format to a signed integer format, using the *round* rounding mode.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Floating-Point Reciprocal Square Root

VRSQRT.fmt



Assembly

$$\text{vrsqrt} \left\{ \begin{matrix} \cdot \text{s} \\ \cdot \text{d} \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VRSQRT.fmt(cont.)**Vector Floating-Point Reciprocal Square Root****Operation (Single Precision)**

```

    if (vl > mvl) {
        Raise vIVL;
    }
    if (vpw < 2) {
        Raise vIUI;
    }
    raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
    for (vp = 0; vp < mvl; vp++) {
        I = U = O = Z = V = E = false;
        if (vp < vl && VF[mask][vp]) {
            x = VR[src][vp];
            z = rsqrt_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
            if (!speculative) {
                write = true;
                if (I && I_enabled) { raise_I = true; write = false; }
                if (U && U_enabled) { raise_U = true; write = false; }
                if (O && O_enabled) { raise_O = true; write = false; }
                if (Z && Z_enabled) { raise_Z = true; write = false; }
                if (V && V_enabled) { raise_V = true; write = false; }
                if (E) { raise_E = true; write = false; }
                if (!write) {
                    z = VR[dest][vp];
                }
            }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }

```

Vector Floating-Point Reciprocal Square Root

VRSQRT.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = rsqrt_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VRSQRT.fmt(cont.)**Vector Floating-Point Reciprocal Square Root****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the floating-point reciprocal square root of \mathbf{vr}_{src} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Register Sync

VRSYNC



Assembly

$$\text{vrsync} \left\{ \begin{matrix} .sav \\ .vav \\ .vp \end{matrix} \right\} \left\{ \begin{matrix} \phi \\ .raw \\ .war \\ .waw \end{matrix} \right\} \text{vr}_{\text{dest}}$$

Operation

`/* operation is a nop in a simple implementation */`

Description

Enforces scalar-after-vector, vector-after-scalar, vector-after-vector, or intra-vp memory ordering. If none of raw/war/waw are specified, then all are in effect. The most recent memory reference to vr_{dest} preceeding the sync must appear to execute before all relevant memory references following the sync. Order is defined as program order.

Exceptions


None.

Notes

None.

VSADD

Signed Saturating Add



Assembly

```
vsadd { .vv[.1] vr_dest, vr_src1, vr_src2
       .sv[.1] vr_dest, vs_src1, vr_src2
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = signed_saturating_add (x, y, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}
```

Signed Saturating Add**VSADD(cont.)****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the signed integer sum of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. The sum saturates to the VP width instead of overflowing.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VSADD.U

Unsigned Saturating Add



Assembly

```
vsadd.u { .vv[.1] vr_dest, vr_src1, vr_src2
         .sv[.1] vr_dest, vs_src1, vr_src2
```

Operation

```
if (vl > mvl) {
    Raise vIVL;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = unsigned_saturating_add (x, y, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}
```

Unsigned Saturating Add**VSADD.U(cont.)****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the unsigned integer sum of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. The sum saturates to the VP width instead of overflowing.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VSAT

Signed Vector Saturate



Assembly

$$\text{vsat} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Operation

```

assert (sat_width == B || sat_width == H || sat_width == W);
if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2 && sat_width == W || vpw < 1 && sat_width == B) {
    Raise vIUI;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = signed_saturate (x, sat_width, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Signed Vector Saturate

VSAT(cont.)

Description

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of saturating \mathbf{vr}_{src} to a signed integer narrower than the VP width. The result is sign-extended to the VP width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

It is illegal to saturate to wider than the VP width.

VSAT.SU

Signed to Unsigned Vector Saturate



Assembly

$$\text{vsat.su} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Operation

```

assert (sat_width == B || sat_width == H || sat_width == W || sat_width == L);
if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3 && sat_width == L || vpw < 2 && sat_width == W || vpw < 1 && sat_width == B)
    Raise vIUI;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = signed_to_unsigned_saturate (x, sat_width, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Signed to Unsigned Vector Saturate

VSAT.SU(cont.)

Description

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of saturating \mathbf{vr}_{src} from a signed VP width value to an unsigned value that is as wide or narrower than the VP width. The result is zero-extended to the VP width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

It is illegal to saturate to wider than the VP width.

VSAT.U

Unsigned Vector Saturate



Assembly

$$\text{vsat.u} \left\{ \begin{matrix} .b \\ .h \\ .w \end{matrix} \right\} [.1] \text{ vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Operation

```

assert (sat_width == B || sat_width == H || sat_width == W);
if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2 && sat_width == W || vpw < 1 && sat_width == B) {
    Raise vIUI;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = unsigned_saturate (x, sat_width, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Unsigned Vector Saturate

VSAT.U(cont.)

Description

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of saturating \mathbf{vr}_{src} to an unsigned integer narrower than the VP width. The result is zero-extended to the VP width.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

It is illegal to saturate to wider than the VP width.

VSATVL

Saturate Vector Length



Assembly

vsatvl

Operation

VC[vc_vl] = vl > mvl ? mvl : vl;

Description

The vector length register is saturated to the maximum vector length.

Exceptions


None.

Notes

None.

Vector Shift Left Logical

VSL



Assembly

$$\text{vsl} \left\{ \begin{array}{l} \text{.vv}[\text{.1}] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] \text{ vr}_{\text{dest}}, \text{ vs}_{\text{src1}}, \text{ vr}_{\text{src2}} \\ \text{.vs}[\text{.1}] \text{ vr}_{\text{dest}}, \text{ vr}_{\text{src1}}, \text{ vs}_{\text{src2}} \end{array} \right.$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        y &= shamt_mask;
        z = x << y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ left-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is zero-filled.

VSLL(cont.)

Vector Shift Left Logical

Exceptions


vIVL Invalid Vector Length Exception

Notes

Only the lower b bits of the register from which the shift amount is taken are used (2^b equals the VP width in bits).

Signed Saturating Left Shift

VSLs



Assembly

```
vsls[.1] vr_dest, vr_src
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = signed_shift_left_and_sat (x, shamt, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Description

Each unmasked VP writes into `vr_dest` the signed saturating left shift of `vr_src`. The shift amount is taken from `vc_vshamt`.

VSLS(cont.)**Signed Saturating Left Shift****Exceptions**


v_{IVL} Invalid Vector Length Exception

Notes

None.

Unsigned Saturating Left Shift

VSL.S.U



Assembly

```
vsls.u[.1] vr_dest, vr_src
```

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = unsigned_shift_left_and_sat (x, shamt, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Description

Each unmasked VP writes into vr_{dest} the unsigned saturating left shift of vr_{src} . The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

VSL.S.U(cont.)**Unsigned Saturating Left Shift****Exceptions**


v_{IVL} Invalid Vector Length Exception

Notes

None.

Vector Floating-Point Square Root

VSQRT.fmt



Assembly

$$\text{vsqrt}\left\{\begin{smallmatrix} \cdot \text{s} \\ \cdot \text{d} \end{smallmatrix}\right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

VSQRT.fmt(cont.)

Vector Floating-Point Square Root

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = sqrt_s (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Square Root

VSQRT.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = sqrt_d (x, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VSQRT.fmt(cont.)**Vector Floating-Point Square Root****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the floating-point square root of \mathbf{vr}_{src} .

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Arithmetic Right Shift

VSRA



Assembly

$$\text{vsra} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        y &= shamt_mask;
        z = shift_right_arith (x, y);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ right-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is sign-extended.

VSRA(cont.)

Vector Arithmetic Right Shift

Exceptions


vIVL Invalid Vector Length Exception

Notes

Only the lower b bits of the register from which the shift amount is taken are used (2^b equals the VP width in bits).

Vector Shift Right Logical

VSRL



Assembly

$$\text{vsrl} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        y &= shamt_mask;
        z = x >> y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP writes into vr_{dest} the signed integer contents of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ right-shifted by the number of bits specified by $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where at least one source is a vector. The shift amount is the b -bit unsigned integer taken from the low-order end of $\text{vs}_{\text{src2}}/\text{vr}_{\text{src2}}$, where 2^b equals the VP width in bits. The result is zero-extended.

VSRL(cont.)

Vector Shift Right Logical

Exceptions


v_{IVL} Invalid Vector Length Exception

Notes

VSRLnotes

Signed Shift Right And Round

VSRR



Assembly

```
vsrr[.1] vr_dest, vr_src
```

Operation

```
assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = signed_shift_right_and_round (x, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
```

Description

Each unmasked VP writes into `vr_dest` the right arithmetic shift of `vr_src`. The result is rounded as per the fixed-point rounding mode. The shift amount is taken from `vc_vshamt`.

Exceptions


`vIVL` Invalid Vector Length Exception

Notes

None.

VSRR.U

Unsigned Shift Right And Round



Assembly

```
vsrr.u[.1] vr_dest, vr_src
```

Operation

```
assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = unsigned_shift_right_and_round (x, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
```

Description

Each unmasked VP writes into vr_{dest} the right logical shift of vr_{src} . The result is rounded as per the fixed-point rounding mode. The shift amount is taken from vc_{vshamt} .

Exceptions


$vIVL$ Invalid Vector Length Exception

Notes

None.

Signed Saturating Subtract

VSSUB



Assembly

$$\text{vssub} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = signed_saturating_sub (x, y, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VSSUB(cont.)**Signed Saturating Subtract****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the signed integer subtraction of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vs}_{\text{src2}}/\mathbf{vr}_{\text{src2}}$, where at least one source is a vector. The difference saturates to the VP width instead of overflowing.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Signed Saturating Subtract

VSSUB.U



Assembly

$$\text{vssub.u} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.vs[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vs}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = unsigned_saturating_sub (x, y, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VSSUB.U(cont.)**Signed Saturating Subtract****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the unsigned integer subtraction of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vs}_{\text{src2}}/\mathbf{vr}_{\text{src2}}$, where at least one source is a vector. The difference saturates to zero instead of underflowing.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Unit Stride Vector Store

VST



Assembly

$$\text{vst} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{src}}, \text{vbase}, [\text{vinc}]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (WRITE, nbytes, va, &pa[vp]);
        }
        if (fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        x = VR[src][vp];
        store (pa[vp], nbytes, x);
    }
}
VC[vc_base] = base + inc;

```

VST(cont.)

Unit Stride Vector Store

Description

The VPs perform a contiguous vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The register value is truncated from the VP width to the memory width. The VPs access memory in order.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdES}	Vector Address Error (Store)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBS}	Vector TLB Exception (Store)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

All addresses are checked and translated before any values are stored.

Variable Stride Vector Store

VSTS



Assembly

$$\text{vsts} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{src}}, \text{vbase}, \text{vstride}, \text{vinc}]]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + nbytes * stride * vp;
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (WRITE, nbytes, va, &pa[vp]);
        }
        if (fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        x = VR[src][vp];
        store (pa[vp], nbytes, x);
    }
}
VC[vc_base] = base + inc;

```

VSTS(cont.)

Variable Stride Vector Store

Description

The VPs perform a contiguous vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0), and must be aligned to the width of the data in memory. The *signed* stride is given by $\mathbf{vstride}$ (default is $\mathbf{vstride}_0$). The stride is in terms of elements, not in terms of bytes. The signed increment in \mathbf{vinc} (default is \mathbf{vinc}_0) is added to \mathbf{vbase} as a side-effect. The width of each element in memory is given by the opcode. The register value is truncated from the VP width to the memory width. The VPs access memory in order.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdES}	Vector Address Error (Store)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBS}	Vector TLB Exception (Store)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

All addresses are checked and translated before any values are stored.

Unordered Indexed Vector Store

VSTX



Assembly

$$\text{vstx} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{ vr}_{\text{src}}, \text{vr}_{\text{offsets}}[, \text{vbase}]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + VR[offsets][vp];
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (WRITE, nbytes, va, &pa[vp]);
        }
        if (fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        x = VR[src][vp];
        store (pa[vp], nbytes, x);
    }
}

```

VSTX(cont.)

Unordered Indexed Vector Store

Description

The VPs perform an indexed vector store of \mathbf{vr}_{src} . The base address is given by \mathbf{vbase} (default is \mathbf{vbase}_0). The *signed* offsets are given by $\mathbf{vr}_{offsets}$. The offsets are in units of bytes, not in units of elements. The effective addresses must be aligned to the width of the data in memory. The register value is truncated from the VP width to the memory width. The stores may be performed in any order.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception
\mathbf{vAdES}	Vector Address Error (Store)
\mathbf{vMod}	Vector TLB Modification Exception
\mathbf{vTLBS}	Vector TLB Exception (Store)
\mathbf{vDBEA}	Vector Bus Error Exception (Address)
\mathbf{vDBED}	Vector Bus Error Exception (Data)
\mathbf{vWatch}	Watch Address Exception

Notes

All addresses are checked and translated before any values are stored. An implementation is free to perform the stores in any order. The implementation shown performs them in order.

Ordered Indexed Vector Store

VSTXO



Assembly

$$\text{vstxo} \left\{ \begin{matrix} .b \\ .h \\ .w \\ .l \end{matrix} \right\} [.1] \text{vr}_{\text{src}}, \text{vr}_{\text{offsets}}, [\text{vbase}]$$

Operation

```

assert (nbytes == 1 || nbytes == 2 || nbytes == 4 || nbytes == 8);
if (vl > mvl) {
    Raise vIVL;
}
if (nbytes > (1 << vpw)) {
    Raise vIUI;
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        va = base + VR[offsets][vp];
        if (va & (nbytes - 1)) {
            fault[vp] = vAdEL;
        } else {
            fault[vp] = translate (WRITE, nbytes, va, &pa[vp]);
        }
        if (fault[vp] != Exc_none) {
            Raise fault[vp];
        }
    }
}
for (vp = 0; vp < vl; vp++) {
    if (VF[mask][vp]) {
        x = VR[src][vp];
        store (pa[vp], nbytes, x);
    }
}

```

VSTXO(cont.)**Ordered Indexed Vector Store****Description**

Operation is identical to `vstx`, except that the VPs access memory in order.


Exceptions

<code>vIVL</code>	Invalid Vector Length Exception
<code>vIUI</code>	Illegal Use of Instruction Exception
<code>vAdES</code>	Vector Address Error (Store)
<code>vMod</code>	Vector TLB Modification Exception
<code>vTLBS</code>	Vector TLB Exception (Store)
<code>vDBEA</code>	Vector Bus Error Exception (Address)
<code>vDBED</code>	Vector Bus Error Exception (Data)
<code>vWatch</code>	Watch Address Exception

Notes

Signed Vector Integer Subtract

VSUB



Assembly

$$\text{vsub} \begin{cases} .vv[.1] & \text{VR}_{\text{dest}}, \text{VR}_{\text{src1}}, \text{VR}_{\text{src2}} \\ .sv[.1] & \text{VR}_{\text{dest}}, \text{VS}_{\text{src1}}, \text{VR}_{\text{src2}} \\ .vs[.1] & \text{VR}_{\text{dest}}, \text{VR}_{\text{src1}}, \text{VS}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
raise_F = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = x - y;
        if (sub_overflow (x, y, z)) {
            f = true;
            if (!speculative && F_enabled) {
                raise_F = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_F][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_F][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_F][vp] = f;
}
if (raise_F) {
    Raise vAri;
}

```

VSUB(cont.)**Signed Vector Integer Subtract****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the signed integer subtraction of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vs}_{\text{src2}}/\mathbf{vr}_{\text{src2}}$, where at least one source is a vector.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Unsigned Vector Integer Subtract

VSUB.U



Assembly

$$\text{vsub.u} \begin{cases} \text{.vv}[\text{.1}] & \text{VR}_{\text{dest}}, \text{VR}_{\text{src1}}, \text{VR}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{VR}_{\text{dest}}, \text{VS}_{\text{src1}}, \text{VR}_{\text{src2}} \\ \text{.vs}[\text{.1}] & \text{VR}_{\text{dest}}, \text{VR}_{\text{src1}}, \text{VS}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = x - y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

VSUB.U(cont.)**Unsigned Vector Integer Subtract****Description**

Each unmasked VP writes into $\mathbf{vr}_{\text{dest}}$ the unsigned integer subtraction of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vs}_{\text{src2}}/\mathbf{vr}_{\text{src2}}$, where at least one source is a vector.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

The only difference between \mathbf{vsub} and $\mathbf{vsub.u}$ is that \mathbf{vsub} can overflow, while $\mathbf{vsub.u}$ cannot.

Vector Floating-Point Subtract

VSUB.fmt



Assembly

$$\text{vsub} \begin{Bmatrix} .s \\ .d \end{Bmatrix} \begin{cases} .vv[.1] & \text{vr}_{\text{dest}} , & \text{vr}_{\text{src1}} , & \text{vr}_{\text{src2}} \\ .sv[.1] & \text{vr}_{\text{dest}} , & \text{vs}_{\text{src1}} , & \text{vr}_{\text{src2}} \\ .vs[.1] & \text{vr}_{\text{dest}} , & \text{vr}_{\text{src1}} , & \text{vs}_{\text{src2}} \end{cases}$$

VSUB.fmt(cont.)

Vector Floating-Point Subtract

Operation (Single Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = sub_s(x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Subtract

VSUB.fmt(cont.)

Operation (Double Precision)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else if (sv) {
            x = VS[src1];
            y = VR[src2][vp];
        } else {
            x = VR[src1][vp];
            y = VS[src2];
        }
        z = sub_d (x, y, FS, RM, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
        } else {
            z = VR[dest][vp];
        }
        VR[dest][vp] = z;
        VF[vfe_I][vp] = I || VF[vfe_I][vp];
        VF[vfe_U][vp] = U || VF[vfe_U][vp];
        VF[vfe_O][vp] = O || VF[vfe_O][vp];
        VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
        VF[vfe_V][vp] = V || VF[vfe_V][vp];
        VF[vfe_E][vp] = E || VF[vfe_E][vp];
    }
    if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
        Raise vAri;
    }
}

```

VSUB.fmt(cont.)**Vector Floating-Point Subtract****Description**

Each unmasked VP places the floating-point difference of $\mathbf{vr}_{src1}/\mathbf{vs}_{src1}$ and $\mathbf{vr}_{src2}/\mathbf{vf}_{src2}$ into \mathbf{vr}_{dest} , where at least one source is a vector.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

Vector Sync

VSYNC



Assembly

vsync $\left\{ \begin{matrix} .sav \\ .vas \\ .vav \\ .vp \end{matrix} \right\} \left\{ \begin{matrix} \phi \\ .raw \\ .war \\ .waw \end{matrix} \right\}$

Operation

```
/* operation is a nop in a simple implementation */
```

Description

Enforces scalar-after-vector, vector-after-scalar, vector-after-vector, or intra-vp memory ordering. If none of raw/war/waw are specified, then all are in effect. All relevant memory references preceeding the sync must appear to execute before all relevant memory references following the sync. Order is defined as program order.

Exceptions


None.

Notes

None.

VTLBP

Vector TLB Probe



Assembly

vtlbp

Operation

```
fprintf (stderr, "vtlbp not yet implementedn");
```

Description

This instruction is not yet implemented.

Exceptions


None.

Notes

None.

Vector TLB Read

VTLBR



Assembly

vtlbr

Operation

```
fprintf (stderr, "vtlbr not yet implementedn");
```

Description

This instruction is not yet implemented.

Exceptions


None.

Notes

None.

VTLBWI

Vector TLB Write Indexed



Assembly

vtlbwi

Operation

```
fprintf (stderr, "vtlbwi not yet implementedn");
```

Description

This instruction is not yet implemented.

Exceptions


None.

Notes

None.

Vector TLB Write Random

VTLBWR



Assembly

vtlbwr

Operation

```
fprintf (stderr, "vtlbwr not yet implementedn");
```

Description

This instruction is not yet implemented.

Exceptions


None.

Notes

None.

VTRUNC

Vector Floating-Point Truncate



Assembly

$$\text{vtrunc} \left\{ \begin{matrix} \cdot \text{w} \\ \cdot \text{l} \end{matrix} \right\} \left\{ \begin{matrix} \cdot \text{s} \\ \cdot \text{d} \end{matrix} \right\} [.1] \text{vr}_{\text{dest}}, \text{vr}_{\text{src}}$$

Vector Floating-Point Truncate

VTRUNC(cont.)

Operation (Single to Word)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 2) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = trunc_s_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VTRUNC(cont.)**Vector Floating-Point Truncate****Operation (Double to Word)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = trunc_d_to_w (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Truncate

VTRUNC(cont.)

Operation (Single to Long)

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = trunc_s_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

VTRUNC(cont.)**Vector Floating-Point Truncate****Operation (Double to Long)**

```

if (vl > mvl) {
    Raise vIVL;
}
if (vpw < 3) {
    Raise vIUI;
}
raise_I = raise_U = raise_O = raise_Z = raise_V = raise_E = false;
for (vp = 0; vp < mvl; vp++) {
    I = U = O = Z = V = E = false;
    if (vp < vl && VF[mask][vp]) {
        x = VR[src][vp];
        z = trunc_d_to_l (x, FS, &I, &U, &O, &Z, &V, &E);
        if (!speculative) {
            write = true;
            if (I && I_enabled) { raise_I = true; write = false; }
            if (U && U_enabled) { raise_U = true; write = false; }
            if (O && O_enabled) { raise_O = true; write = false; }
            if (Z && Z_enabled) { raise_Z = true; write = false; }
            if (V && V_enabled) { raise_V = true; write = false; }
            if (E) { raise_E = true; write = false; }
            if (!write) {
                z = VR[dest][vp];
            }
        }
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_I][vp] = I || VF[vfe_I][vp];
    VF[vfe_U][vp] = U || VF[vfe_U][vp];
    VF[vfe_O][vp] = O || VF[vfe_O][vp];
    VF[vfe_Z][vp] = Z || VF[vfe_Z][vp];
    VF[vfe_V][vp] = V || VF[vfe_V][vp];
    VF[vfe_E][vp] = E || VF[vfe_E][vp];
}
if (raise_I || raise_U || raise_O || raise_Z || raise_V || raise_E) {
    Raise vAri;
}

```

Vector Floating-Point Truncate**VTRUNC(cont.)****Description**

Each unmasked VP places into $\mathbf{vr}_{\text{dest}}$ the result of converting \mathbf{vr}_{src} from a floating-point format to a signed integer format, using the *truncate* rounding mode.

Exceptions


\mathbf{vIVL}	Invalid Vector Length Exception
\mathbf{vIUI}	Illegal Use of Instruction Exception

Notes

None.

VXLMADD

Signed Multiply Add Lower Halves



Assembly

$$\text{vxlmadd} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

Signed Multiply Add Lower Halves

VXLMADD(cont.)

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = signed_multiply_lower_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
        z = signed_saturating_add (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VXLMADD(cont.)**Signed Multiply Add Lower Halves****Description**

Each unmasked VP computes the signed integer product of the lower halves of $\mathbf{vs}_{\text{src1}}$ / $\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. This result is added into $\mathbf{vr}_{\text{dest}}$ after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\mathbf{vc}_{\text{vshamt}}$.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

If the source registers do not contain valid half-width signed integers, then the operation is not sensible.

Unsigned Multiply Add Lower Halves

VXLMADD.U



Assembly

$$\text{vxlmadd.u} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

VXLMADD.U(cont.)**Unsigned Multiply Add Lower Halves****Operation**

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = unsigned_multiply_lower_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
        z = unsigned_saturating_add (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Unsigned Multiply Add Lower Halves

VXLMADD.U(cont.)

Description

Each unmasked VP computes the unsigned integer product of the lower halves of $\mathbf{vs}_{src1}/\mathbf{vr}_{src1}$ and \mathbf{vr}_{src2} . This result is added into \mathbf{vr}_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from \mathbf{vc}_{vshamt} .

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VXLMSUB

Signed Multiply Subtract Lower Halves



Assembly

$$\text{vxlmsub} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

Signed Multiply Subtract Lower Halves

VXLMSUB(cont.)

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = signed_multiply_lower_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
        z = signed_saturating_sub (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VXLMSUB(cont.)**Signed Multiply Subtract Lower Halves****Description**

Each unmasked VP computes the signed integer product of the lower halves of $\mathbf{v}_{\mathbf{s}_{\text{src1}}}/\mathbf{v}_{\mathbf{r}_{\text{src1}}}$ and $\mathbf{v}_{\mathbf{r}_{\text{src2}}}$. This result is subtracted from $\mathbf{v}_{\mathbf{r}_{\text{dest}}}$ after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\mathbf{v}_{\mathbf{c}_{\text{vshamt}}}$.

Exceptions


\mathbf{v}_{IVL} Invalid Vector Length Exception

Notes

If the source registers do not contain valid half-width signed integers, then the operation is not sensible.

Unsigned Multiply Subtract Lower Halves

VXLMSUB.U



Assembly

$$\text{vxlmsub.u} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

VXLMSUB.U(cont.)**Unsigned Multiply Subtract Lower Halves****Operation**

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = unsigned_multiply_lower_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
        z = unsigned_saturating_sub (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Unsigned Multiply Subtract Lower Halves

VXLMSUB.U(cont.)

Description

Each unmasked VP computes the unsigned integer product of the lower halves of $\mathbf{vs}_{src1}/\mathbf{vr}_{src1}$ and \mathbf{vr}_{src2} . This result is subtracted from \mathbf{vr}_{dest} after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from \mathbf{vc}_{vshamt} .

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VXLMUL

Signed Multiply Lower Halves



Assembly

$$\text{vxlmul} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = signed_multiply_lower_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP computes the signed integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after an arithmetic right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

Signed Multiply Lower Halves**VXLMUL(cont.)****Exceptions**


v_{IVL} Invalid Vector Length Exception

Notes

If the source registers do not contain valid half-width signed integers, then the operation is not sensible.

VXLMUL.U

Unsigned Multiply Lower Halves



Assembly

$$\text{vxlmul.u} \begin{cases} \text{.vv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = unsigned_multiply_lower_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP computes the unsigned integer product of the lower halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after a logical right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

Unsigned Multiply Lower Halves**VXLMUL.U(cont.)****Exceptions**


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VXOR

Vector Xor



Assembly

$$\text{vxor} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

if (vl > mvl) {
    Raise vIVL;
}
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = x ^ y;
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each VP writes into vr_{dest} the bit-wise logical *xor* of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} .

Exceptions


vIVL Invalid Vector Length Exception

Notes

None.

Signed Multiply Add Upper Halves

VXUMADD



Assembly

$$\text{vxumadd} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

VXUMADD(cont.)**Signed Multiply Add Upper Halves****Operation**

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = signed_multiply_upper_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
        z = signed_saturating_add (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Signed Multiply Add Upper Halves
VXUMADD(cont.)**Description**


Each unmasked VP computes the signed integer product of the upper halves of $\mathbf{vs}_{src1}/\mathbf{vr}_{src1}$ and \mathbf{vr}_{src2} . This result is added into \mathbf{vr}_{dest} after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from \mathbf{vc}_{vshamt} .

Exceptions

\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VXUMADD.U
Unsigned Multiply Add Upper Halves**Assembly**

$$\text{vxumadd.u} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

Unsigned Multiply Add Upper Halves

VXUMADD.U(cont.)

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = unsigned_multiply_upper_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
        z = unsigned_saturating_add (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VXUMADD.U(cont.)**Unsigned Multiply Add Upper Halves****Description**

Each unmasked VP computes the unsigned integer product of the upper halves of $\mathbf{vs}_{\text{src1}}$ / $\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. This result is added into $\mathbf{vr}_{\text{dest}}$ after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\mathbf{vc}_{\text{vshamt}}$.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Signed Multiply Subtract Upper Halves

VXUMSUB



Assembly

$$\text{vxumsub} \begin{cases} \text{.vv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vr}_{\text{src1}} , \text{vr}_{\text{src2}} \\ \text{.sv}[\text{.1}] & \text{vr}_{\text{dest}} , \text{vs}_{\text{src1}} , \text{vr}_{\text{src2}} \end{cases}$$

VXUMSUB(cont.)**Signed Multiply Subtract Upper Halves****Operation**

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = signed_multiply_upper_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
        z = signed_saturating_sub (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

Signed Multiply Subtract Upper Halves

VXUMSUB(cont.)

Description


Each unmasked VP computes the signed integer product of the upper halves of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. This result is subtracted from $\mathbf{vr}_{\text{dest}}$ after an arithmetic right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\mathbf{vc}_{\text{vshamt}}$.

Exceptions

\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

VXUMSUB.U**Unsigned Multiply Subtract Upper Halves****Assembly**

$$\text{vxumsub.u} \begin{cases} \text{.vv}[\cdot] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Unsigned Multiply Subtract Upper Halves

VXUMSUB.U(cont.)

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
raise_S = false;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        w = VR[src3][vp];
        z = unsigned_multiply_upper_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
        z = unsigned_saturating_sub (w, z, &saturated);
        if (saturated) {
            f = true;
            if (!speculative && S_enabled) {
                raise_S = true;
                z = VR[dest][vp]; /* don't write result */
            }
        } else {
            f = VF[vfe_S][vp];
        }
    } else {
        z = VR[dest][vp];
        f = VF[vfe_S][vp];
    }
    VR[dest][vp] = z;
    VF[vfe_S][vp] = f;
}
if (raise_S) {
    Raise vAri;
}

```

VXUMSUB.U(cont.)**Unsigned Multiply Subtract Upper Halves****Description**

Each unmasked VP computes the unsigned integer product of the upper halves of $\mathbf{vs}_{\text{src1}}/\mathbf{vr}_{\text{src1}}$ and $\mathbf{vr}_{\text{src2}}$. This result is subtracted from $\mathbf{vr}_{\text{dest}}$ after a logical right shift and fixed-point round. The final result is saturated. The shift amount is taken from $\mathbf{vc}_{\text{vshamt}}$.

Exceptions


\mathbf{vIVL} Invalid Vector Length Exception

Notes

None.

Signed Multiply Upper Halves

VXUMUL



Assembly

$$\text{vxumul} \begin{cases} \text{.vv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv}[\cdot 1] & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```

assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = signed_multiply_upper_halves (x, y);
        z = signed_shift_right_and_round (z, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}

```

Description

Each unmasked VP computes the signed integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after an arithmetic right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

VXUMUL(cont.)**Signed Multiply Upper Halves****Exceptions**


v_{IVL} Invalid Vector Length Exception

Notes

None.

Unsigned Multiply Upper Halves

VXUMUL.U



Assembly

$$\text{vxumul.u} \begin{cases} \text{.vv[.1]} & \text{vr}_{\text{dest}}, \text{vr}_{\text{src1}}, \text{vr}_{\text{src2}} \\ \text{.sv[.1]} & \text{vr}_{\text{dest}}, \text{vs}_{\text{src1}}, \text{vr}_{\text{src2}} \end{cases}$$

Operation

```
assert (RM == TRC || RM == UP || RM == NRE || RM == JAM);
if (vl > mvl) {
    Raise vIVL;
}
shamt_mask = (1 << (vpw + 3)) - 1;
shamt = vshamt & shamt_mask;
for (vp = 0; vp < mvl; vp++) {
    if (vp < vl && VF[mask][vp]) {
        if (vv) {
            x = VR[src1][vp];
            y = VR[src2][vp];
        } else {
            x = VS[src1];
            y = VR[src2][vp];
        }
        z = unsigned_multiply_upper_halves (x, y);
        z = unsigned_shift_right_and_round (z, shamt, RM);
    } else {
        z = VR[dest][vp];
    }
    VR[dest][vp] = z;
}
```

Description

Each unmasked VP computes the unsigned integer product of the upper halves of $\text{vs}_{\text{src1}}/\text{vr}_{\text{src1}}$ and vr_{src2} . This result is written into vr_{dest} after a logical right shift and fixed-point round. The shift amount is taken from $\text{vc}_{\text{vshamt}}$.

VXUMUL.U(cont.)**Unsigned Multiply Upper Halves****Exceptions**

v_{IVL} Invalid Vector Length Exception

Notes

None.